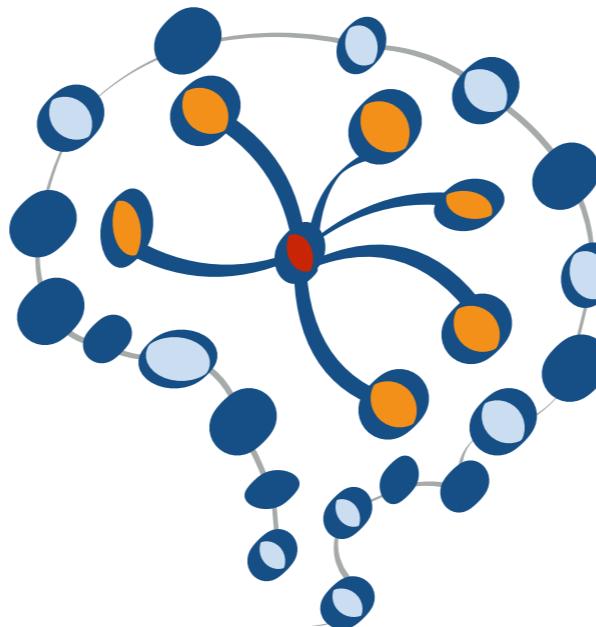


# STAT 453: Introduction to Deep Learning and Generative Models

Sebastian Raschka

<http://stat.wisc.edu/~sraschka/teaching>



## Lecture 06

# Automatic Differentiation with PyTorch

# Today

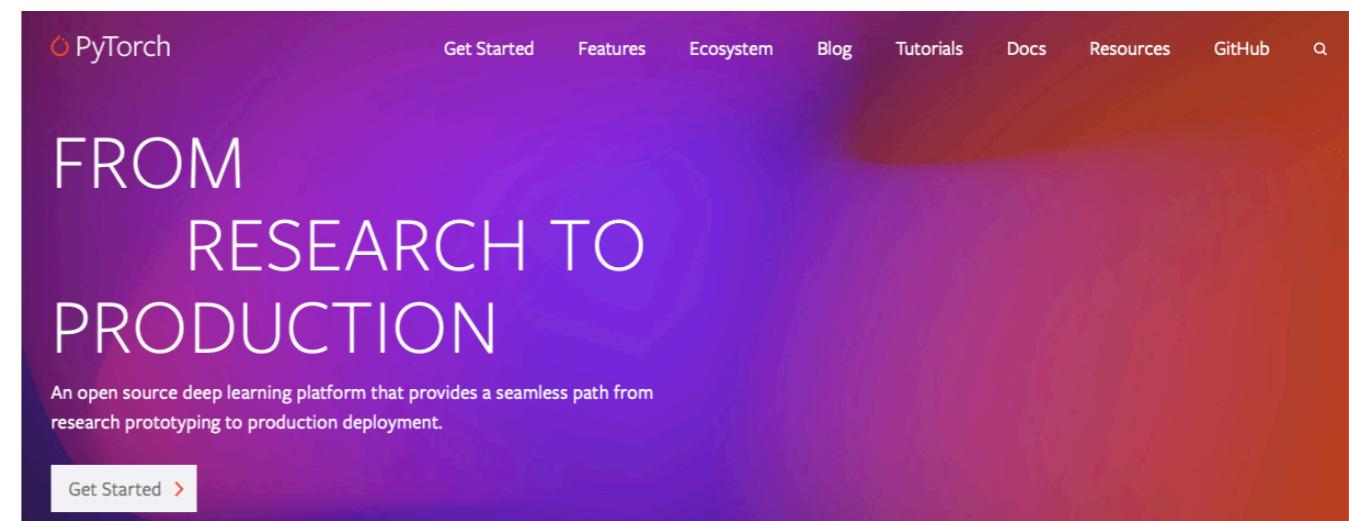
Computing partial derivatives more easily  
(and automatically) with PyTorch

# Lecture Overview

1. PyTorch Resources
2. Computation Graphs
3. Automatic Differentiation in PyTorch
4. Training ADALINE Manually Vs Automatically in PyTorch
5. A Closer Look at the PyTorch API

# Learning More About PyTorch

- 1. PyTorch Resources**
2. Computation Graphs
3. Automatic Differentiation in PyTorch
4. Training ADALINE Manually Vs Automatically in PyTorch
5. A Closer Look at the PyTorch API

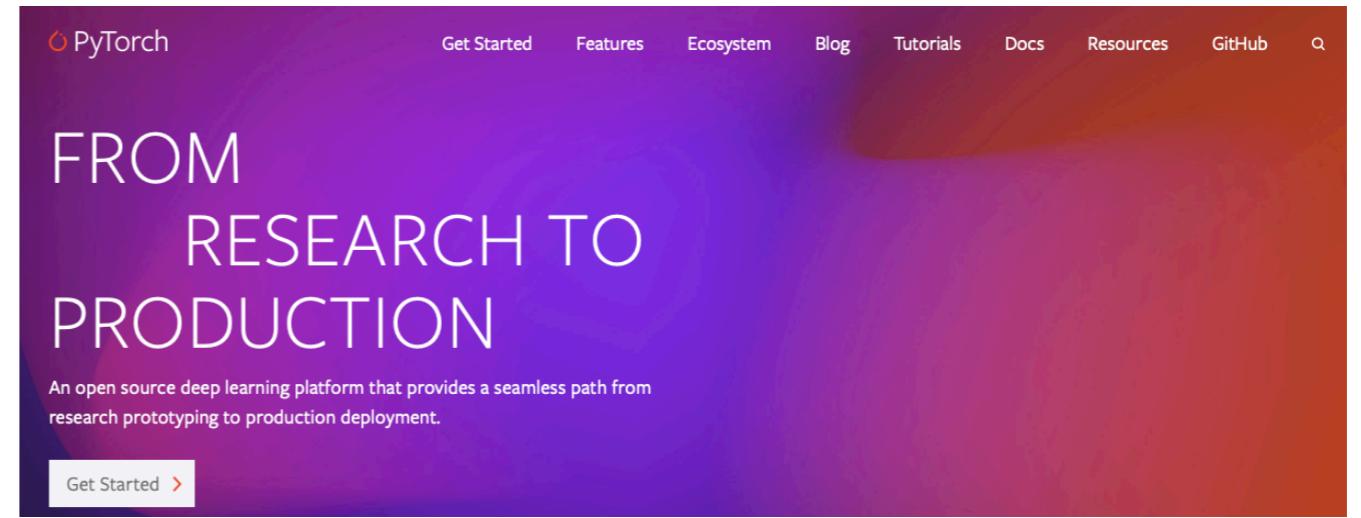


The screenshot shows the official PyTorch website. At the top, there's a navigation bar with links for "Get Started", "Features", "Ecosystem", "Blog", "Tutorials", "Docs", "Resources", "GitHub", and a search icon. The main title "PyTorch" is at the top left. The central heading reads "FROM RESEARCH TO PRODUCTION" in large white capital letters. Below it, a subtitle says "An open source deep learning platform that provides a seamless path from research prototyping to production deployment." A "Get Started" button is visible at the bottom left of the main content area.

<https://pytorch.org/>

## At a Glance:

- Based on Torch 7, which was based on Lua and inspired by Lush
- PyTorch started in 2016
- Focuses on flexibility and minimizing cognitive overhead
- Dynamic nature of autograd API inspired by Chainer
- Core features
  - Automatic differentiation
  - Dynamic computation graphs
  - NumPy integration
- written in C++ and CUDA (CUDA is like C++ for the GPU)
- Python is the usability glue



The banner features the PyTorch logo at the top left. To its right is a navigation bar with links: Get Started, Features, Ecosystem, Blog, Tutorials, Docs, Resources, GitHub, and a search icon. The main title 'FROM RESEARCH TO PRODUCTION' is centered in large white capital letters. Below it is a subtitle: 'An open source deep learning platform that provides a seamless path from research prototyping to production deployment.' At the bottom is a 'Get Started' button with a right-pointing arrow.

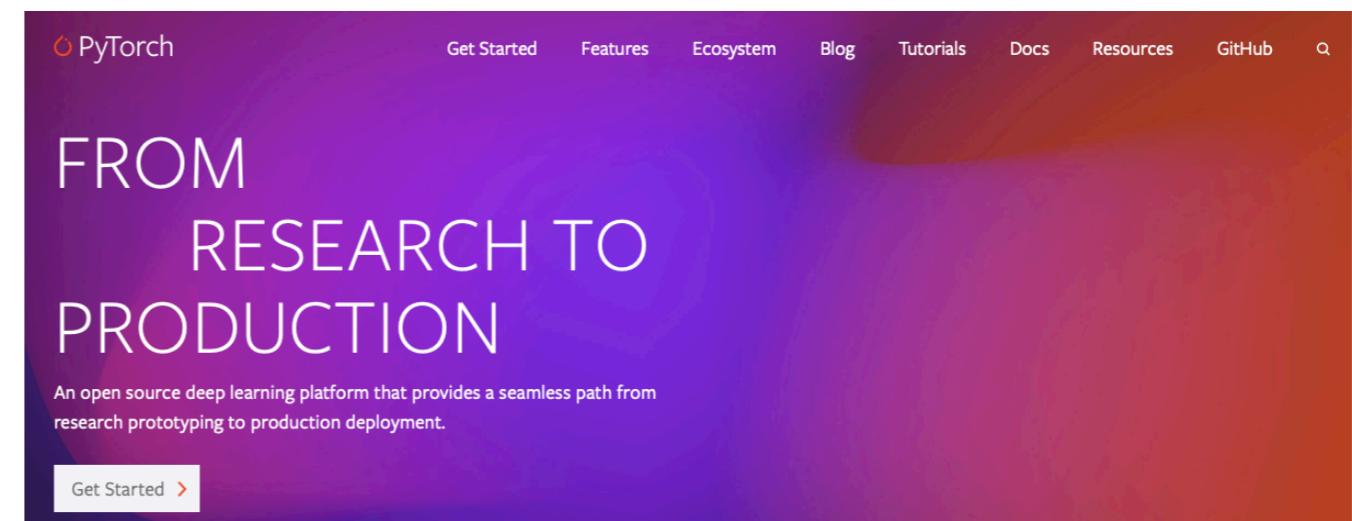
<https://pytorch.org/>

## At a Glance:

- Based on Torch 7, which was based on Lua and inspired by Lush
- PyTorch started in 2016
- Focuses on flexibility and minimizing cognitive overhead
- Dynamic nature of autograd API inspired by Chainer
- Core features
  - Automatic differentiation
  - Dynamic computation graphs
  - NumPy integration
- written in C++ and CUDA (CUDA is optional)
- Python is the usability glue

### PyTorch vs NumPy

- Support GPU
- distribute ops across multiple devices
- keep track of computation graph and ops that created them



The screenshot shows the official PyTorch website. At the top, there's a navigation bar with links for "Get Started", "Features", "Ecosystem", "Blog", "Tutorials", "Docs", "Resources", "GitHub", and a search icon. The main title "PyTorch" is at the top left. Below it, the text "FROM RESEARCH TO PRODUCTION" is displayed in large white letters. A subtitle below reads "An open source deep learning platform that provides a seamless path from research prototyping to production deployment." A "Get Started" button is visible at the bottom left of the main content area.

<https://pytorch.org/>

## At a Glance:

- Based on Torch 7, which was based on Lua and inspired by Lush
- PyTorch started in 2016
- Focuses on flexibility and minimizing cognitive overhead
- Dynamic nature of autograd API inspired by Chainer
- Core features
  - Automatic differentiation

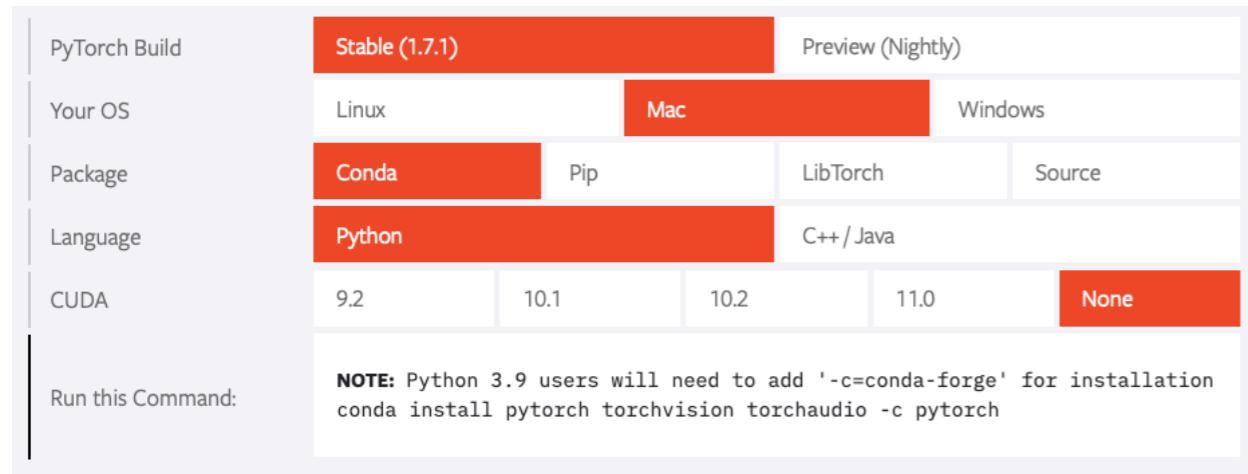
**"the speedup gained by taking Python out of  
the computation is 10% or less"**

**-- Stevens *et al.*: Deep Learning with PyTorch**

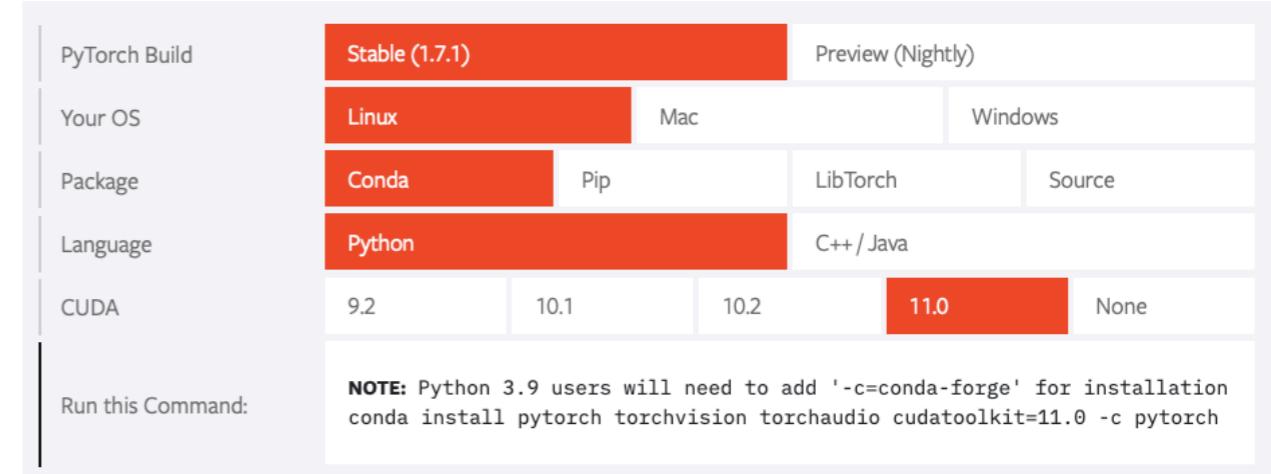
- written in C++ and CUDA (CUDA is like C++ for the GPU)
- Python is the usability glue

# Installation

## Recommendation for Laptop (e.g., MacBook)



## Recommendation for Desktop (Linux) with GPU



<https://pytorch.org/>

## As mention in the installation tips on Canvas

And don't forget that you import PyTorch as "import torch," not "import pytorch" :)

```
[In [1]: import torch

[In [2]: torch.__version__
Out[2]: '1.7.0'

In [3]: ]
```

# Many Useful Tutorials (recommend that you read some of them)

The screenshot shows the PyTorch Tutorials homepage. At the top, there is a navigation bar with links: Get Started, Ecosystem, Mobile, Blog, Tutorials (which is highlighted in red), Docs (with a dropdown arrow), Resources (with a dropdown arrow), and Github. Below the navigation bar, the version '1.7.1' is displayed. On the left side, there is a sidebar with sections: PyTorch Recipes (Search Tutorials button), Learning PyTorch (Deep Learning with PyTorch: A 60 Minute Blitz, Learning PyTorch with Examples, What is `torch.nn` really?, Visualizing Models, Data, and Training with TensorBoard), and Image/Video (TorchVision Object Detection Finetuning Tutorial, Transfer Learning for Computer Vision Tutorial, Adversarial Example Generation). The main content area features a large title 'WELCOME TO PYTORCH TUTORIALS'. Below it, there are two main sections: 'New to PyTorch?' and 'PyTorch Recipes'. The 'New to PyTorch?' section includes a description of the 60-minute blitz tutorial and a button to 'Start 60-min blitz'. The 'PyTorch Recipes' section includes a description of bite-size examples and a button to 'Explore Recipes'. At the bottom of the main content area, there are several category buttons: All, Audio, Best Practice, C++, CUDA, Frontend APIs, Getting Started, Image/Video, Interpretability, and Memory Format.

<https://pytorch.org/tutorials/>

# Many Useful Tutorials (recommend that you read some of them)



Get Started   Ecosystem   Mobile   Blog   **Tutorials**   Docs ▾   Resources ▾   Github

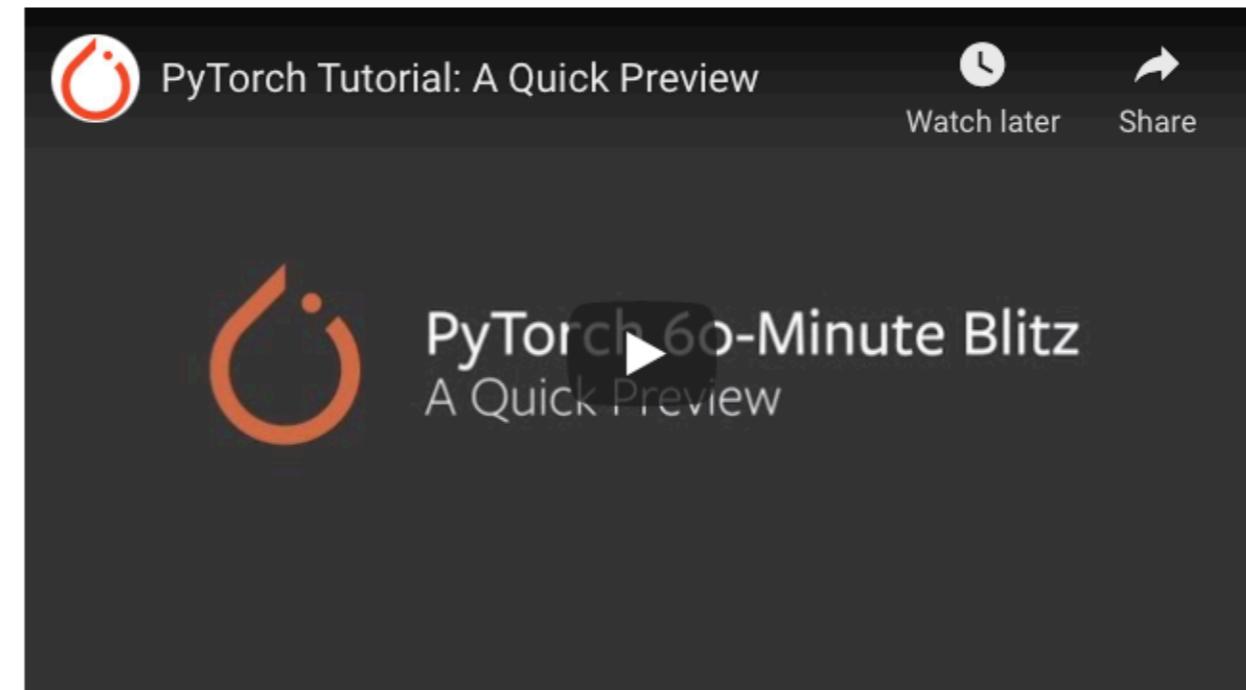
1.7.1

[Tutorials](#) > Deep Learning with PyTorch: A 60 Minute Blitz

[Shortcuts](#)

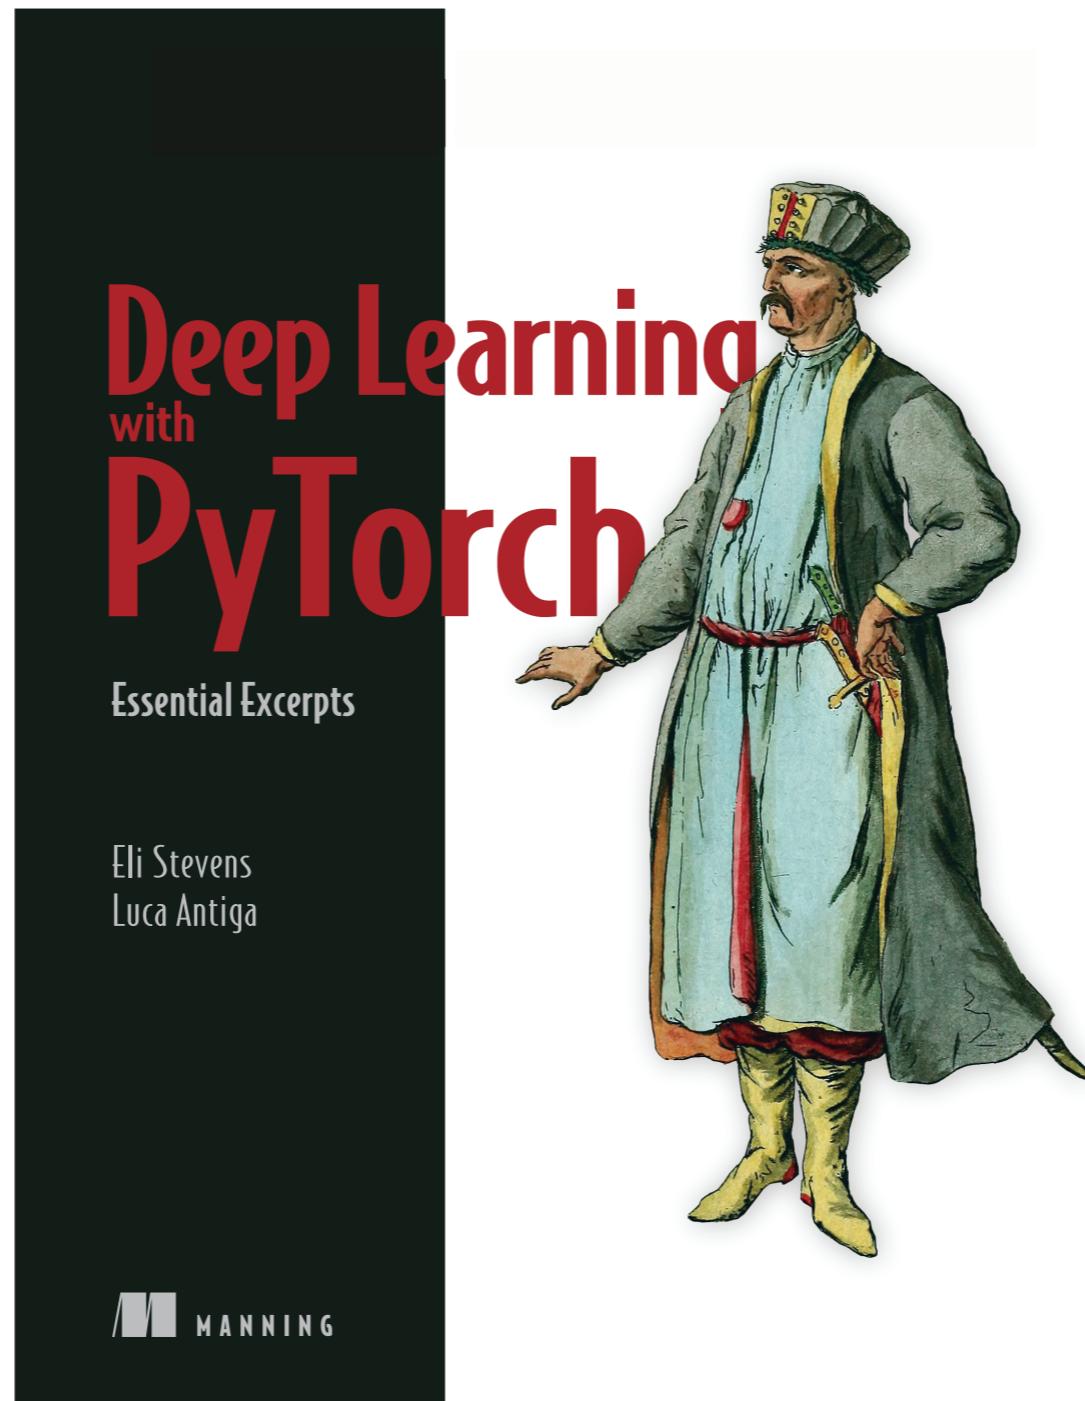
## DEEP LEARNING WITH PYTORCH: A 60 MINUTE BLITZ

**Author:** Soumith Chintala



### What is PyTorch?

[https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)



<https://pytorch.org/assets/deep-learning/Deep-Learning-with-PyTorch.pdf>

# Very Active & Friendly Community and Help/Discussion Forum

PyTorch

Do you want live notifications when people reply to your posts? [Enable Notifications](#) X

all categories ▾ all ▾ **Latest** New (47) Unread (104) Top Categories + New Topic

Topic		Replies	Views	Activity
<input checked="" type="checkbox"/> Using MSELoss instead of CrossEntropy for Ordinal Regression/Classification vision		2	83	1h
Optimizer.load_state_dict() weird behaviour with Adam optimizer vision		7	2.0k	1h
Is there a way to train 3 dataloaders using multiprocessing? •		0	11	2h
<input checked="" type="checkbox"/> Getting different feature vectors from frozen layers after training vision		5	86	2h
<input checked="" type="checkbox"/> Libtorch_cuda.so is too large (>2GB) deployment		22	346	2h
Undo pruning - How to 'unmask' pruned weights • vision		0	8	2h
If input.dim() == 2 and bias is not None: AttributeError: 'tuple' object has no attribute 'dim' •		3	41	2h
Export unsupported/compound ops to ONNX • deployment		0	9	2h

<https://discuss.pytorch.org>

# Understanding Automatic Differentiation via Computation Graphs

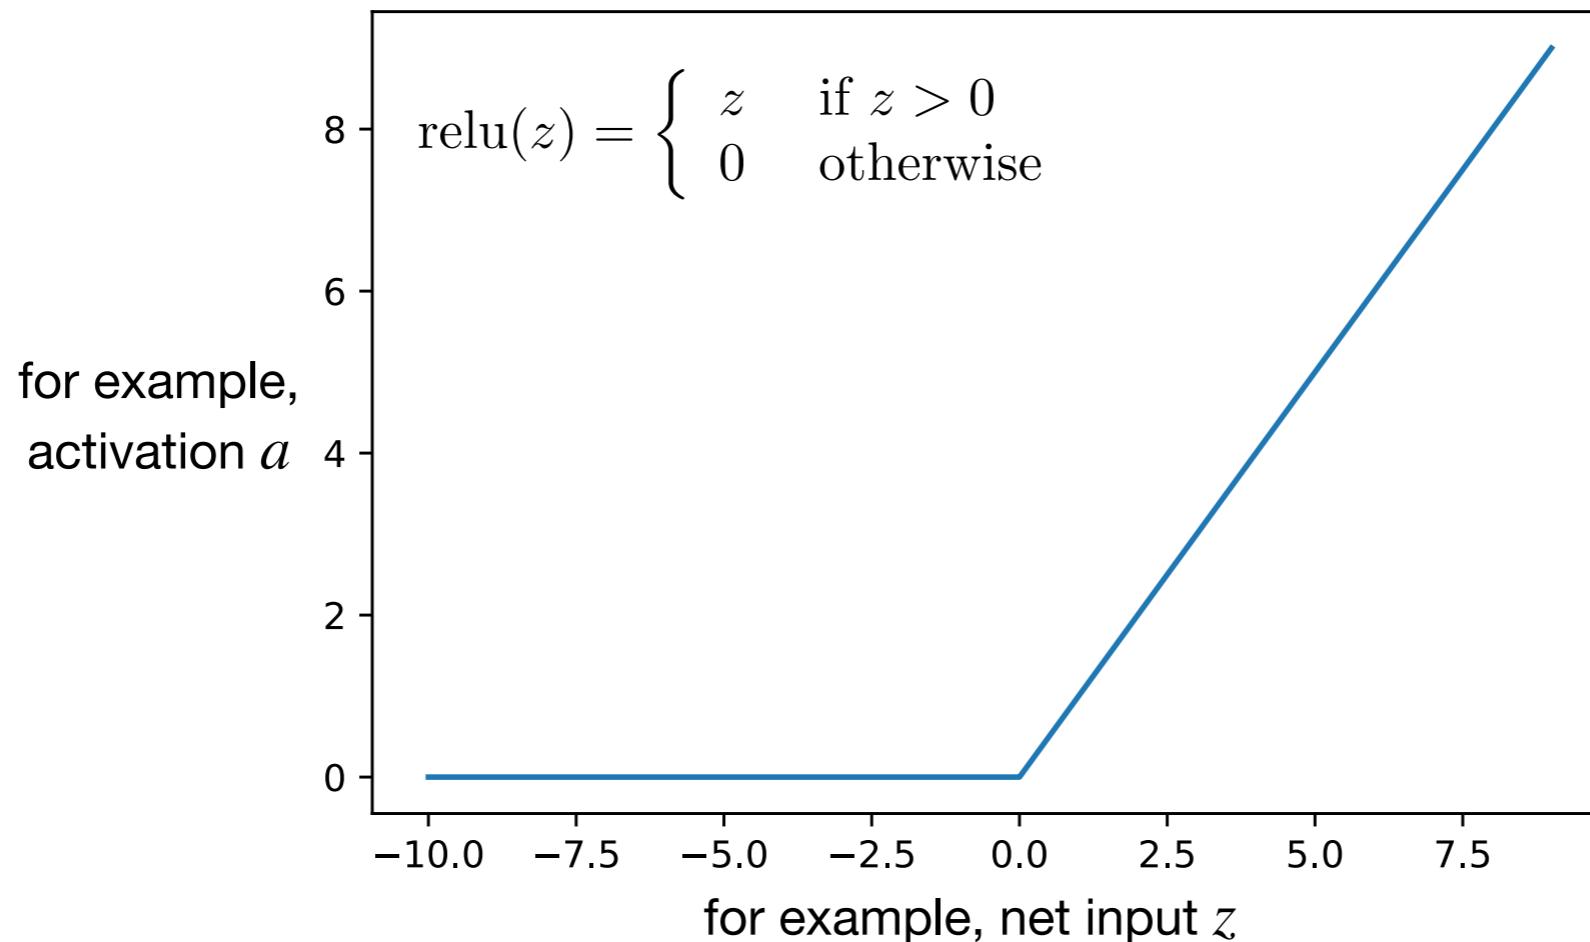
1. PyTorch Resources
2. **Computation Graphs**
3. Automatic Differentiation in PyTorch
4. Training ADALINE Manually Vs Automatically in PyTorch
5. A Closer Look at the PyTorch API

In the context of deep learning (and PyTorch)  
it is helpful to think about neural networks  
as computation graphs

# Computation Graphs

Suppose we have the following activation function:

$$a(x, w, b) = \text{relu}(w \cdot x + b)$$



**ReLU = Rectified Linear Unit**  
(prob. the most commonly used activation function in DL)

# Side-note about ReLU Function

You may note that

$$\sigma'(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z > 0 \\ \text{DNE} & \text{if } z = 0 \end{cases}$$

But in the machine learning--computer science context, for convenience, we can just say

$$\sigma'(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ 1 & \text{if } z > 0 \end{cases}$$

Why not differentiable?

Derivative does not exist (DNE) at 0, because the derivative is different if we approach the limit from the left or right:

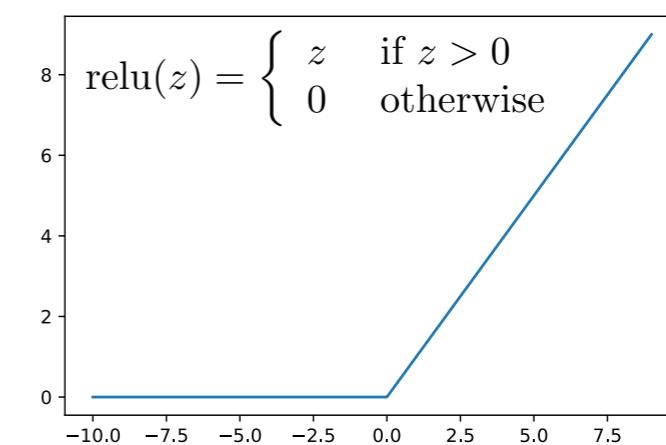
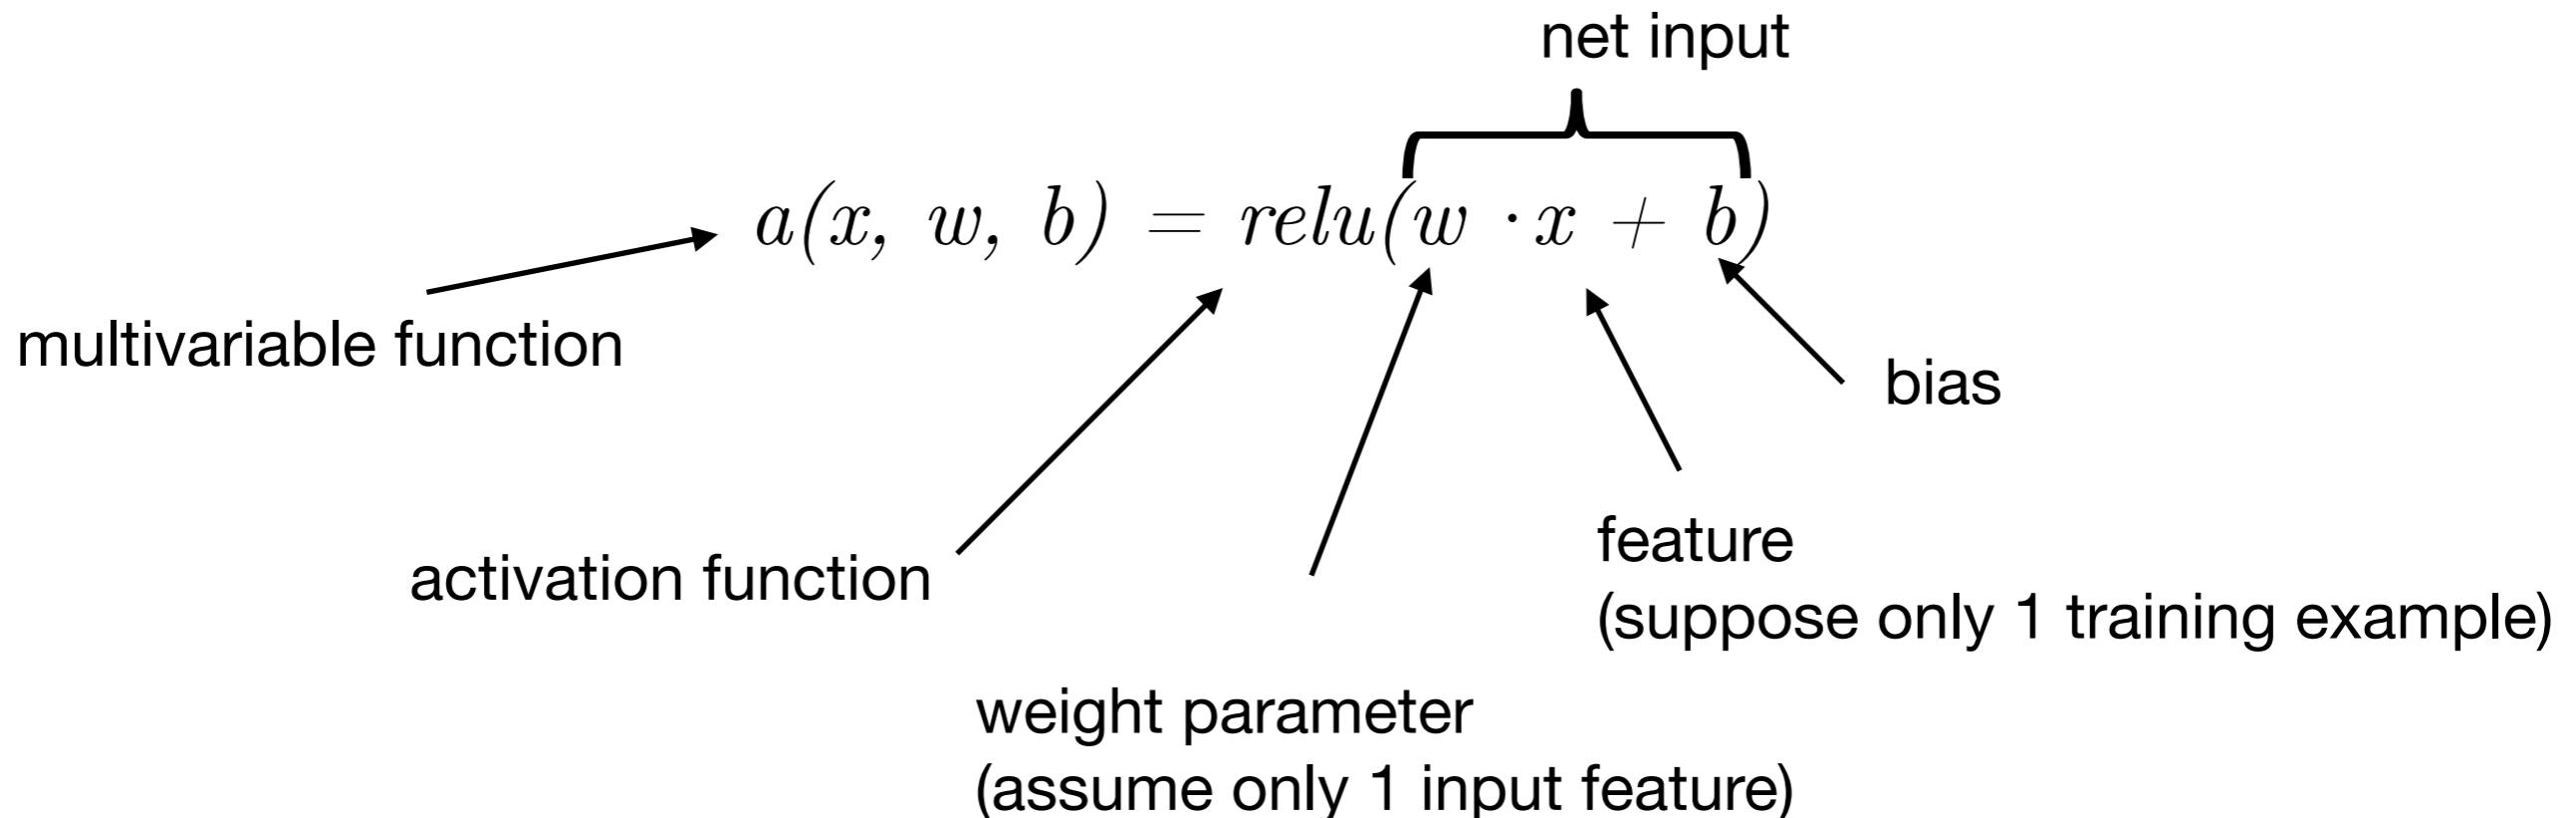
$$\sigma'(z) = \lim_{z \rightarrow 0} \frac{\max(0, z + \Delta z) - \max(0, z)}{\Delta z}$$

$$\sigma'(0) = \lim_{z \rightarrow 0^+} \frac{0 + \Delta z - 0}{\Delta z} = 1$$

$$\sigma'(0) = \lim_{z \rightarrow 0^-} \frac{0 - 0}{\Delta z} = 0$$

# Computation Graphs

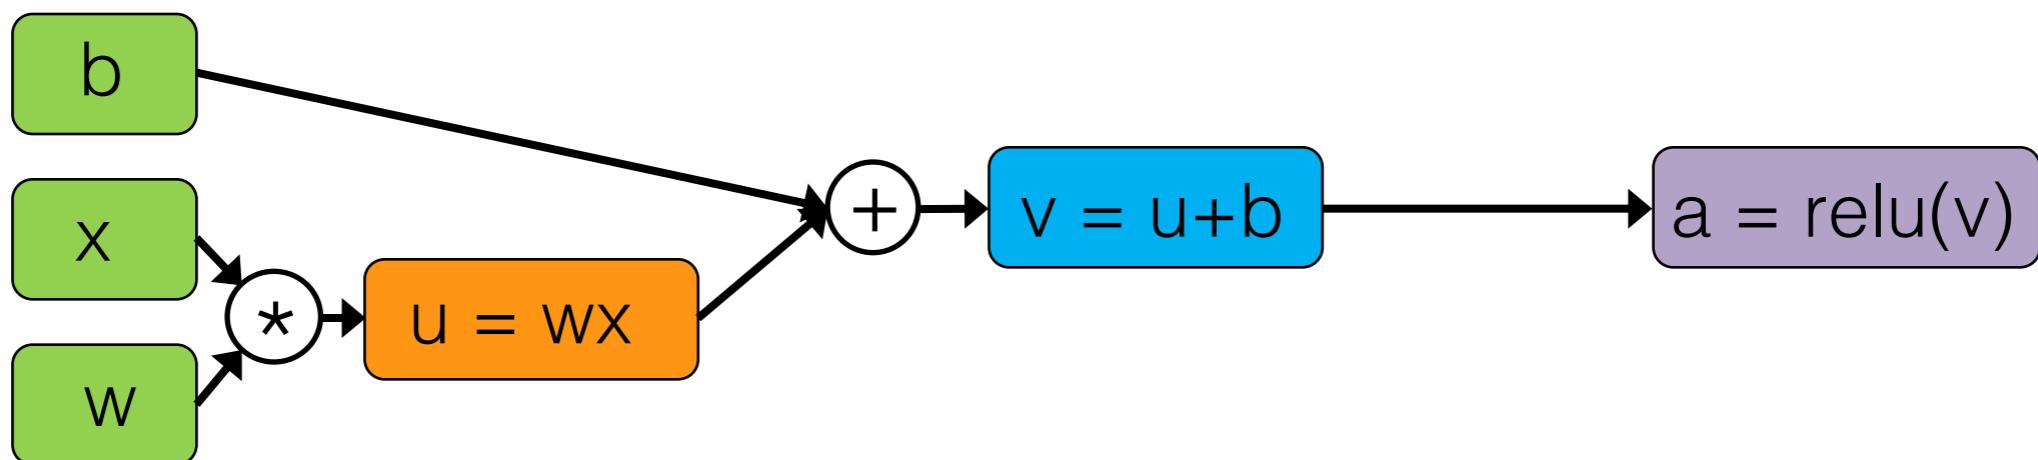
Suppose we have the following activation function:



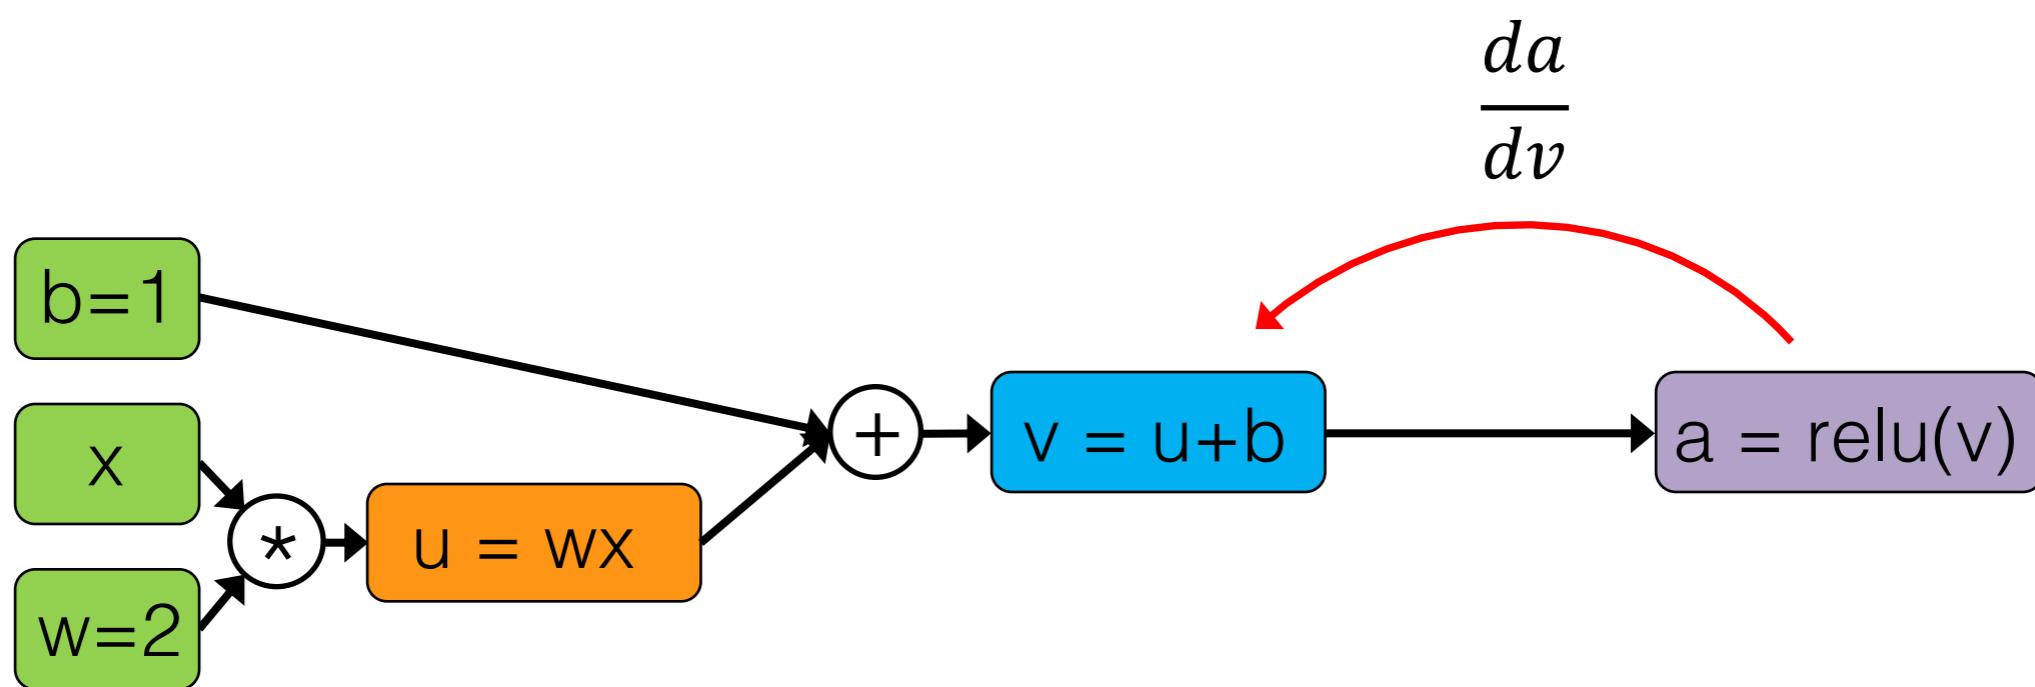
# Computation Graphs

$$a(x, w, b) = \text{relu}(w \cdot x + b)$$

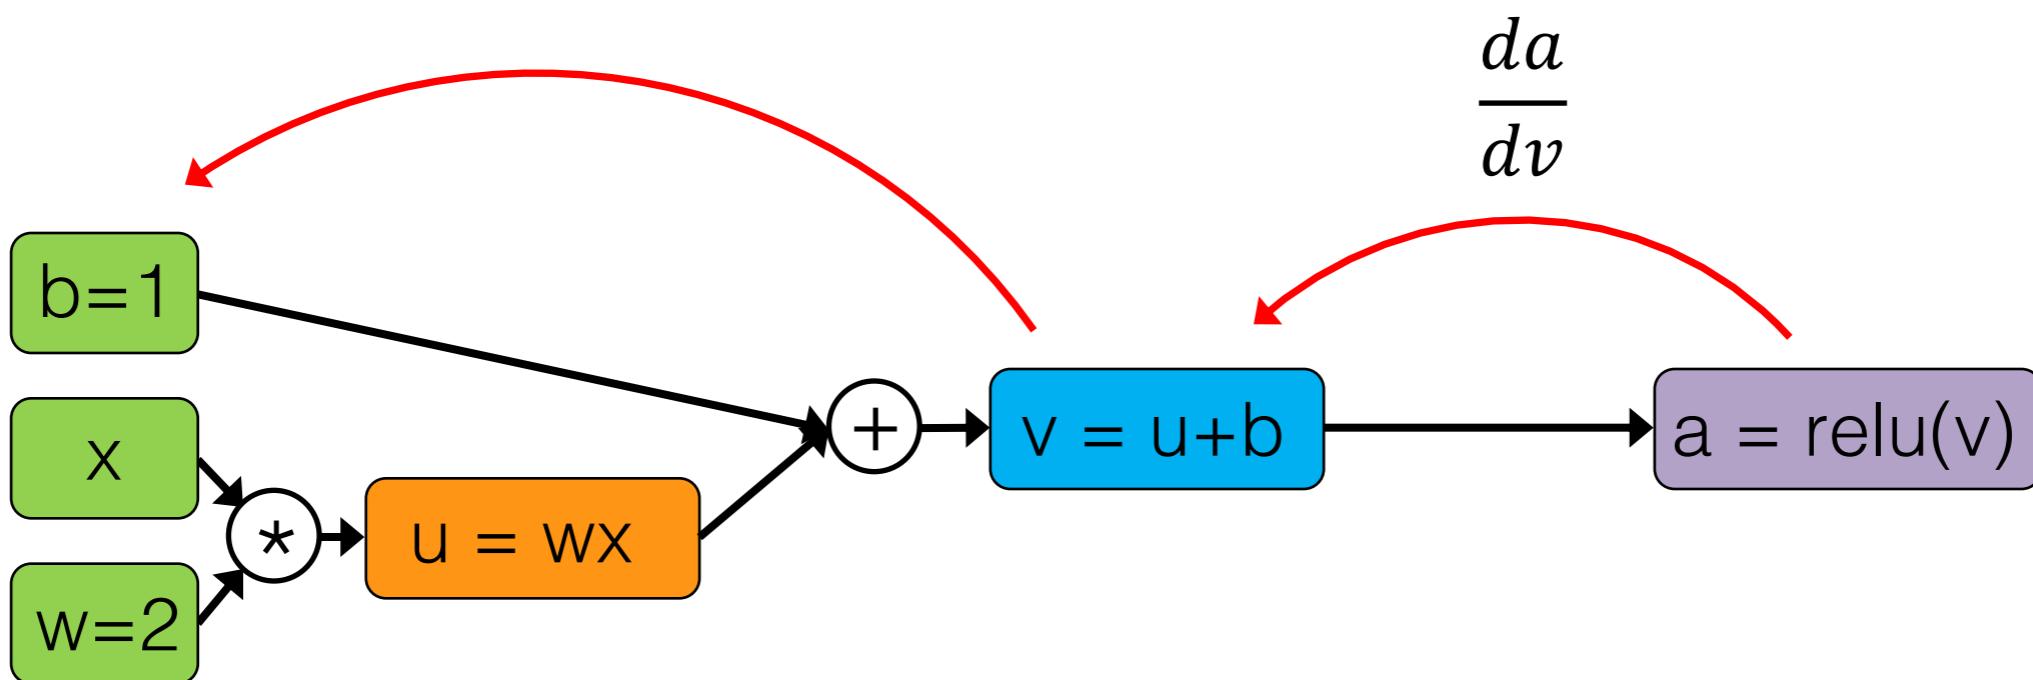
$\underbrace{\phantom{w \cdot x}}$   
 $u$   
 $\underbrace{\phantom{w \cdot x + b}}$   
 $v$



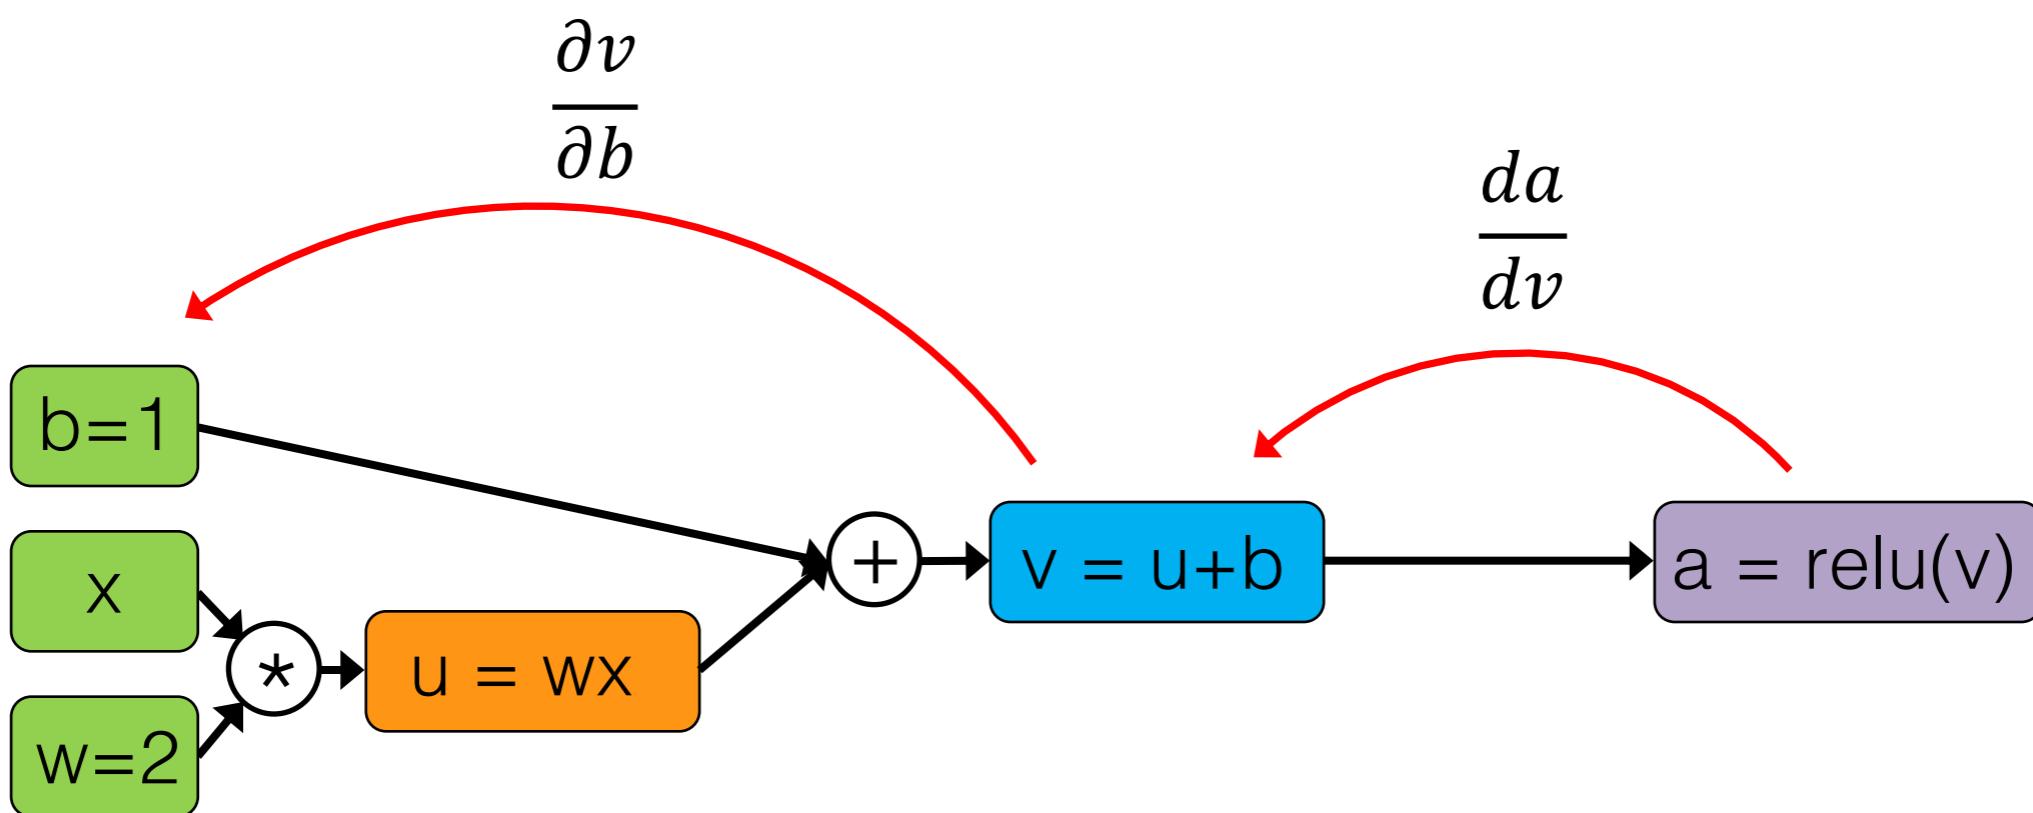
# Computation Graphs



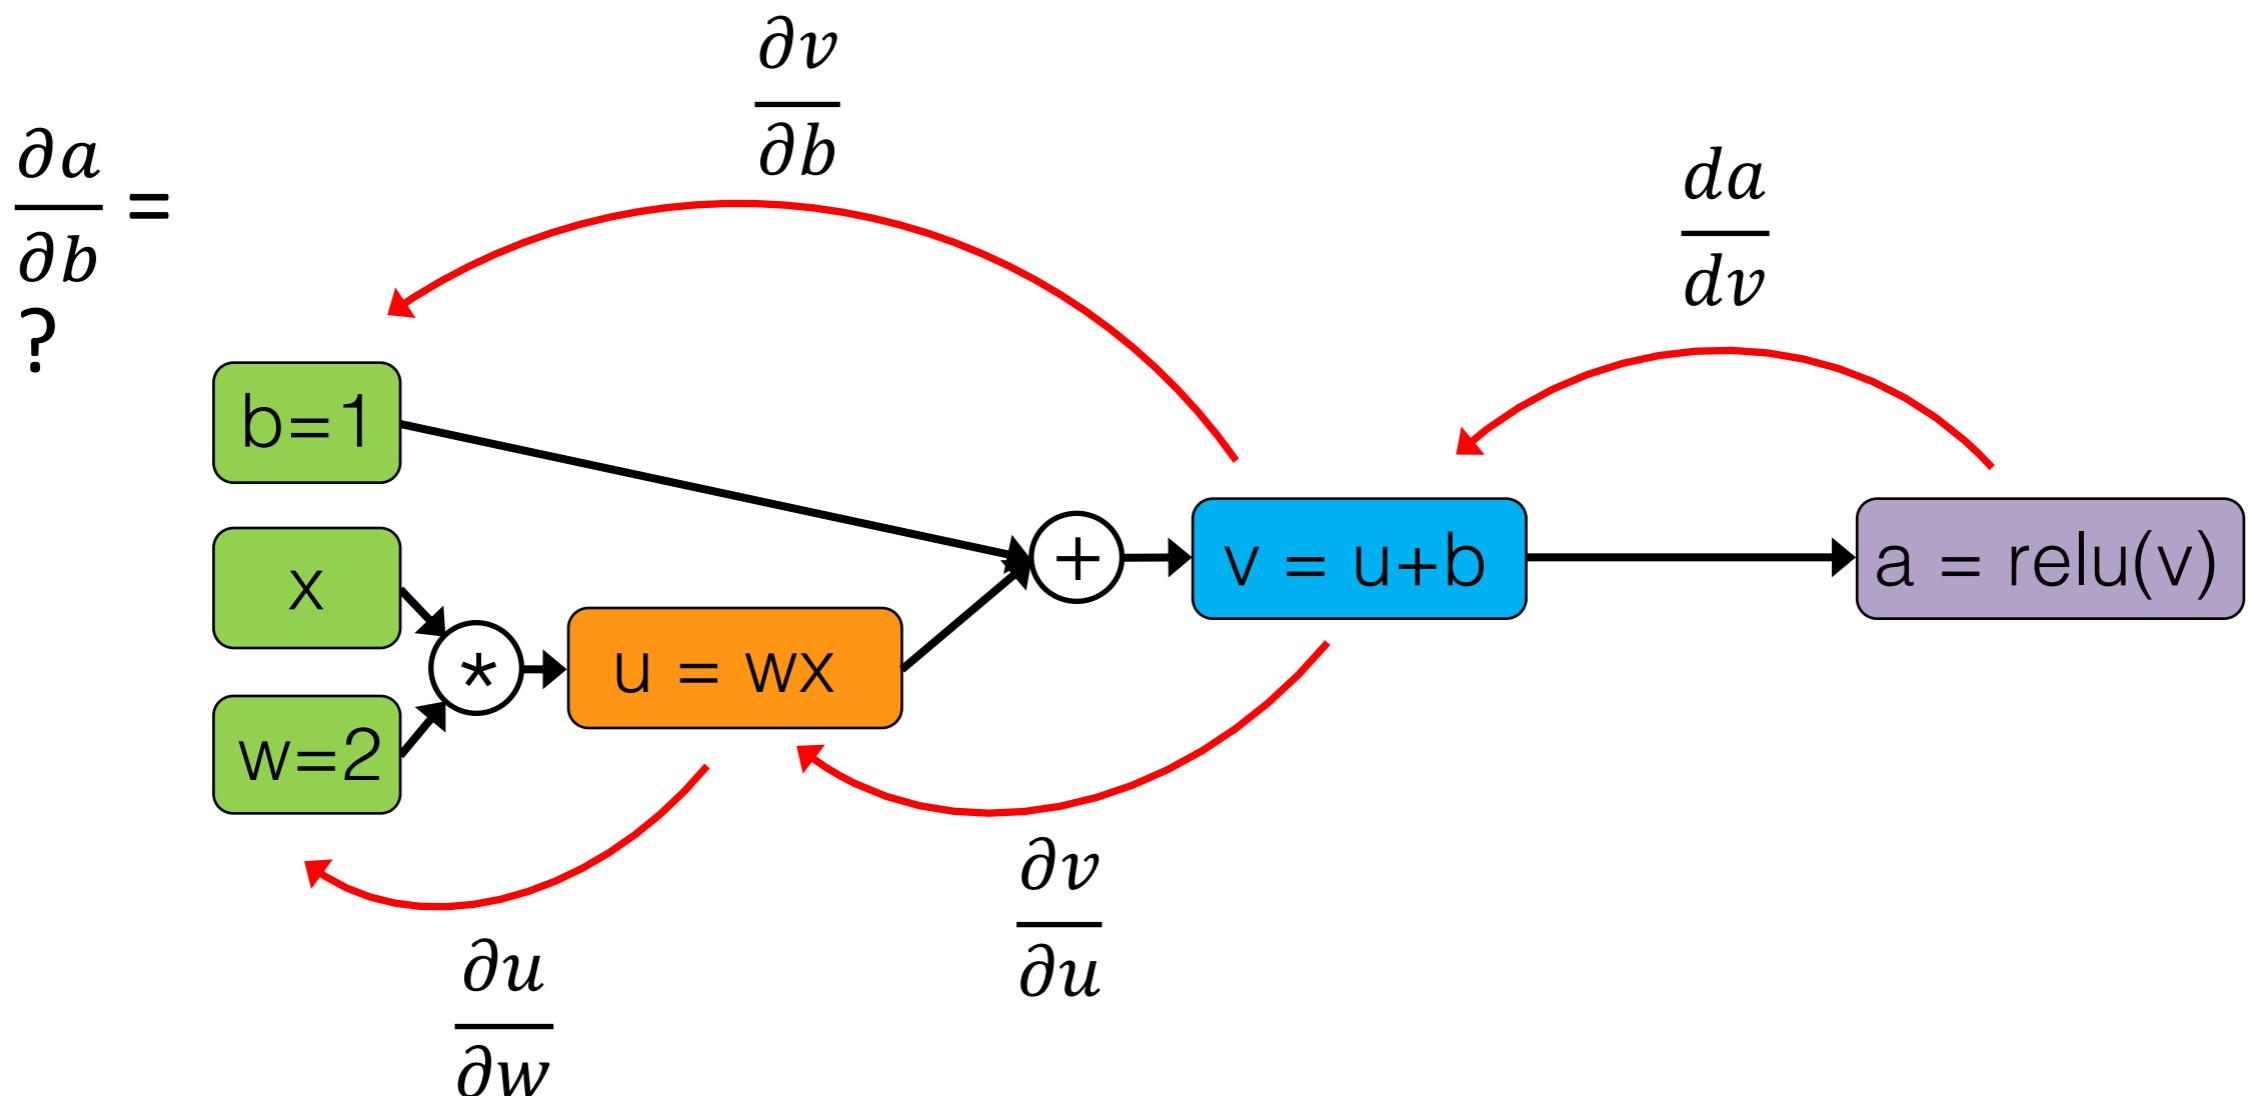
# Computation Graphs



# Computation Graphs

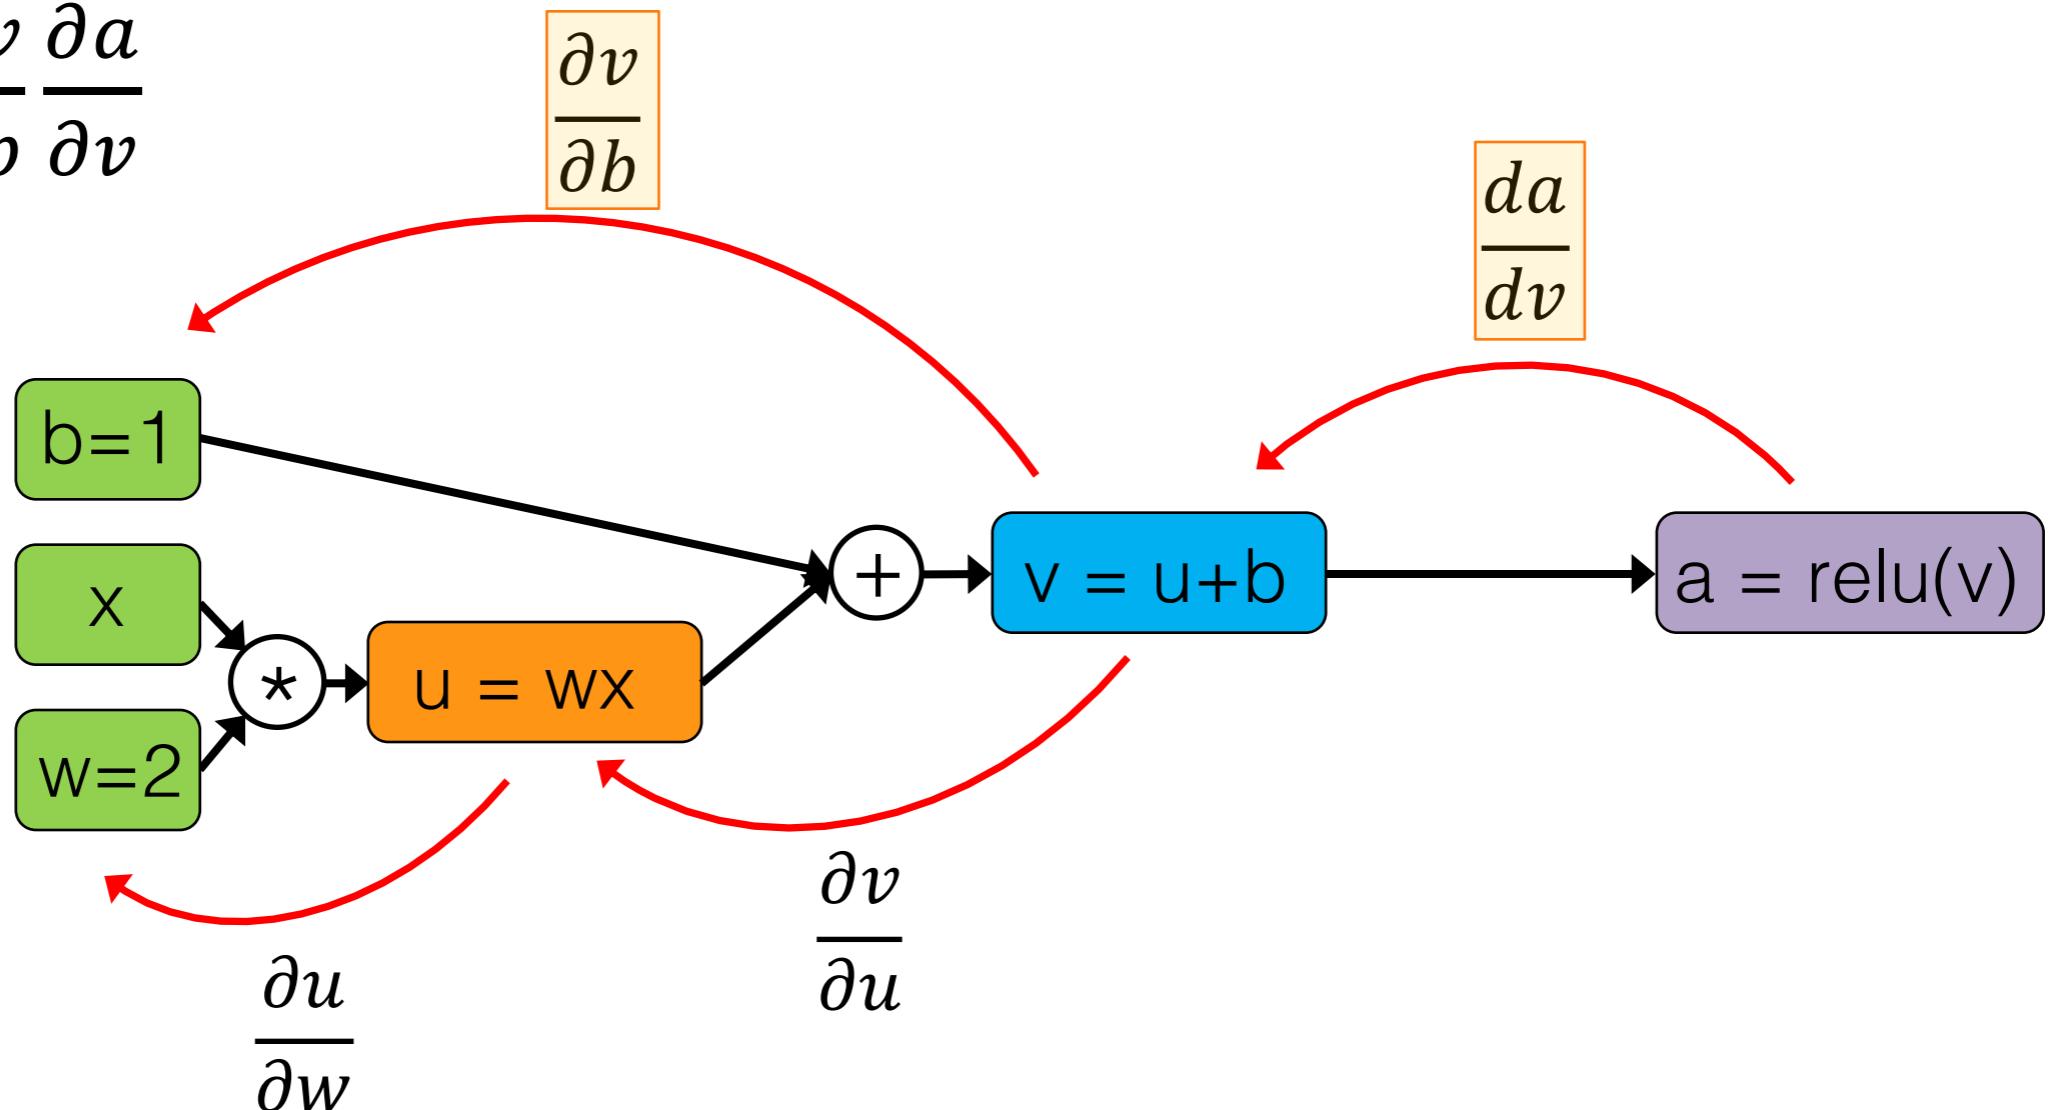


# Computation Graphs



# Computation Graphs

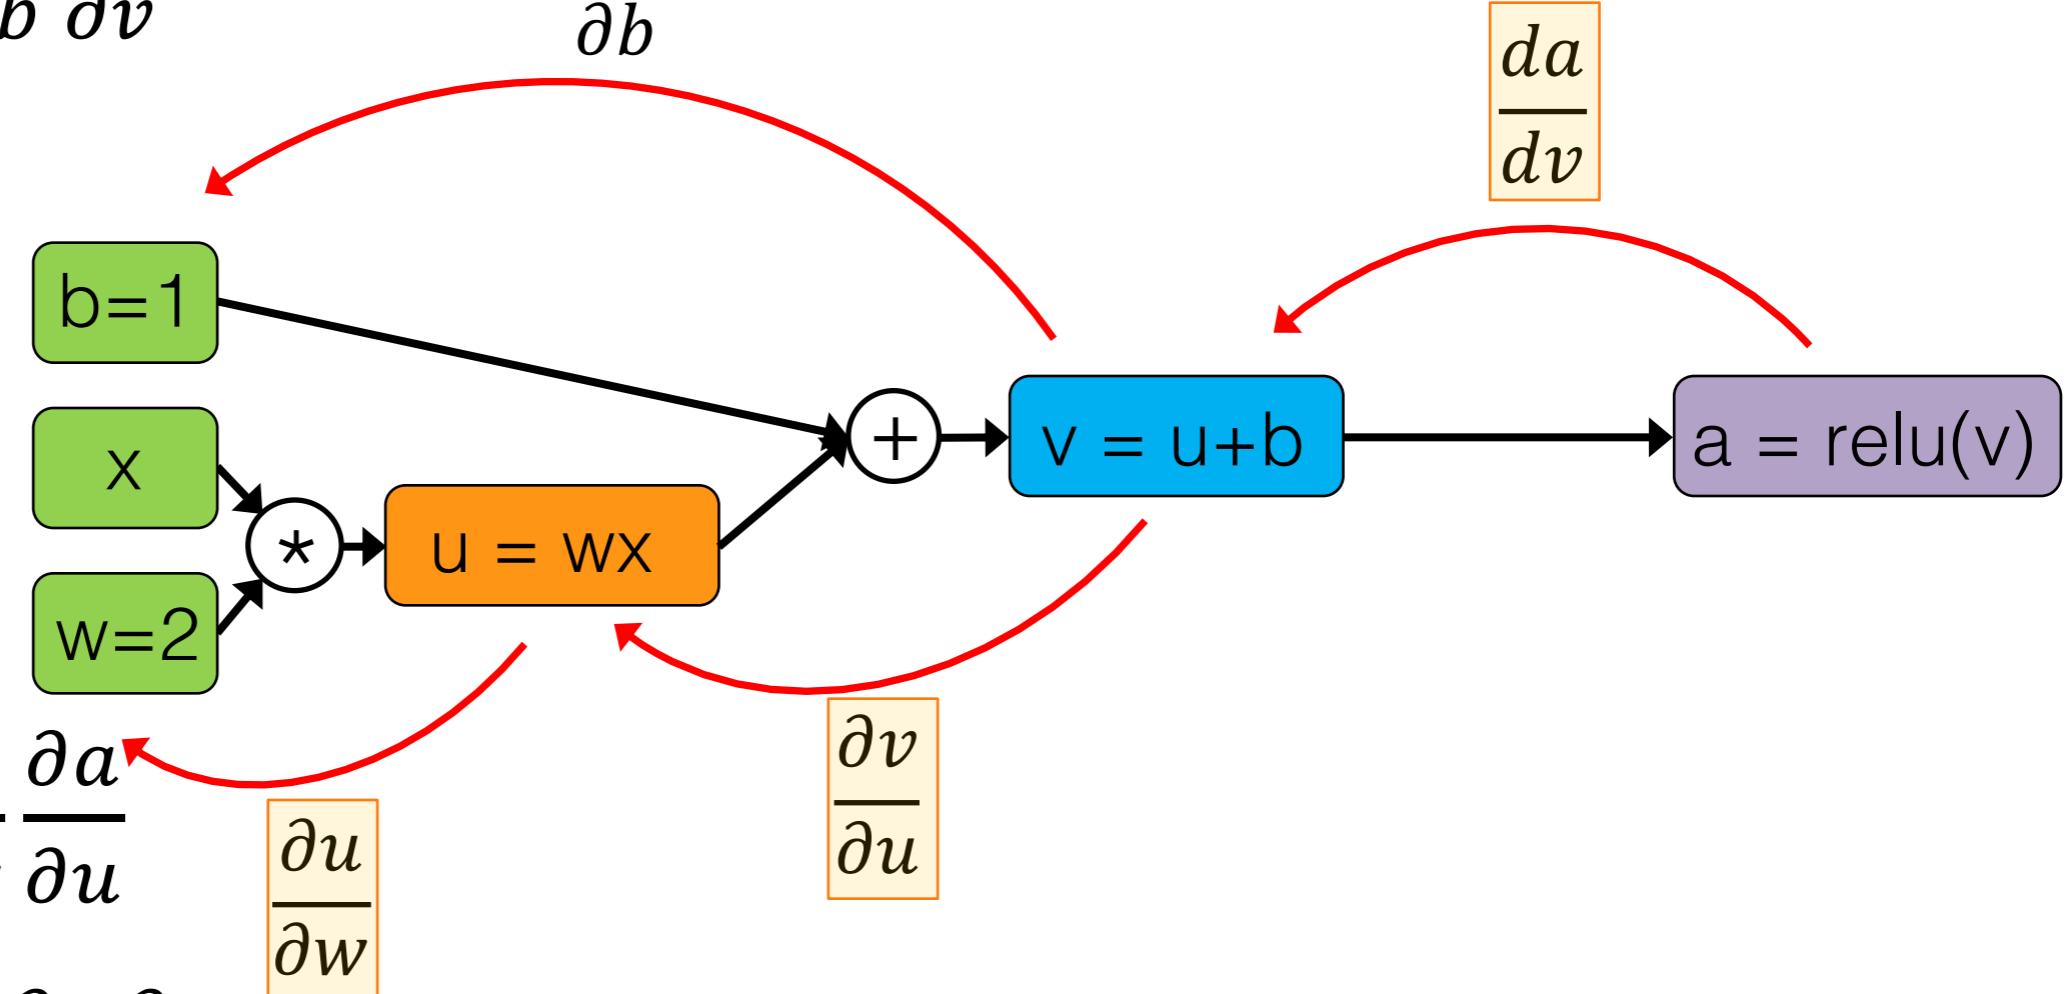
$$\frac{\partial a}{\partial b} = \frac{\partial v}{\partial b} \frac{\partial a}{\partial v}$$



# Computation Graphs

$$\frac{\partial a}{\partial b} = \frac{\partial v}{\partial b} \frac{\partial a}{\partial v}$$

$$\frac{\partial v}{\partial b}$$



$$\frac{\partial a}{\partial w} = \frac{\partial u}{\partial w} \frac{\partial a}{\partial u}$$

$$\frac{\partial u}{\partial w}$$

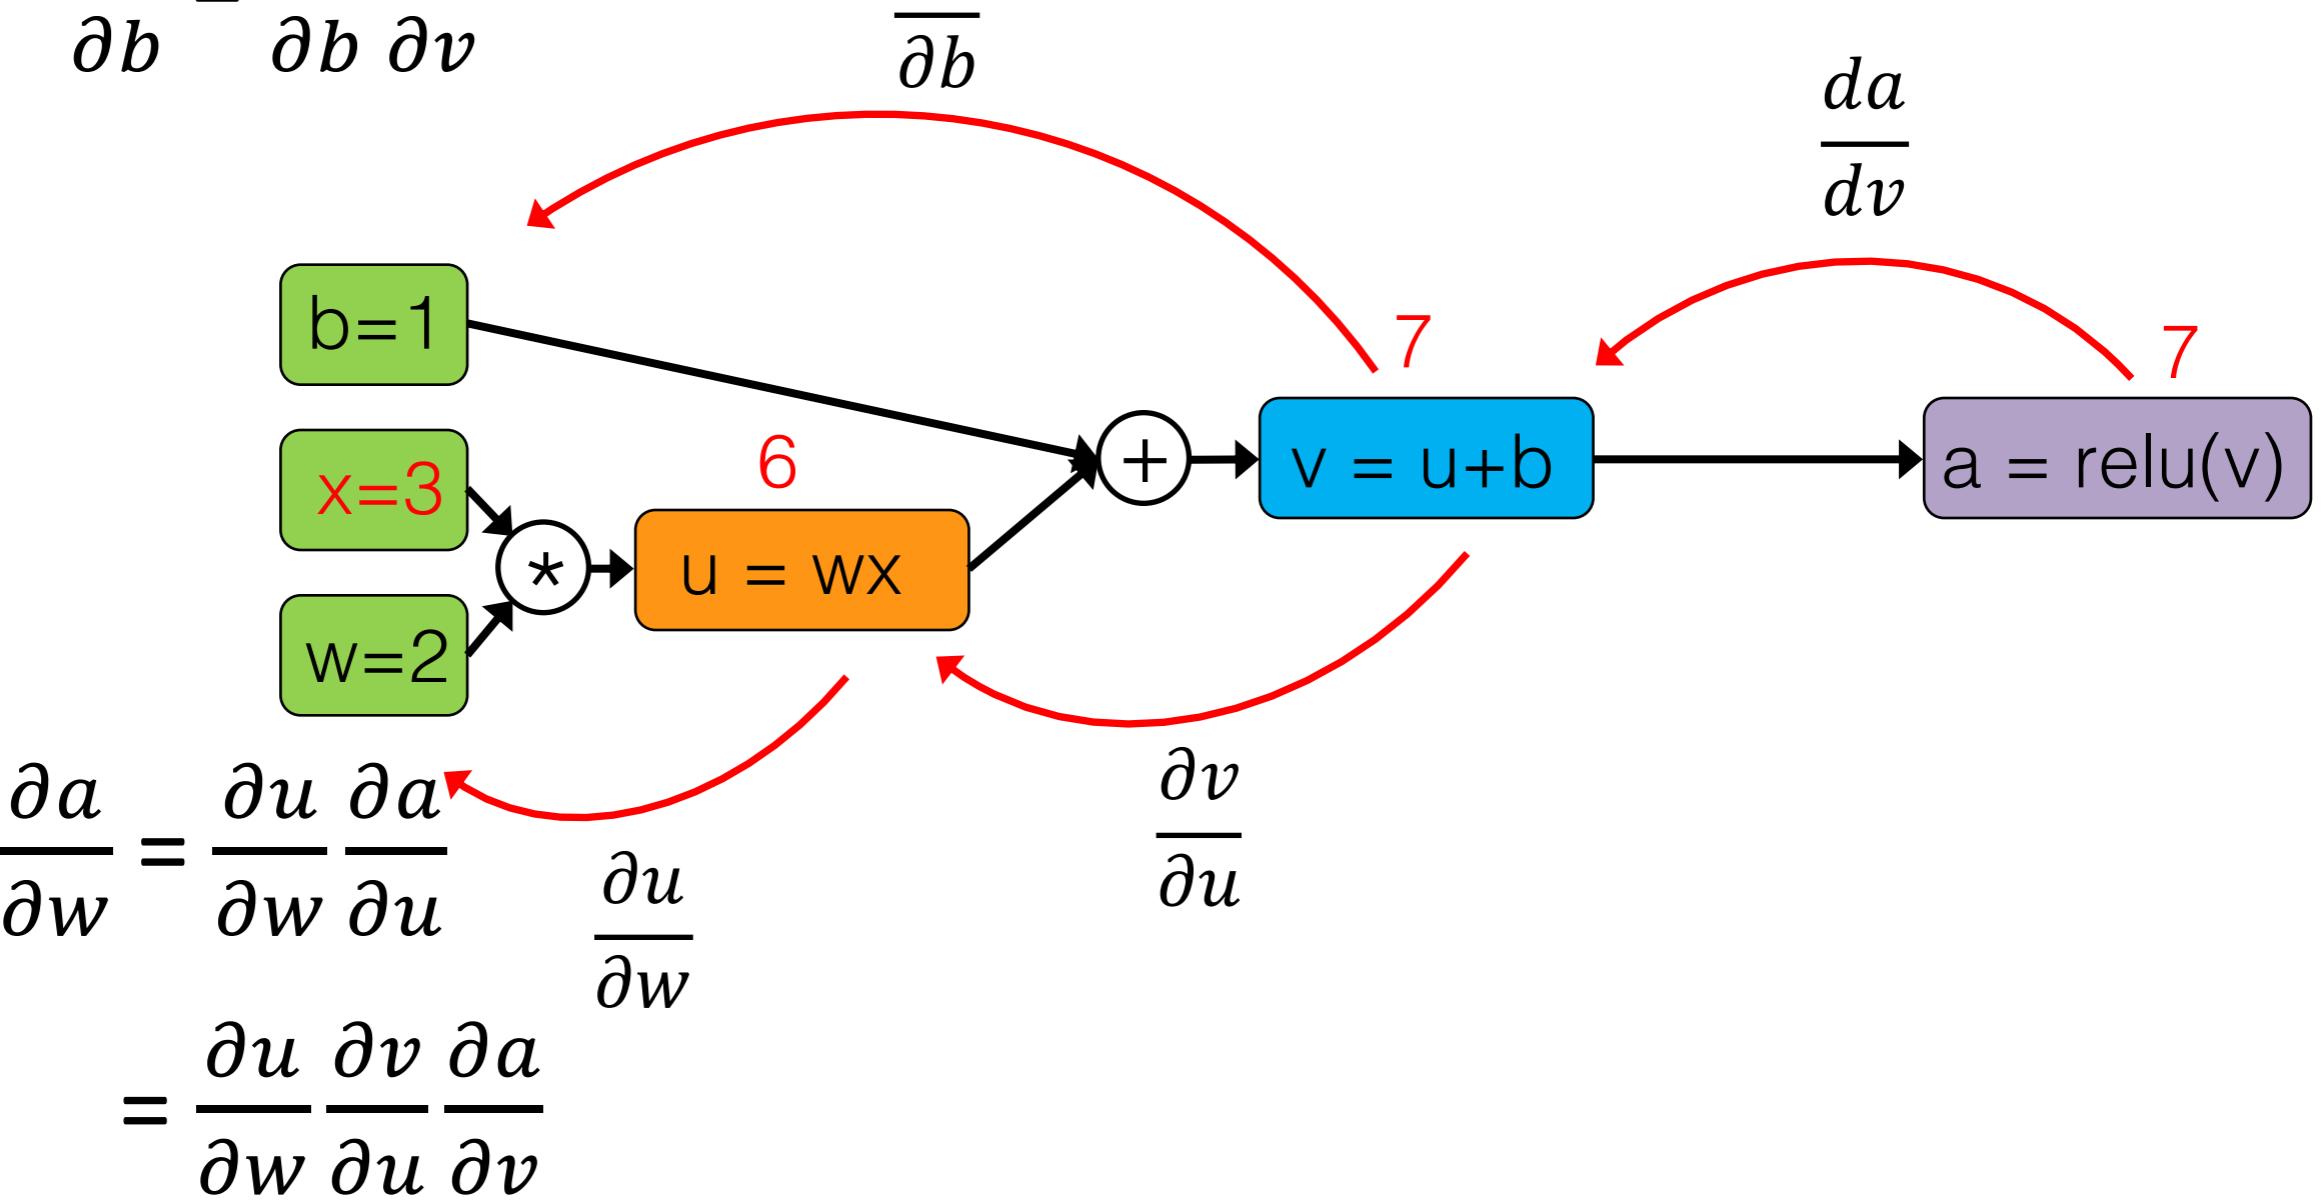
$$= \frac{\partial u}{\partial w} \frac{\partial v}{\partial u} \frac{\partial a}{\partial v}$$

# Computation Graphs

$$\frac{\partial a}{\partial b} = \frac{\partial v}{\partial b} \frac{\partial a}{\partial v}$$

$$\frac{\partial v}{\partial b}$$

$$\frac{da}{dv}$$



$$\frac{\partial a}{\partial w} = \frac{\partial u}{\partial w} \frac{\partial a}{\partial u}$$

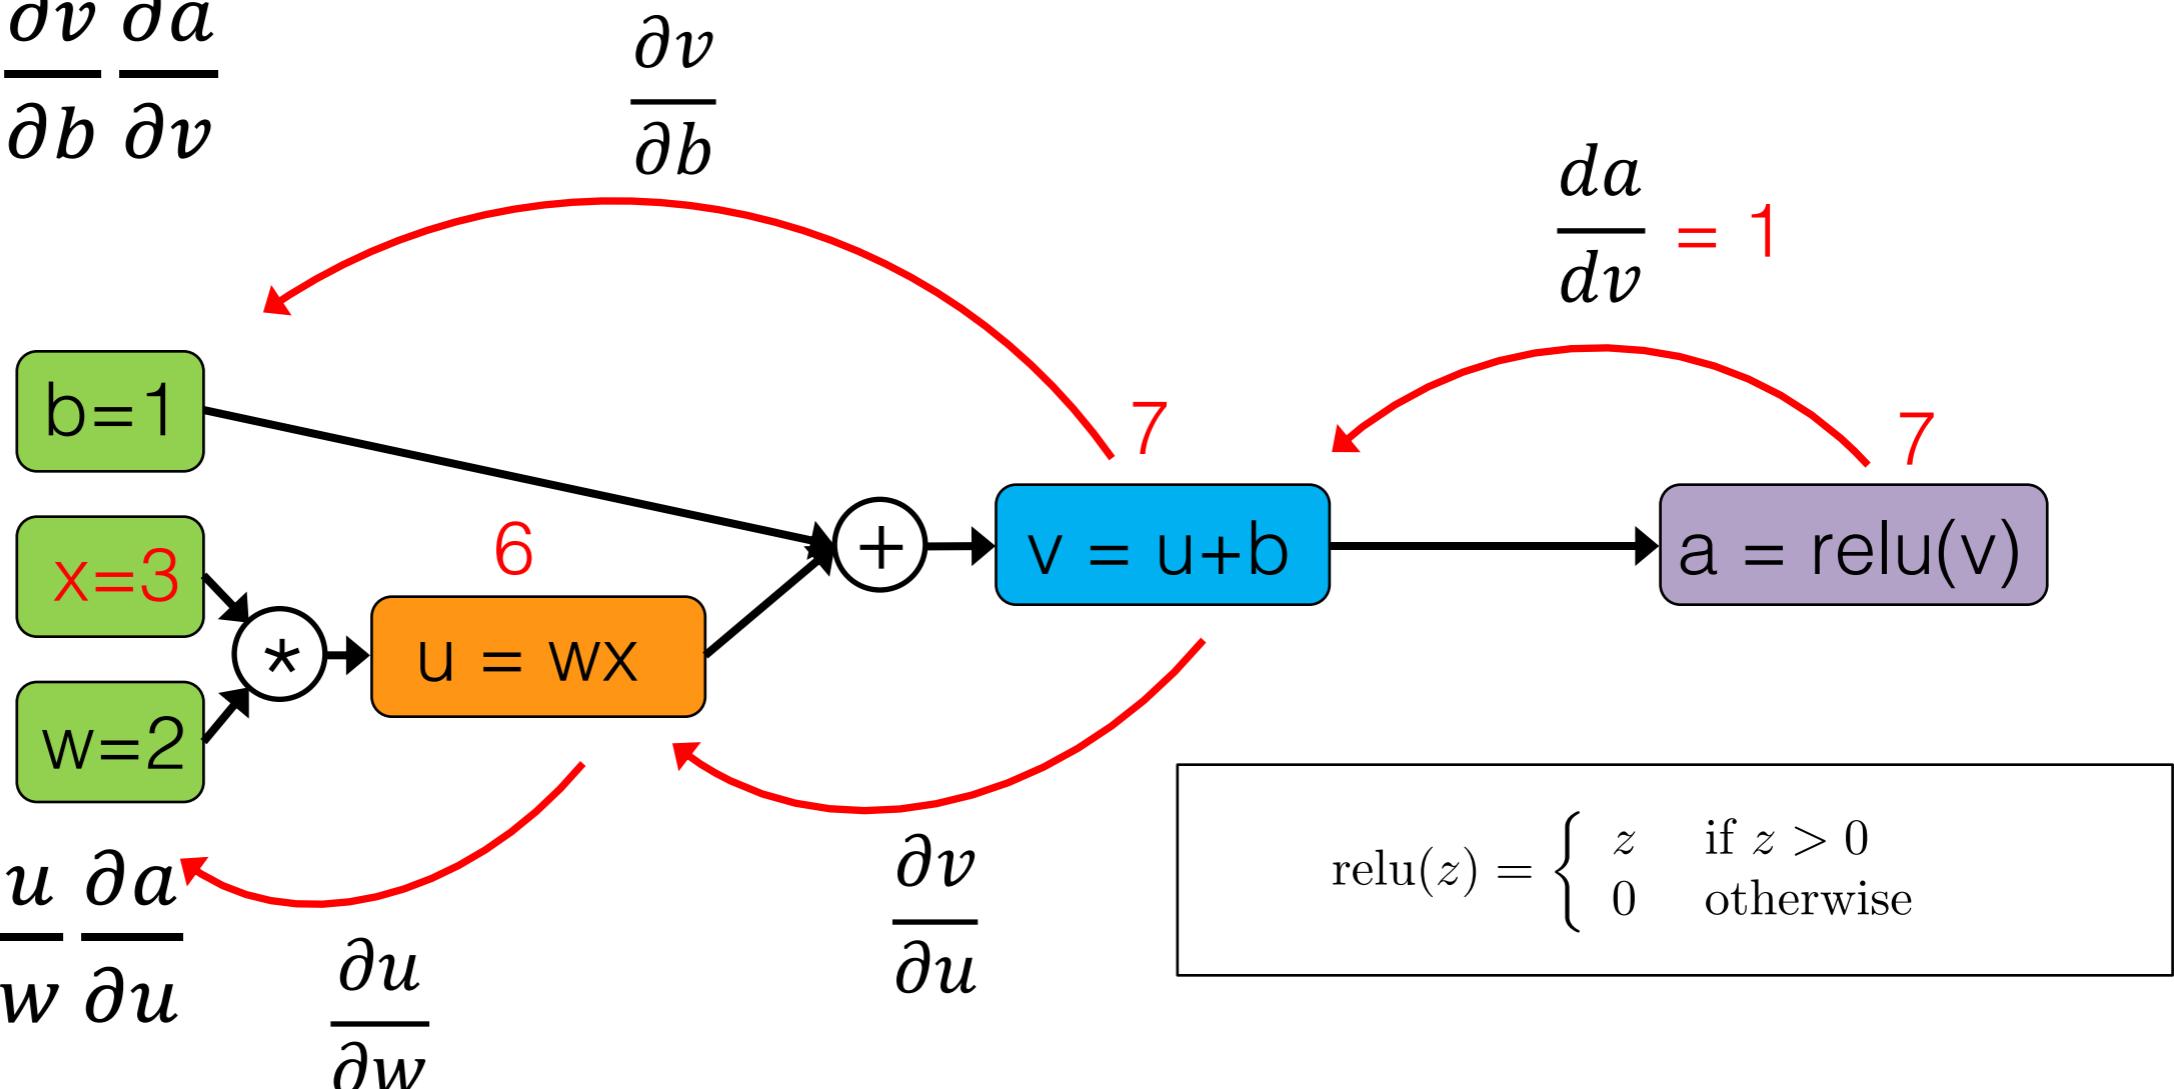
$$\frac{\partial u}{\partial w}$$

$$= \frac{\partial u}{\partial w} \frac{\partial v}{\partial u} \frac{\partial a}{\partial v}$$

# Computation Graphs

$$\frac{\partial a}{\partial b} = \frac{\partial v}{\partial b} \frac{\partial a}{\partial v}$$

$$\begin{aligned}\frac{\partial a}{\partial w} &= \frac{\partial u}{\partial w} \frac{\partial a}{\partial u} \\ &= \frac{\partial u}{\partial w} \frac{\partial v}{\partial u} \frac{\partial a}{\partial v}\end{aligned}$$

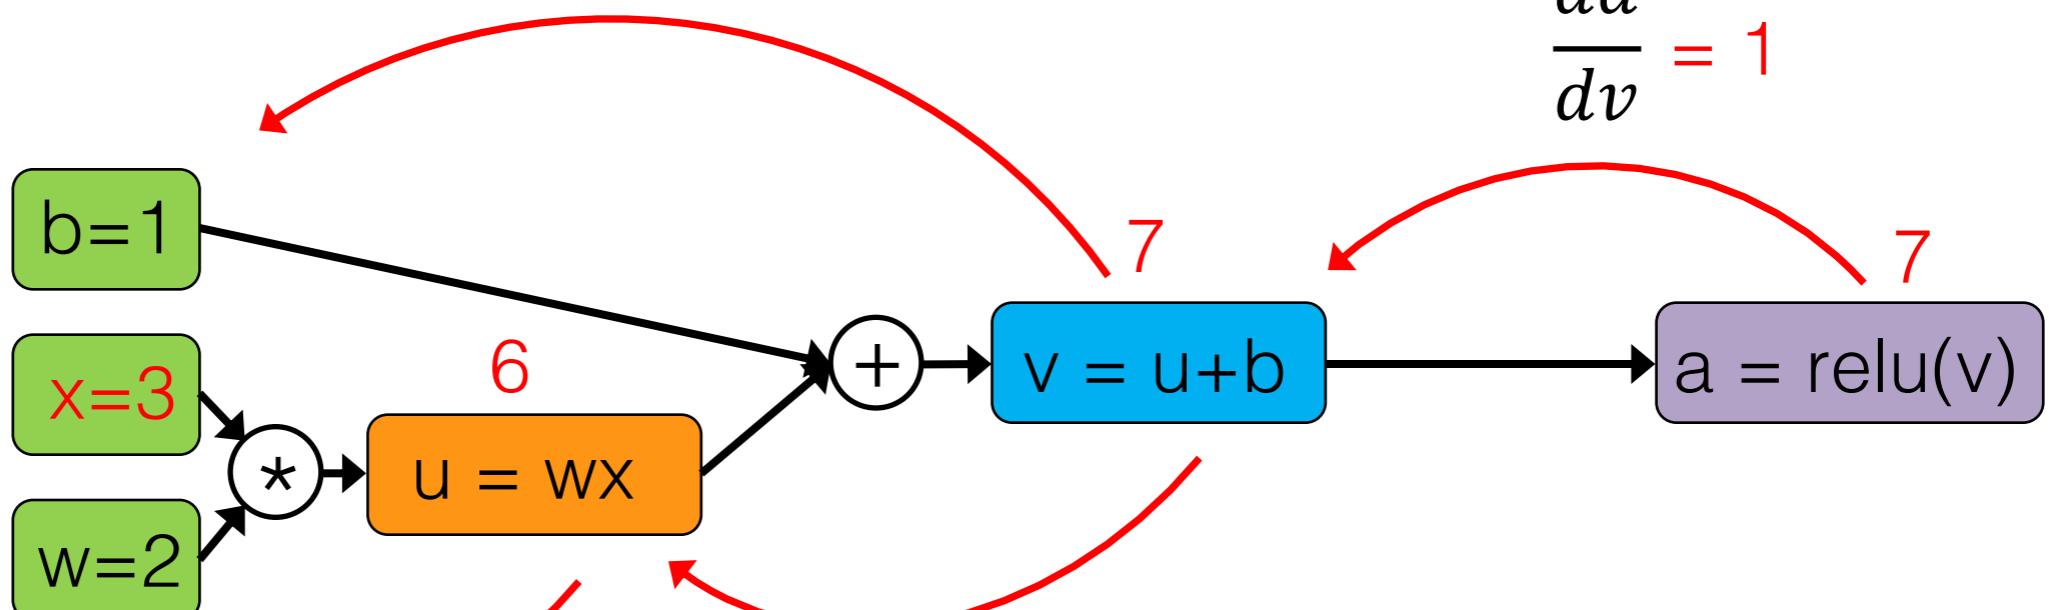


# Computation Graphs

$$\frac{\partial a}{\partial b} = \frac{\partial v}{\partial b} \frac{\partial a}{\partial v}$$

$$\frac{\partial v}{\partial b} = ?$$

$$\frac{da}{dv} = 1$$



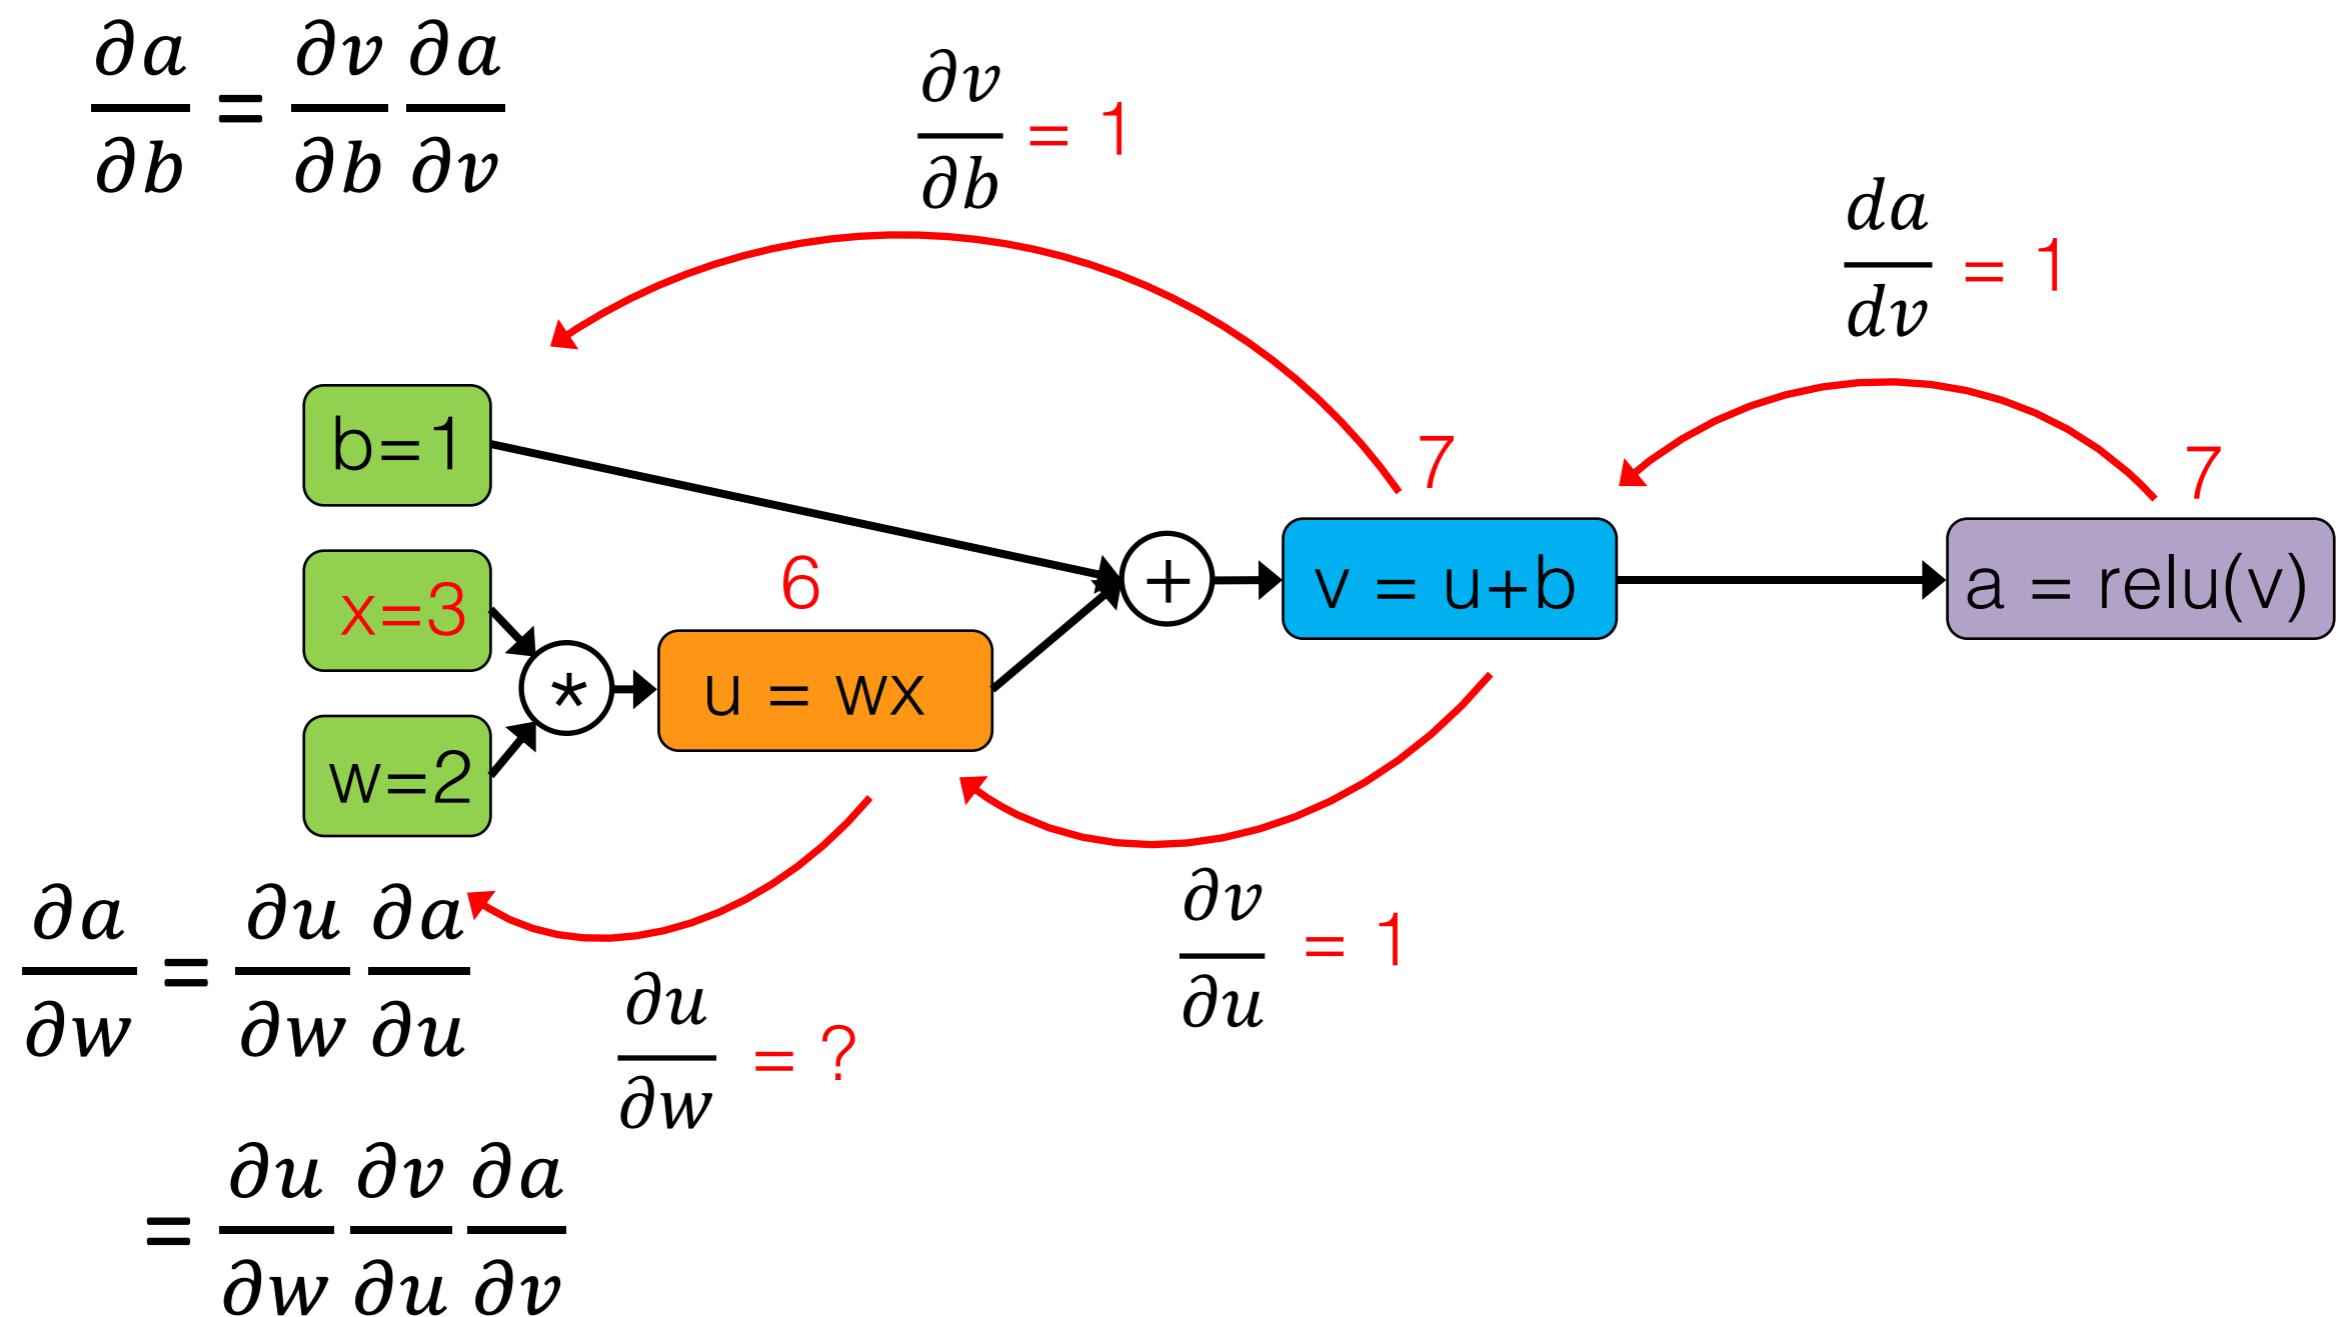
$$\frac{\partial a}{\partial w} = \frac{\partial u}{\partial w} \frac{\partial a}{\partial u}$$

$$\frac{\partial u}{\partial w}$$

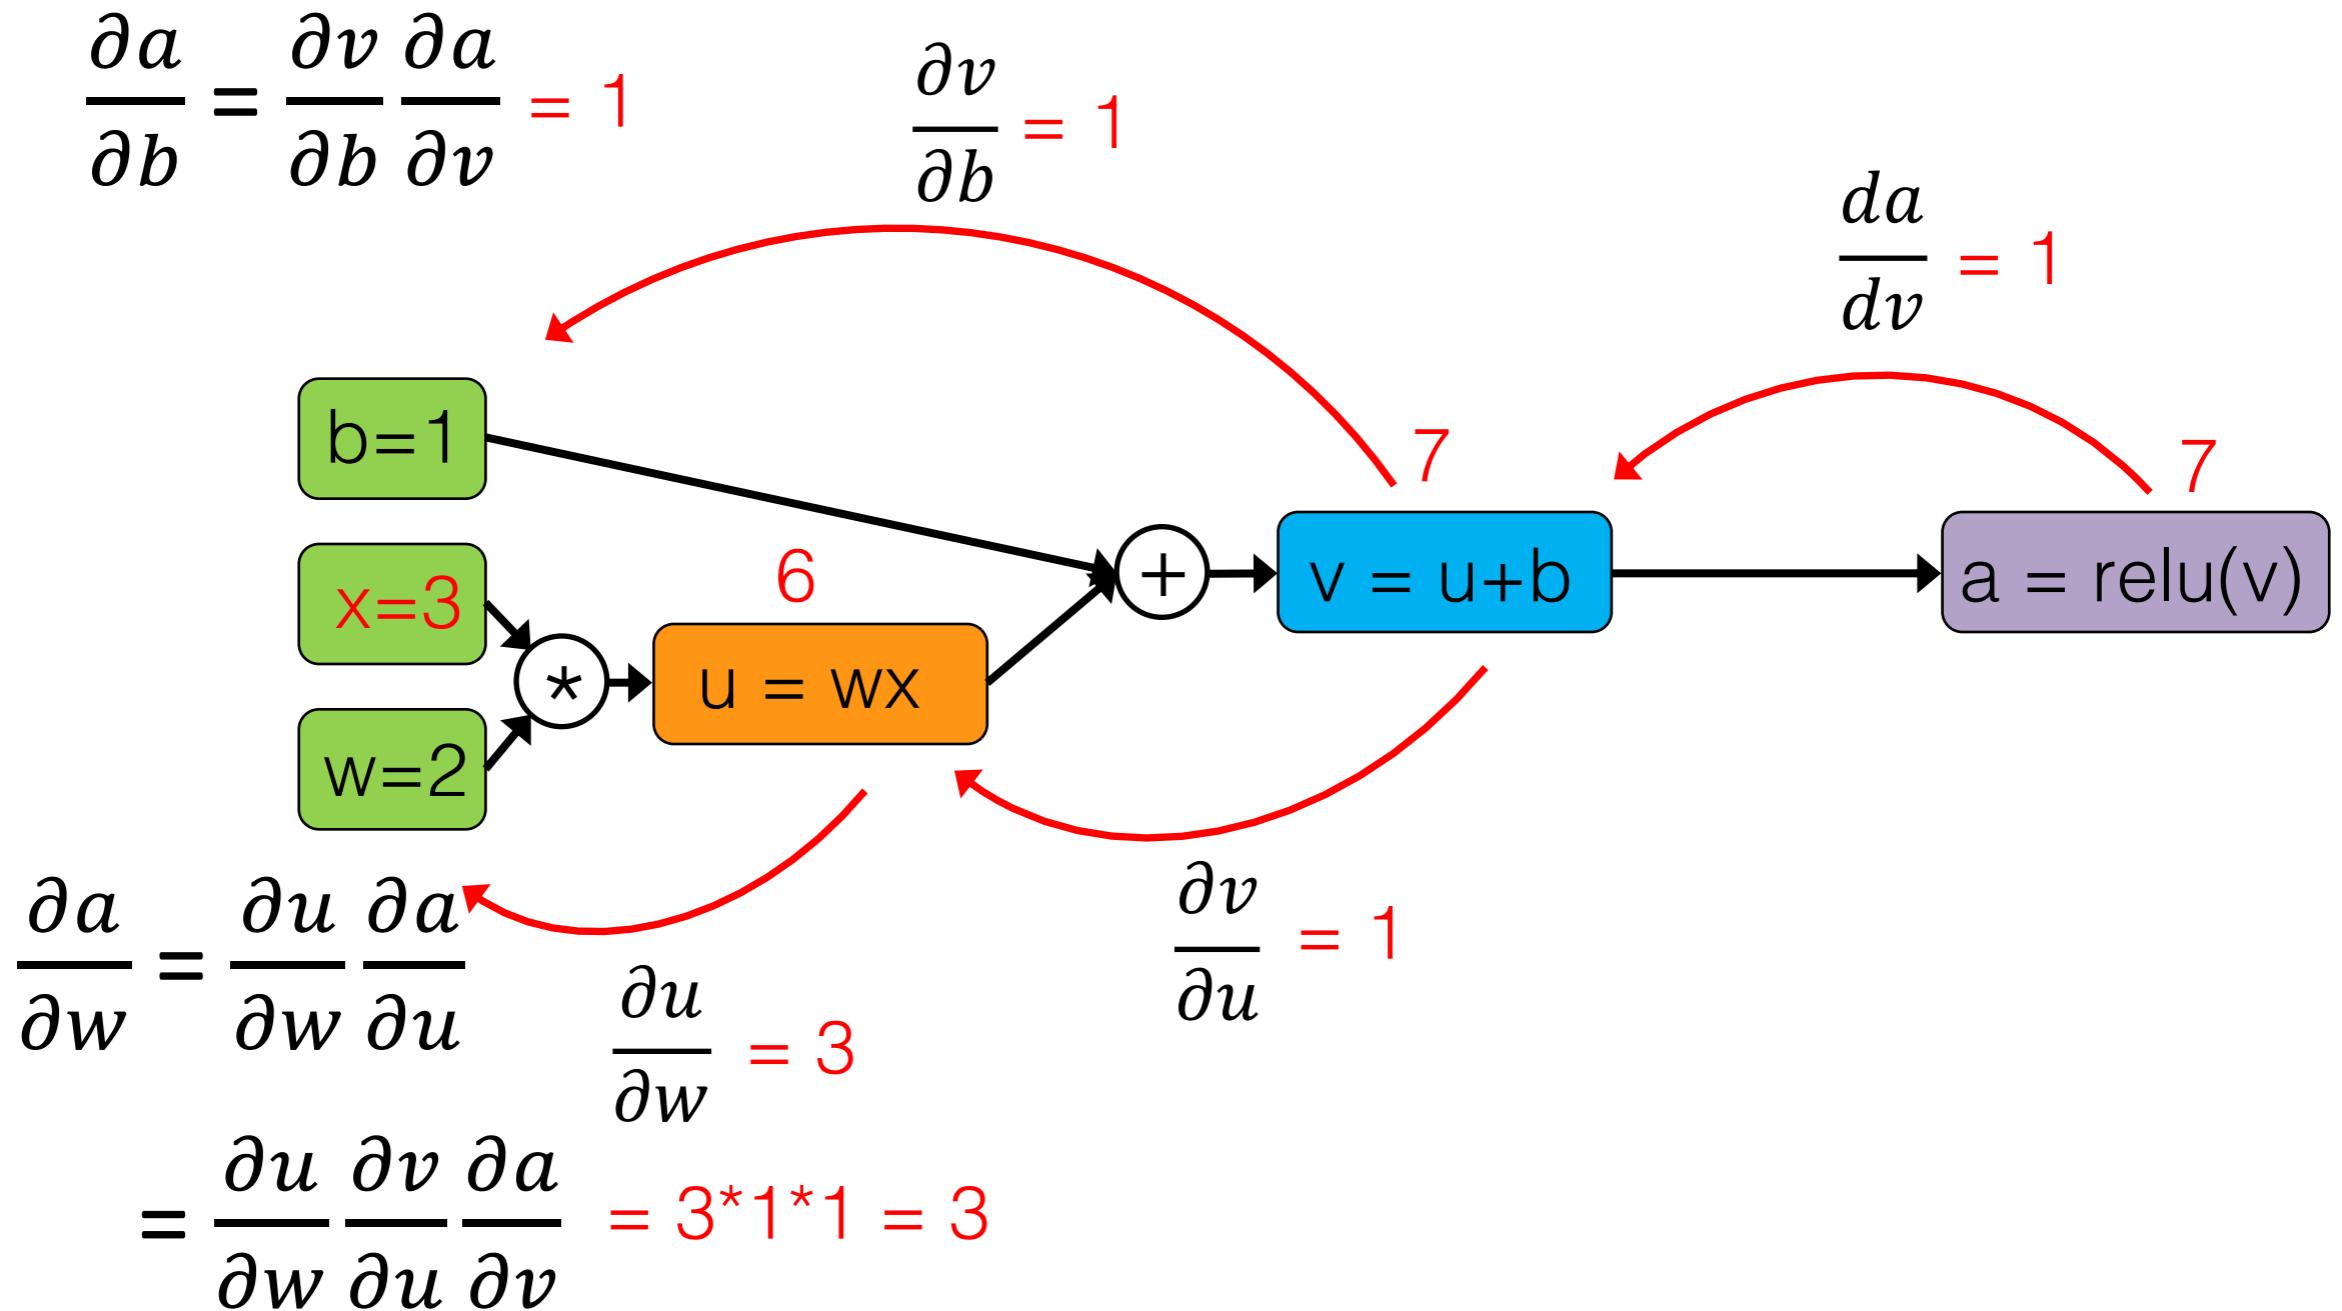
$$= \frac{\partial u}{\partial w} \frac{\partial v}{\partial u} \frac{\partial a}{\partial v}$$

Function	Derivative
$f(x) + g(x)$	$f'(x) + g'(x)$

# Computation Graphs

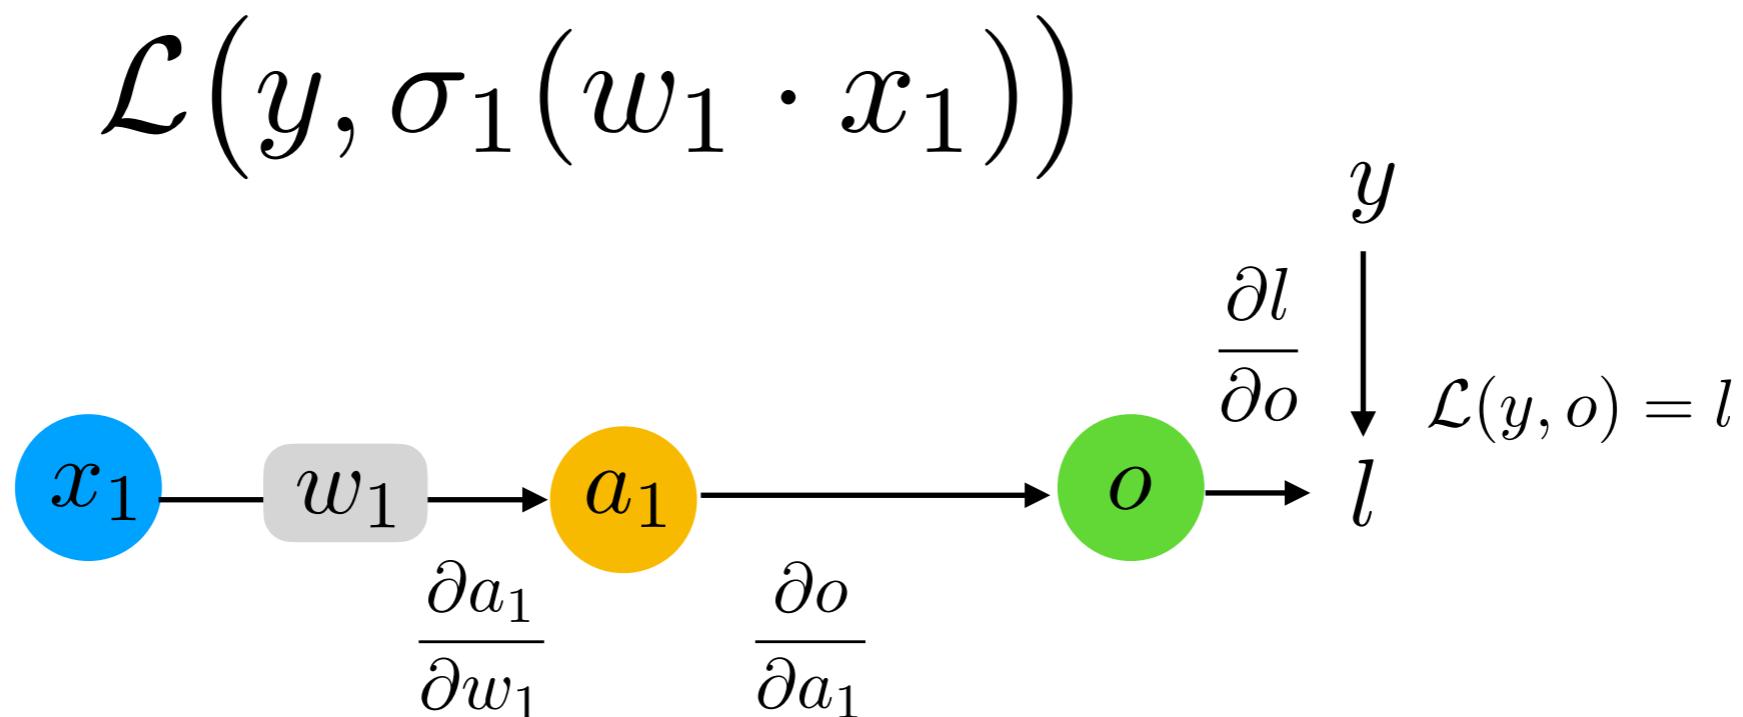


# Computation Graphs



# Some More Computation Graphs

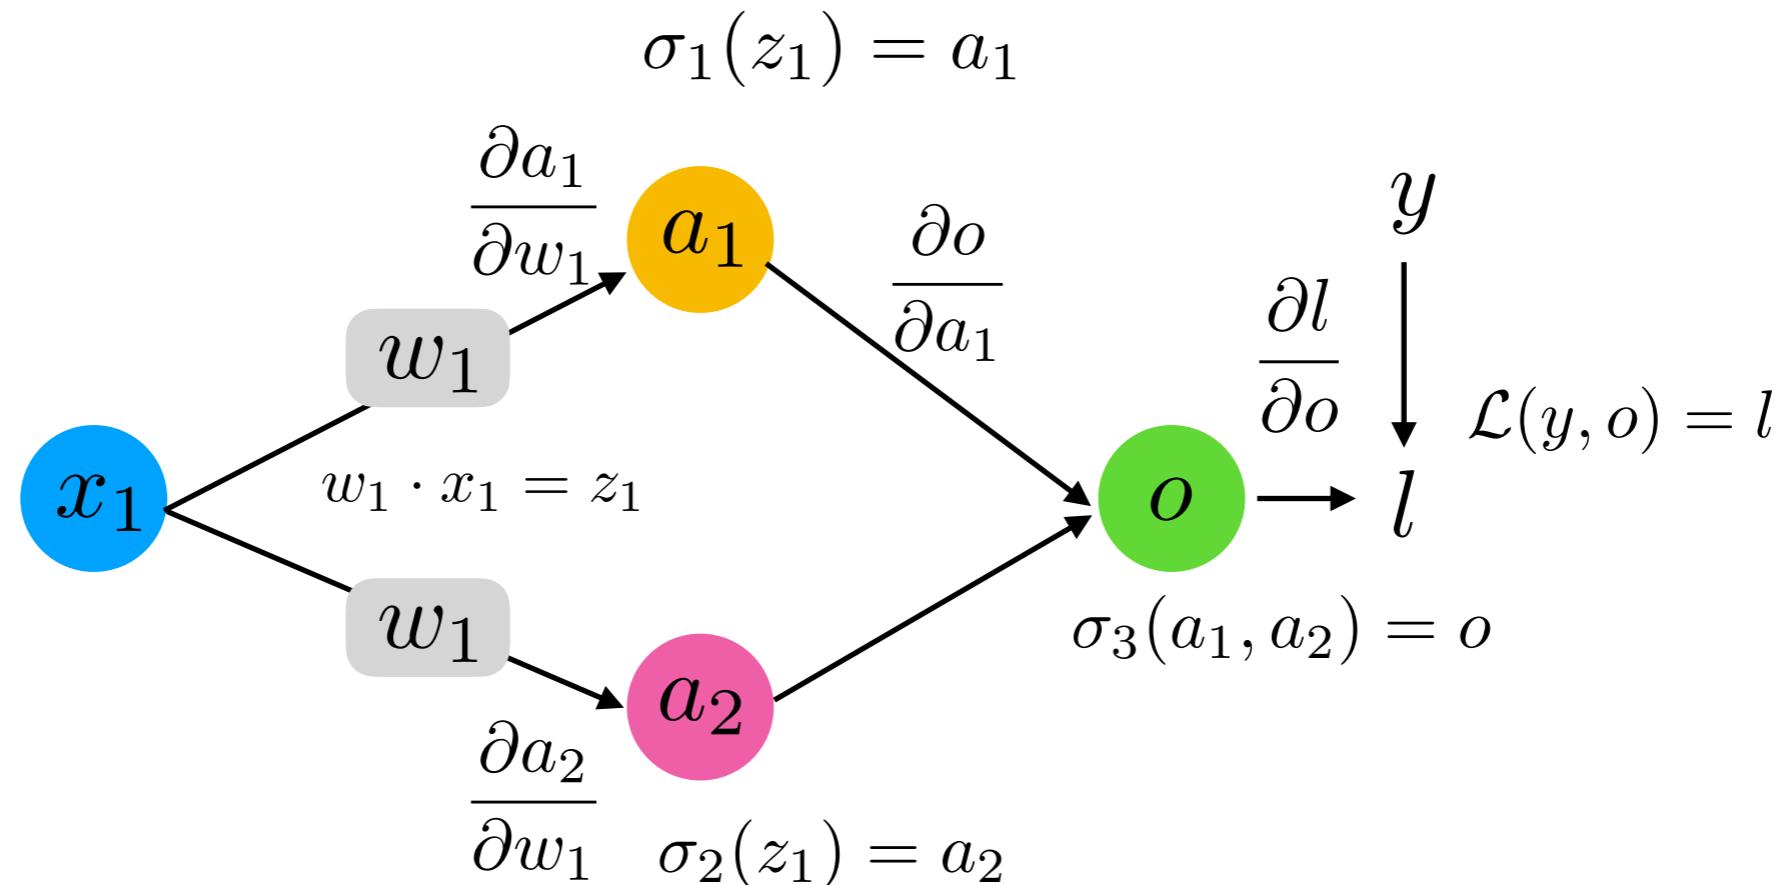
# Graph with Single Path



$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_1} \cdot \frac{\partial a_1}{\partial w_1} \quad (\text{univariate chain rule})$$

# Graph with Weight Sharing

$$\mathcal{L}(y, \sigma_3[\sigma_1(w_1 \cdot x_1), \sigma_2(w_1 \cdot x_1)])$$

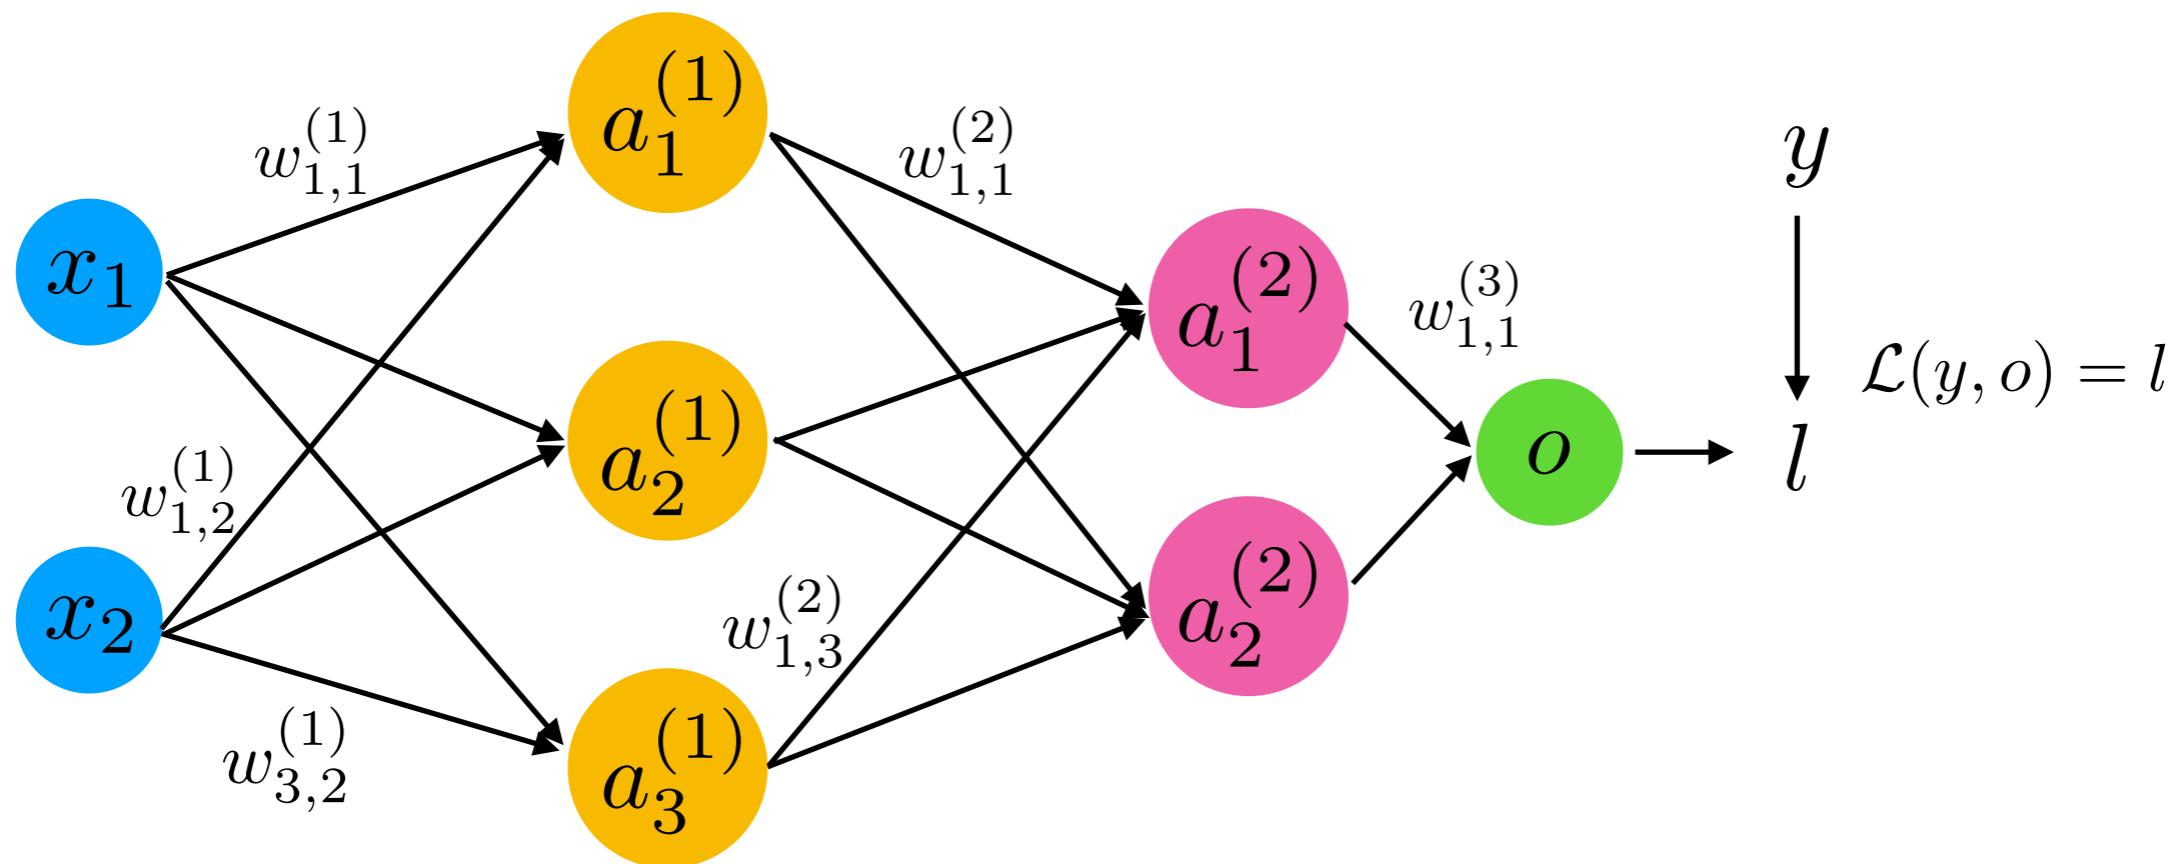


Upper path

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_1} \cdot \frac{\partial a_1}{\partial w_1} + \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_2} \cdot \frac{\partial a_2}{\partial w_1} \quad (\text{multivariable chain rule})$$

Lower path

# Graph with Fully-Connected Layers (later in this course)



$$\begin{aligned}\frac{\partial l}{\partial w_{1,1}^{(1)}} &= \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_1^{(2)}} \cdot \frac{\partial a_1^{(2)}}{\partial a_1^{(1)}} \cdot \frac{\partial a_1^{(1)}}{\partial w_{1,1}^{(1)}} \\ &\quad + \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_2^{(2)}} \cdot \frac{\partial a_2^{(2)}}{\partial a_1^{(1)}} \cdot \frac{\partial a_1^{(1)}}{\partial w_{1,1}^{(1)}}\end{aligned}$$

# Automatic Differentiation with PyTorch

## -- An Autograd Example

1. PyTorch Resources
2. Computation Graphs
- 3. Automatic Differentiation in PyTorch**
4. Training ADALINE Manually Vs Automatically in PyTorch
5. A Closer Look at the PyTorch API

# PyTorch Autograd Example

<https://github.com/rasbt/stat453-deep-learning-ss21/tree/master/L06/code/pytorch-autograd.ipynb>

# Training an Adaptive Linear Neuron in PyTorch

1. PyTorch Resources
2. Computation Graphs
3. Automatic Differentiation in PyTorch
- 4. Training ADALINE Manually Vs Automatically in PyTorch**
5. A Closer Look at the PyTorch API

simplify to `super().__init__()` in constructor

# PyTorch ADALINE (neuron model) Example

<https://github.com/rasbt/stat453-deep-learning-ss21/tree/master/L06/code/adaline-with-autograd.ipynb>

# Using PyTorch: A Closer Look at the Object-Oriented and Functional APIs

1. PyTorch Resources
2. Computation Graphs
3. Automatic Differentiation in PyTorch
4. Training ADALINE Manually Vs Automatically in PyTorch
- 5. A Closer Look at the PyTorch API**

# PyTorch Usage: Step 1 (Definition)

```
class MultilayerPerceptron(torch.nn.Module): ←  
  
    def __init__(self, num_features, num_classes):  
        super(MultilayerPerceptron, self).__init__()  
  
        ### 1st hidden layer  
        self.linear_1 = torch.nn.Linear(num_feat, num_h1)  
  
        ### 2nd hidden layer  
        self.linear_2 = torch.nn.Linear(num_h1, num_h2)  
  
        ### Output layer  
        self.linear_out = torch.nn.Linear(num_h2, num_classes)  
  
    def forward(self, x):  
        out = self.linear_1(x)  
        out = F.relu(out)  
        out = self.linear_2(out)  
        out = F.relu(out)  
        logits = self.linear_out(out)  
        probas = F.log_softmax(logits, dim=1)  
        return logits, probas
```

Backward will be inferred automatically if we use the nn.Module class!

Define model parameters that will be instantiated when created an object of this class

Define how and it what order the model parameters should be used in the forward pass

# PyTorch Usage: Step 2 (Creation)

```
torch.manual_seed(random_seed)
model = MultilayerPerceptron(num_features=num_features,
                             num_classes=num_classes)    | Instantiate model
                                                               (creates the model parameters)

model = model.to(device)

optimizer = torch.optim.SGD(model.parameters(),
                           lr=learning_rate)        | Define an optimization method
```

# PyTorch Usage: Step 2 (Creation)

```
torch.manual_seed(random_seed)
model = MultilayerPerceptron(num_features=num_features,
                             num_classes=num_classes)

model = model.to(device) ←
optimizer = torch.optim.SGD(model.parameters(),
                           lr=learning_rate)
```

Optionally move model to GPU, where  
device e.g. `torch.device('cuda:0')`

# PyTorch Usage: Step 3 (Training)

```
for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)

        ### FORWARD AND BACK PROP
        logits, probas = model(features)
        cost = F.cross_entropy(probas, targets)
        optimizer.zero_grad()

        cost.backward()

        ### UPDATE MODEL PARAMETERS
        optimizer.step()

model.eval()
with torch.no_grad():
    # compute accuracy
```

Run for a specified number of epochs

Iterate over minibatches in epoch

If your model is on the GPU, data should also be on the GPU

y = model(x) calls `__call__` and then `.forward()`, where some extra stuff is done in `__call__`;  
don't run y = model.forward(x) directly

Gradients at each leaf node are accumulated under the `.grad` attribute, not just stored. This is why we have to zero them before each backward pass

# PyTorch Usage: Step 3 (Training)

```
for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)

        ### FORWARD AND BACK PROP
        logits, probas = model(features) ← This will run the forward() method
        loss = F.cross_entropy(logits, targets) ← Define a loss function to optimize
        optimizer.zero_grad() ← Set the gradient to zero
                               (could be non-zero from a previous forward pass)

        loss.backward() ← Compute the gradients, the backward is
                         automatically constructed by "autograd" based on
                         the forward() method and the loss function

        ### UPDATE MODEL PARAMETERS
        optimizer.step() ← Use the gradients to update the weights according to
                           the optimization method (defined on the previous
                           slide)
                           E.g., for SGD,  $w := w + \text{learning\_rate} \times \text{gradient}$ 

    model.eval()
    with torch.no_grad():
        # compute accuracy
```

# PyTorch Usage: Step 3 (Training)

```
for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)

        ### FORWARD AND BACK PROP
        logits, probas = model(features)
        loss = F.cross_entropy(logits, targets)
        optimizer.zero_grad()

        loss.backward()

        ### UPDATE MODEL PARAMETERS
        optimizer.step()

    model.eval()
    with torch.no_grad():
        # compute accuracy
```

For evaluation, set the model to eval mode (will be relevant later when we use DropOut or BatchNorm)

This prevents the computation graph for backpropagation from automatically being build in the background to save memory

# PyTorch Usage: Step 3 (Training)

```
for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)

        ### FORWARD AND BACK PROP
        logits, probas = model(features)
        loss = F.cross_entropy(logits, targets)
        optimizer.zero_grad()

        loss.backward()

        ### UPDATE MODEL PARAMETERS
        optimizer.step()

    model.eval()
    with torch.no_grad():
        # compute accuracy
```

logits because of computational efficiency.  
Basically, it internally uses a `log_softmax(logits)` function  
that is more stable than `log(softmax(logits))`.  
More on logits ("net inputs" of the last layer) in the  
next lecture. Please also see

# Objected-Oriented vs Functional\* API

\*Note that with "functional" I mean "functional programming" (one paradigm in CS)

`torch.nn.functional` = api without internal state

```
import torch.nn.functional as F

class MultilayerPerceptron(torch.nn.Module):
    def __init__(self, num_features, num_classes):
        super(MultilayerPerceptron, self).__init__()
        # 1st hidden layer
        self.linear_1 = torch.nn.Linear(num_features,
                                        num_hidden_1)
        # 2nd hidden layer
        self.linear_2 = torch.nn.Linear(num_hidden_1,
                                        num_hidden_2)
        # Output layer
        self.linear_out = torch.nn.Linear(num_hidden_2,
                                         num_classes)

    def forward(self, x):
        out = self.linear_1(x)
        out = F.relu(out)
        out = self.linear_2(out)
        out = F.relu(out)
        logits = self.linear_out(out)
        probas = F.log_softmax(logits, dim=1)
        return logits, probas
```

Unnecessary because these functions  
don't need to store a state but maybe  
helpful for keeping track of order of ops  
(when implementing "forward")

```
class MultilayerPerceptron(torch.nn.Module):
    def __init__(self, num_features, num_classes):
        super(MultilayerPerceptron, self).__init__()
        # 1st hidden layer
        self.linear_1 = torch.nn.Linear(num_features,
                                        num_hidden_1)
        self.relu1 = torch.nn.ReLU()
        # 2nd hidden layer
        self.linear_2 = torch.nn.Linear(num_hidden_1,
                                        num_hidden_2)
        self.relu2 = torch.nn.ReLU()
        # Output layer
        self.linear_out = torch.nn.Linear(num_hidden_2,
                                         num_classes)
        self.softmax = torch.nn.Softmax()

    def forward(self, x):
        out = self.linear_1(x)
        out = self.relu1(out)
        out = self.linear_2(out)
        out = self.relu2(out)
        logits = self.linear_out(out)
        probas = self.softmax(logits, dim=1)
        return logits, probas
```

# Objected-Oriented vs Functional API

Using "Sequential"

```
import torch.nn.functional as F

class MultilayerPerceptron(torch.nn.Module):

    def __init__(self, num_features, num_classes):
        super(MultilayerPerceptron, self).__init__()

        ### 1st hidden layer
        self.linear_1 = torch.nn.Linear(num_features,
                                       num_hidden_1)

        ### 2nd hidden layer
        self.linear_2 = torch.nn.Linear(num_hidden_1,
                                       num_hidden_2)

        ### Output layer
        self.linear_out = torch.nn.Linear(num_hidden_2,
                                         num_classes)

    def forward(self, x):
        out = self.linear_1(x)
        out = F.relu(out)
        out = self.linear_2(out)
        out = F.relu(out)
        logits = self.linear_out(out)
        probas = F.log_softmax(logits, dim=1)
        return logits, probas
```

```
class MultilayerPerceptron(torch.nn.Module):

    def __init__(self, num_features, num_classes):
        super(MultilayerPerceptron, self).__init__()

        self.my_network = torch.nn.Sequential(
            torch.nn.Linear(num_features, num_hidden_1),
            torch.nn.ReLU(),
            torch.nn.Linear(num_hidden_1, num_hidden_2),
            torch.nn.ReLU(),
            torch.nn.Linear(num_hidden_2, num_classes)
        )

    def forward(self, x):
        logits = self.my_network(x)
        probas = F.softmax(logits, dim=1)
        return logits, probas
```

Much more compact and clear, but  
"forward" may be harder to debug if there  
are errors (we cannot simply add  
breakpoints or insert "print" statements

# Objected-Oriented vs Functional API

## Using "Sequential"

1)

```
class MultilayerPerceptron(torch.nn.Module):

    def __init__(self, num_features, num_classes):
        super(MultilayerPerceptron, self).__init__()

        self.my_network = torch.nn.Sequential(
            torch.nn.Linear(num_features, num_hidden),
            torch.nn.ReLU(),
            torch.nn.Linear(num_hidden_1, num_hidden_2),
            torch.nn.ReLU(),
            torch.nn.Linear(num_hidden_2, num_classes)
        )

    def forward(self, x):
        logits = self.my_network(x)
        probas = F.softmax(logits, dim=1)
        return logits, probas
```

Much more compact and clear, but  
"forward" may be harder to debug if there  
are errors (we cannot simply add  
breakpoints or insert "print" statements

2)

However, if you use Sequential, you can  
define "hooks" to get intermediate outputs.  
For example:

```
[7]: model.net
[7]: Sequential(
[7]:     (0): Linear(in_features=784, out_features=128, bias=True)
[7]:     (1): ReLU(inplace)
[7]:     (2): Linear(in_features=128, out_features=256, bias=True)
[7]:     (3): ReLU(inplace)
[7]:     (4): Linear(in_features=256, out_features=10, bias=True)
[7]: )
[7]: If we want to get the output from the 2nd layer during the forward pass, we can register a hook as follows:
[8]: outputs = []
[8]: def hook(module, input, output):
[8]:     outputs.append(output)
[8]: model.net[2].register_forward_hook(hook)
[8]: <torch.utils.hooks.RemovableHandle at 0x7f659c6685c0>
[8]: Now, if we call the model on some inputs, it will save the intermediate results in the "outputs" list:
[9]: _ = model(features)
[9]: print(outputs)
[9]: [tensor([[0.5341, 1.0513, 2.3542, ..., 0.0000, 0.0000, 0.0000],
[9]:         [0.0000, 0.6676, 0.6620, ..., 0.0000, 0.0000, 2.4056],
[9]:         [1.1520, 0.0000, 0.0000, ..., 2.5860, 0.8992, 0.9642],
[9]:         ...,
[9]:         [0.0000, 0.1076, 0.0000, ..., 1.8367, 0.0000, 2.5203],
[9]:         [0.5415, 0.0000, 0.0000, ..., 2.7968, 0.8244, 1.6335],
[9]:         [1.0710, 0.9805, 3.0103, ..., 0.0000, 0.0000, 0.0000]], device='cuda:3', grad_fn=<ThresholdBackward1>)]
```

## Jupyter Notebook vs Python Scripts

In general, we recommend to use jupyter notebooks for initial exploration/ playing around with new models and code. Python scripts should be used as soon as you want to train the model on a bigger dataset where also reproducibility is more important.

**Our recommended workflow:**

1. Start with a jupyter notebook
2. Explore the data and models
3. Build your classes/ methods inside cells of the notebook
4. Move your code to python scripts
5. Train/ deploy on server

Jupyter Notebook	Python Scripts
+ Exploration	+ Running longer jobs without interruption
+ Debugging	+ Easy to track changes with git
- Can become a huge file	- Debugging mostly means rerunning the whole script
- Can be interrupted (don't use for long training)	
- Prone to errors and become a mess	

Type	Convention	Example
Packages & Modules	lower_with_underscores	<code>from prefetch_generator import BackgroundGenerator</code>
Classes	CapWords	<code>class DataLoader</code>
Constants	CAPS_WITH_UNDER	<code>BATCH_SIZE=16</code>
Instances	lower_with_underscores	<code>dataset = Dataset</code>
Methods & Functions	lower_with_underscores()	<code>def visualize_tensor()</code>
Variables	lower_with_underscores	<code>background_color='Blue'</code>

More PyTorch features will be introduced step-by-step later in this course when we start working with more complex networks, including

- Running code on the GPU
- Using efficient data loaders
- Splitting networks across different GPUs