

## Lecture 06

# Decision Trees

STAT 451: Intro to Machine Learning, Fall 2021

Sebastian Raschka

# Lecture 6: Decision Trees

## Topics

### **1. Intro to decision trees**

2. Recursive algorithms & Big-O

3. Types of decision trees

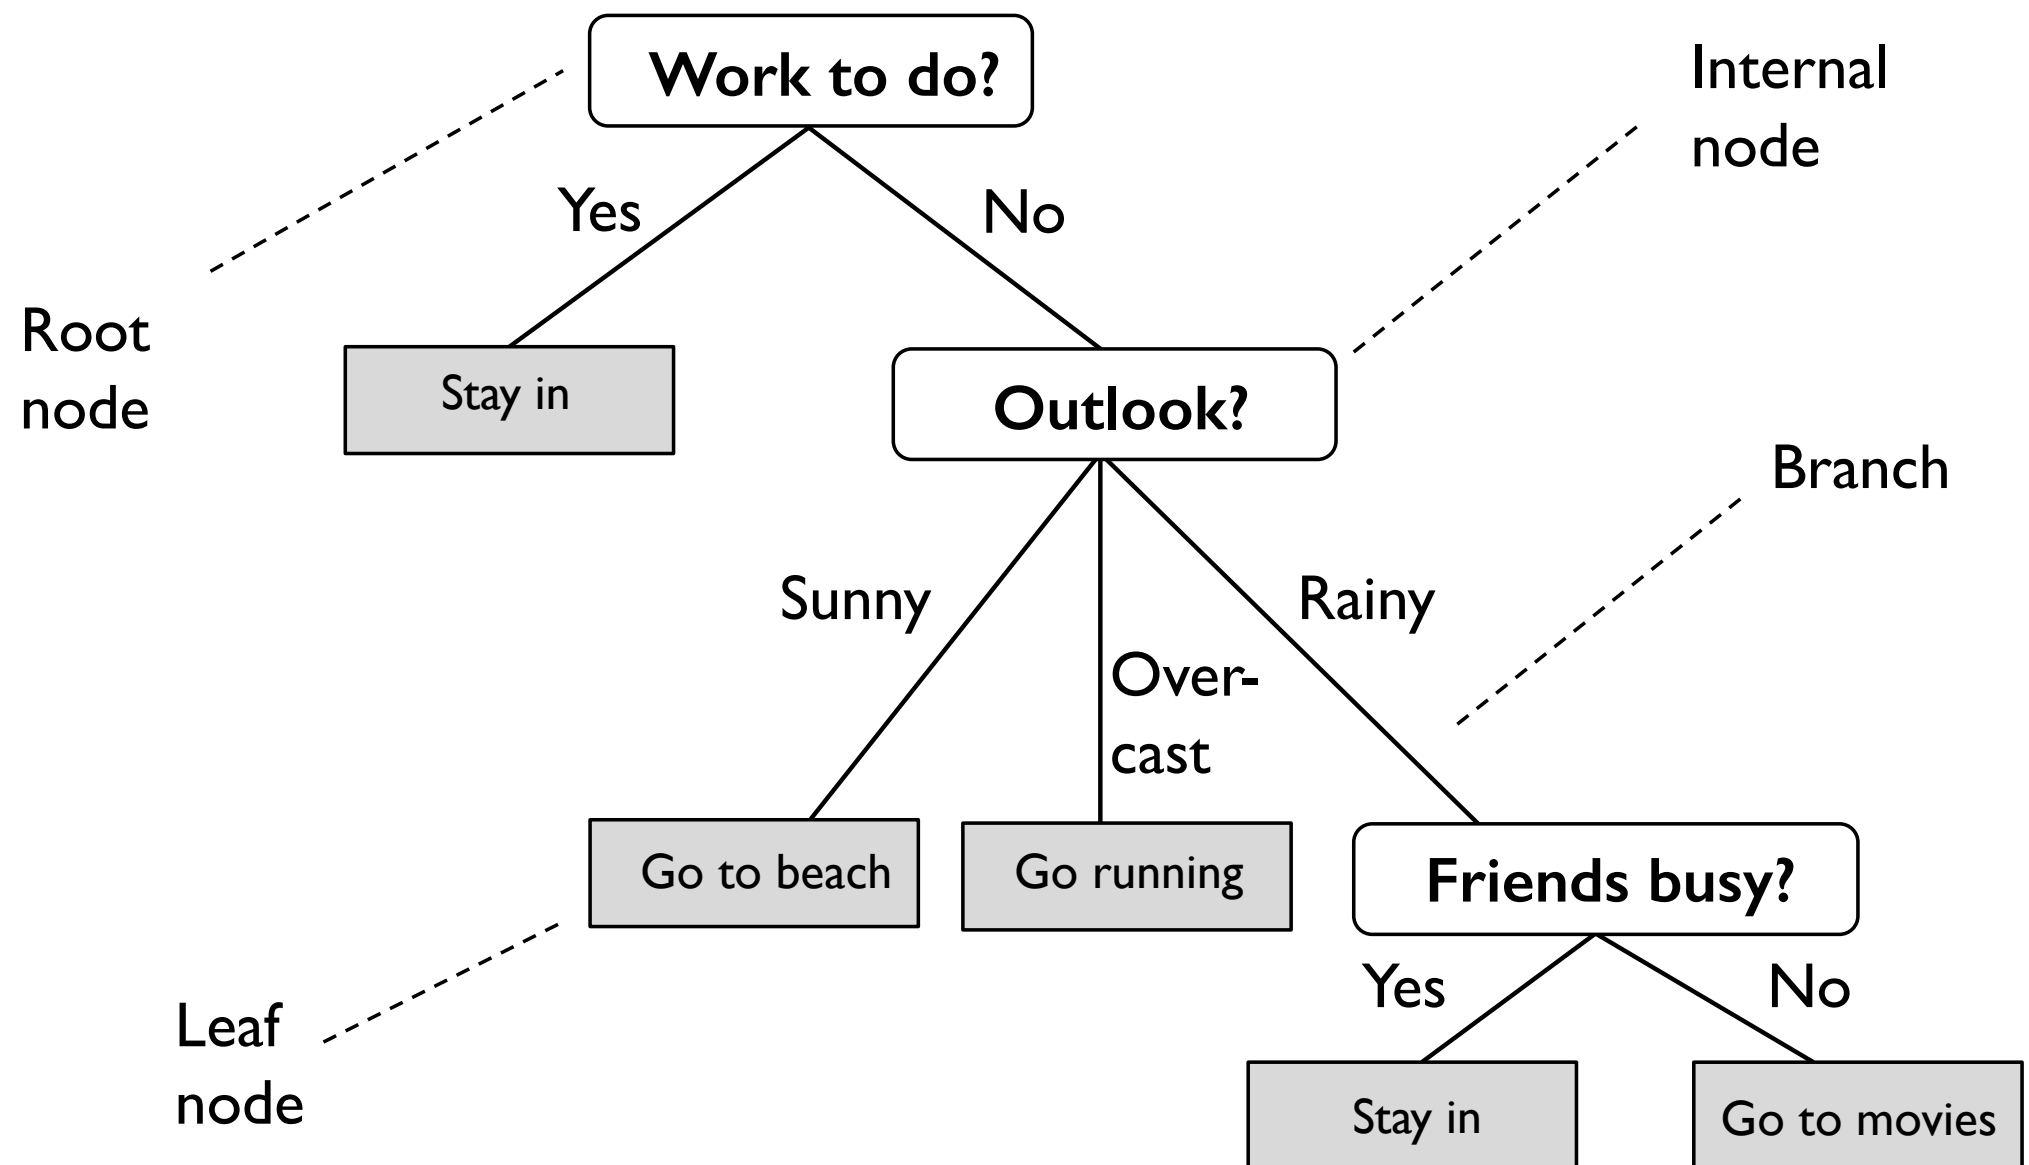
4. Splitting criteria

5. Gini & Entropy vs misclassification error

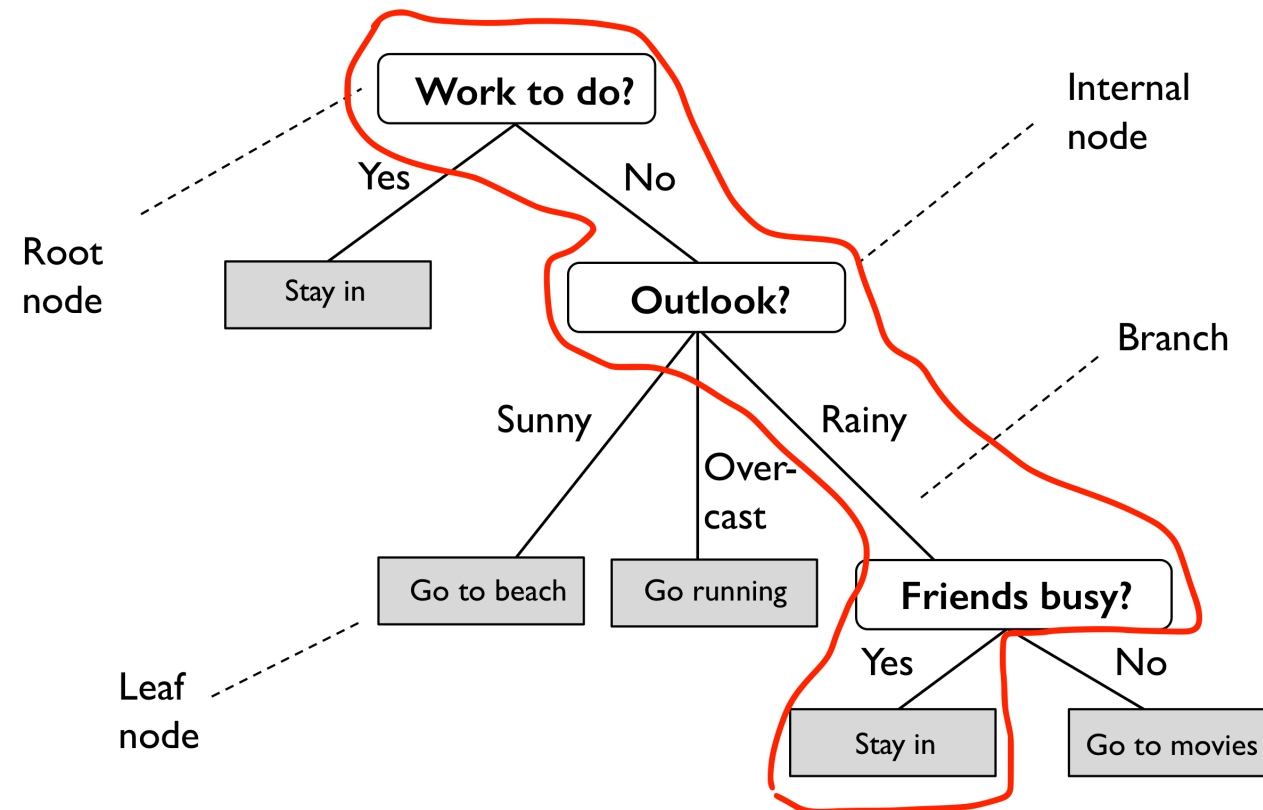
6. Improvements & dealing with overfitting

7. Code example

# Decision Tree Terminology





# Decision Trees as Rulesets




**IF**


**THEN**




"ENE4:" 90.5 %



ZINC72400307: 90.4 %



ZINC03876071: 3.2 %

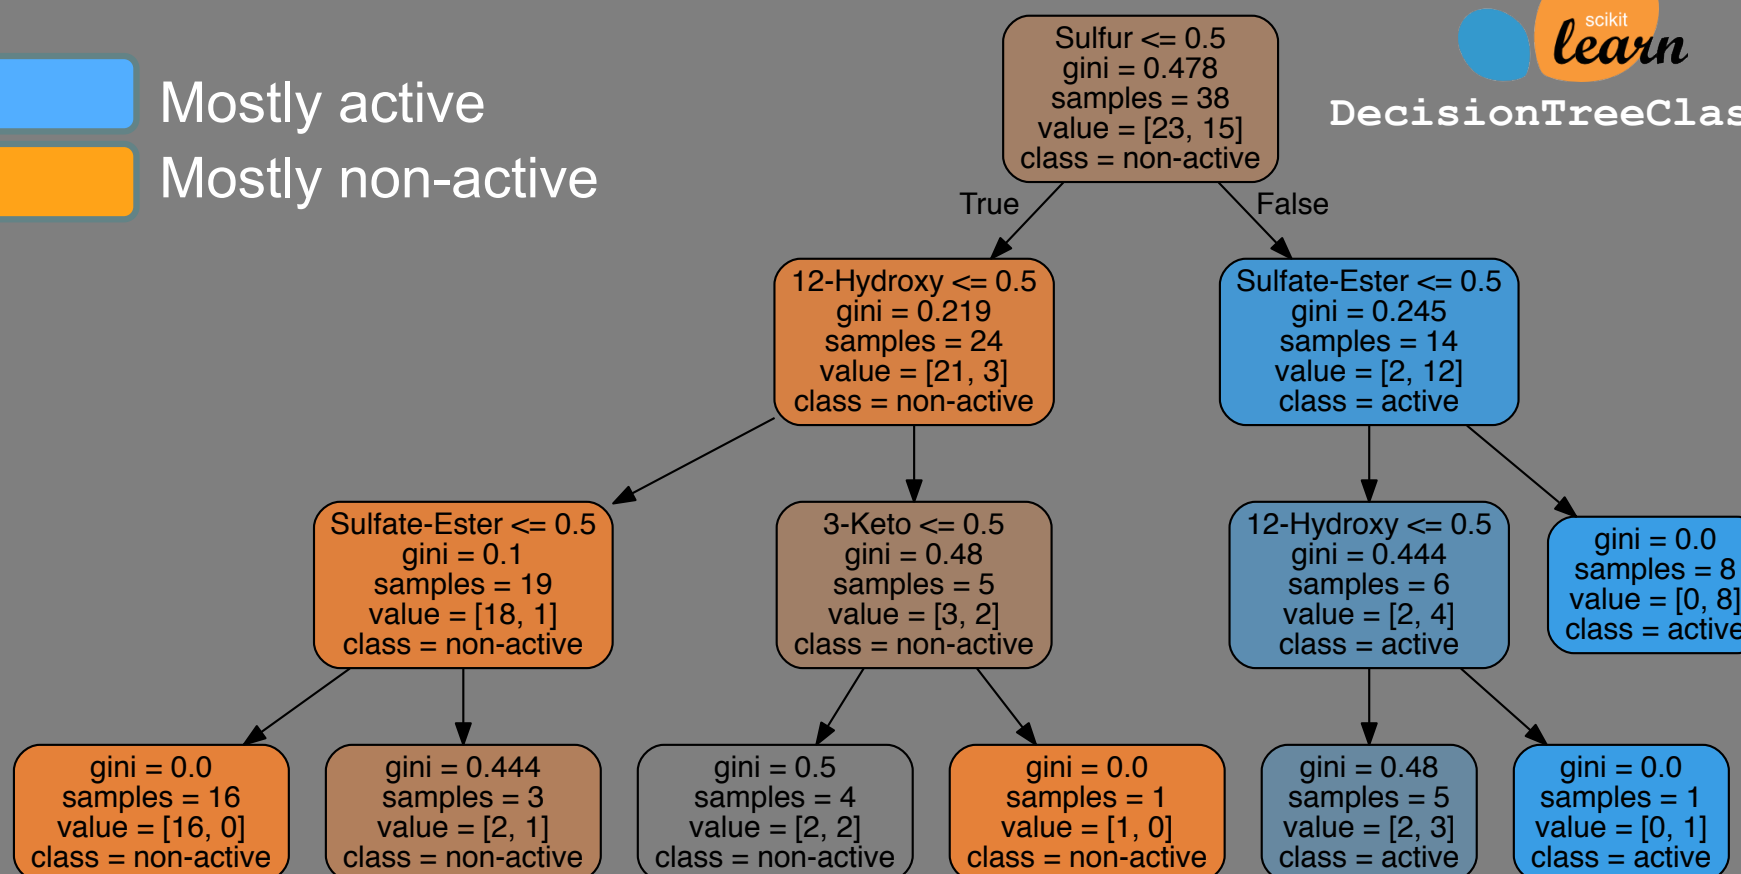


ZINC40576706: 0.0 %

Mostly active

Mostly non-active

DecisionTreeClassifier



# Decision tree-based diagnosis of coronary artery disease: CART model

6

M.M. Ghiasi, S. Zendehboudi and A.A. Mohsenipour / Computer Methods and Programs in Biomedicine 192 (2020) 105400

Mohammad M. Ghiasi <sup>a</sup>✉, Sohrab Zendehboudi <sup>a</sup>, Ali Asghar Mohsenipour <sup>b</sup>

Show more ✓




Fig. 2. Graphical representation of the CART model (using all features) introduced for CAD diagnosis.

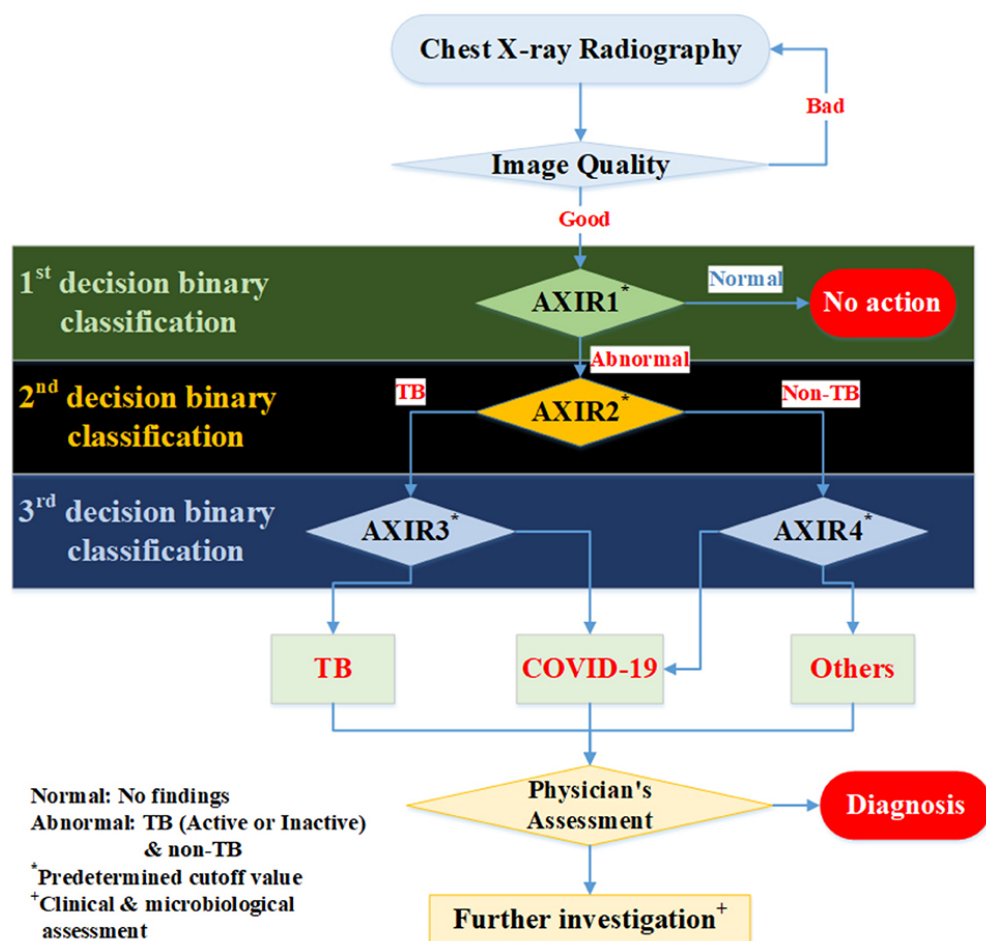
<https://www.sciencedirect.com/science/article/abs/pii/S0169260719308971>



# Deep Learning-Based Decision-Tree Classifier for COVID-19 Diagnosis From Chest X-ray Imaging

Seung Hoon Yoo<sup>1</sup>, Hui Geng<sup>1</sup>, Tin Lok Chiu<sup>1</sup>, Siu Ki Yu<sup>1</sup>, Dae Chul Cho<sup>2</sup>, Jin Heo<sup>2</sup>, Min Sung Choi<sup>2</sup>, Il Hyun Choi<sup>2</sup>, Cong Cung Van<sup>3</sup>, Nguen Viet Nhung<sup>3</sup>, Byung Jun Min<sup>4\*</sup> and Ho Lee<sup>5\*</sup>

<https://www.frontiersin.org/articles/10.3389/fmed.2020.00427/full>







## **Random forests, adaptive boosting, gradient boosting**



# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
- 2. Recursive algorithms & Big-O**
3. Types of decision trees
4. Splitting criteria
5. Gini & Entropy vs misclassification error
6. Improvements & dealing with overfitting
7. Code example

# Recursion / Recursive Algorithms

```
1  def some_fun(x):  
2      if x == []:  
3          return 0  
4      else:  
5          return 1 + some_fun(x[1:])
```

What does this function do?

# Divide & Conquer Algorithms: Quicksort

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12        return quicksort(smaller) + [pivot] + quicksort(bigger)
```

# Divide & Conquer Algorithms: Quicksort

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12        return quicksort(smaller) + [pivot] + quicksort(bigger)
```



# Time complexity of quicksort:

$O(\underline{\hspace{10em}})$

```
1 def quicksort(array):
2     if len(array) < 2:
3         return array
4     else:
5         pivot = array[0]
6         smaller, bigger = [], []
7         for ele in array[1:]:
8             if ele <= pivot:
9                 smaller.append(ele)
10            else:
11                bigger.append(ele)
12        return quicksort(smaller) + [pivot] + quicksort(bigger)
```



# Array Sorting Algorithms

Algorithm	Time Complexity			Space Complexity
	Best	Average	Worst*	
<u>Quicksort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n^2)$	$O(\log(n))$
<u>Mergesort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$
<u>Timsort</u>	$\Omega(n)$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$
<u>Heapsort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(1)$
<u>Bubble Sort</u>	$\Omega(n)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Insertion Sort</u>	$\Omega(n)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Selection Sort</u>	$\Omega(n^2)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
<u>Tree Sort</u>	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$O(n^2)$	$O(n)$
<u>Shell Sort</u>	$\Omega(n \log(n))$	$\Theta(n(\log(n))^2)$	$O(n(\log(n))^2)$	$O(1)$
<u>Bucket Sort</u>	$\Omega(n+k)$	$\Theta(n+k)$	$O(n^2)$	$O(n)$
<u>Radix Sort</u>	$\Omega(nk)$	$\Theta(nk)$	$O(nk)$	$O(n+k)$
<u>Counting Sort</u>	$\Omega(n+k)$	$\Theta(n+k)$	$O(n+k)$	$O(k)$
<u>Cubesort</u>	$\Omega(n)$	$\Theta(n \log(n))$	$O(n \log(n))$	$O(n)$

<http://www.bigocheatsheet.com>

\* "worst" ~ inversely-sorted array

# Decision Tree in Pseudocode

GenerateTree( $\mathcal{D}$ ):

- if  $y = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$  or  $y = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$  :
  - return Tree
- else:
  - Pick best feature  $x_j$ :
    - $\mathcal{D}_0$  at Child<sub>0</sub> :  $x_j = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$
    - $\mathcal{D}_1$  at Child<sub>1</sub> :  $x_j = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$

return Node( $x_j$ , GenerateTree( $\mathcal{D}_0$ ), GenerateTree( $\mathcal{D}_1$ ))

# Time Complexity ("Big-O")

Growing the tree:  $O(\dots)$

Tip: It can be shown that optimal split is on boundary between adjacent examples (similar feature value) with different class labels.  $\longrightarrow$  Consider sorting

Fayyad, Usama Mohammad. "On the induction of decision trees for multiple concept learning." (1992).

Sorting a feature column :  $\mathcal{O}(n \log n)$

Consider/sort  $m$  feature columns:  $\mathcal{O}(m)$

Decision tree has up to  $n$  terminal leaf nodes

There are  $2n - 1$  internal nodes in the tree

The number we split nodes is  $2n - 1 - n = n - 1$

Complexity with re-sorting:  $\mathcal{O}(mn^2 \log n)$

Complexity with sorting once & caching:  $\mathcal{O}(mn \log n)$

# Time Complexity ("Big-O")

Querying the tree:  $\mathcal{O}(\dots)$



# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
2. Recursive algorithms & Big-O
- 3. Types of decision trees**
4. Splitting criteria
5. Gini & Entropy vs misclassification error
6. Improvements & dealing with overfitting
7. Code example

# Decision Tree in Pseudocode

GenerateTree( $\mathcal{D}$ ):

- if  $y = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$  or  $y = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$  :
  - return Tree
- else:
  - Pick best feature  $x_j$ :
    - $\mathcal{D}_0$  at Child<sub>0</sub> :  $x_j = 0 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$
    - $\mathcal{D}_1$  at Child<sub>1</sub> :  $x_j = 1 \ \forall \ \langle \mathbf{x}, \mathbf{y} \rangle \in \mathcal{D}$

return Node( $x_j$ , GenerateTree( $\mathcal{D}_0$ ), GenerateTree( $\mathcal{D}_1$ ))

# Generic Tree Growing Algorithm

- 1)** Pick the feature that, when parent node is split, results in the largest information gain
- 2)** Stop if child nodes are pure or information gain  $\leq 0$
- 3)** Go back to step 1 for each of the two child nodes

# Generic Tree Growing Algorithm

- 1) Pick the feature that, when parent node is split, results in the largest information gain
  - 2) Stop if child nodes are pure or information gain  $\leq 0$
  - 3) Go back to step 1 for each of the two child nodes
- How make predictions if features in dataset are not sufficient to make child nodes pure?

# Design choices

- How to split
  - what measurement/criterion as measure of 'goodness'
  - binary vs multi-category split
- When to stop
  - if leaf nodes contain only examples of the same class
  - feature values are all the same for all examples
  - maximum number of splits
  - score threshold or statistical significance test



# ID3 -- Iterative Dichotomizer 3

- one of the earlier decision tree algorithms
- Quinlan, J. R. 1986. Induction of Decision Trees. Mach. Learn. 1, 1 (Mar. 1986), 81-106.
- cannot handle numeric features
- no pruning, prone to overfitting
- short and wide trees (compared to CART)
- maximizing information gain/minimizing entropy
- discrete features, binary and multi-category features

# C4.5

- continuous and discrete features
- Ross Quinlan 1993, Quinlan, J. R. (1993).  
C4.5: Programming for machine learning.  
*Morgan Kauffmann*, 38, 48.
- continuous is very expensive, because must consider all possible ranges
- handles missing attributes (ignores them in gain compute)
- post-pruning (bottom-up pruning)
- Gain Ratio

# CART

- Breiman, L. (1984). *Classification and regression trees*. Belmont, Calif: Wadsworth International Group.
- continuous and discrete features
- strictly binary splits (taller trees than ID3, C4.5)
- binary splits can generate better trees than C4.5, but tend to be larger and harder to interpret; k-attributes has a ways to create a binary partitioning
- variance reduction in regression trees
- Gini impurity, twoing criteria in classification trees
- cost complexity pruning

# Others

- CHAID (CHi-squared Automatic Interaction Detector); Kass, G. V. (1980). "An exploratory technique for investigating large quantities of categorical data". *Applied Statistics*. 29 (2): 119–127.
- MARS (Multivariate adaptive regression splines); Friedman, J. H. (1991). "Multivariate Adaptive Regression Splines". *The Annals of Statistics*. 19: 1
- C5.0 (patented)
- ...

# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
2. Recursive algorithms & Big-O
3. Types of decision trees
- 4. Splitting criteria**
5. Gini & Entropy vs misclassification error
6. Improvements & dealing with overfitting
7. Code example





# Finding a Decision Rule

$x_1$	$x_2$	$x_3$	$y$
6 cm	8 cm	9 cm	1
4 cm	11 cm	2 cm	0
6 cm	12 cm	4 cm	0
10 cm	9 cm	3 cm	1
5 cm	7 cm	8 cm	0
8 cm	9 cm	3 cm	1
3 cm	11 cm	5 cm	0

# Drawing a Decision Boundary

x <sub>1</sub>	x <sub>2</sub>	x <sub>3</sub>	y
6 cm	8 cm	9 cm	1
4 cm	11 cm	2 cm	0
6 cm	12 cm	4 cm	0
10 cm	9 cm	3 cm	1
5 cm	7 cm	8 cm	0
8 cm	9 cm	3 cm	1
3 cm	11 cm	5 cm	0





$x_1 \leq 5.5$   
 entropy=0.985  
 samples=7  
 value=[4,3]  
 class=Class0

True


False


entropy=0.0  
 samples=3  
 value=[3,0]  
 class=Class0

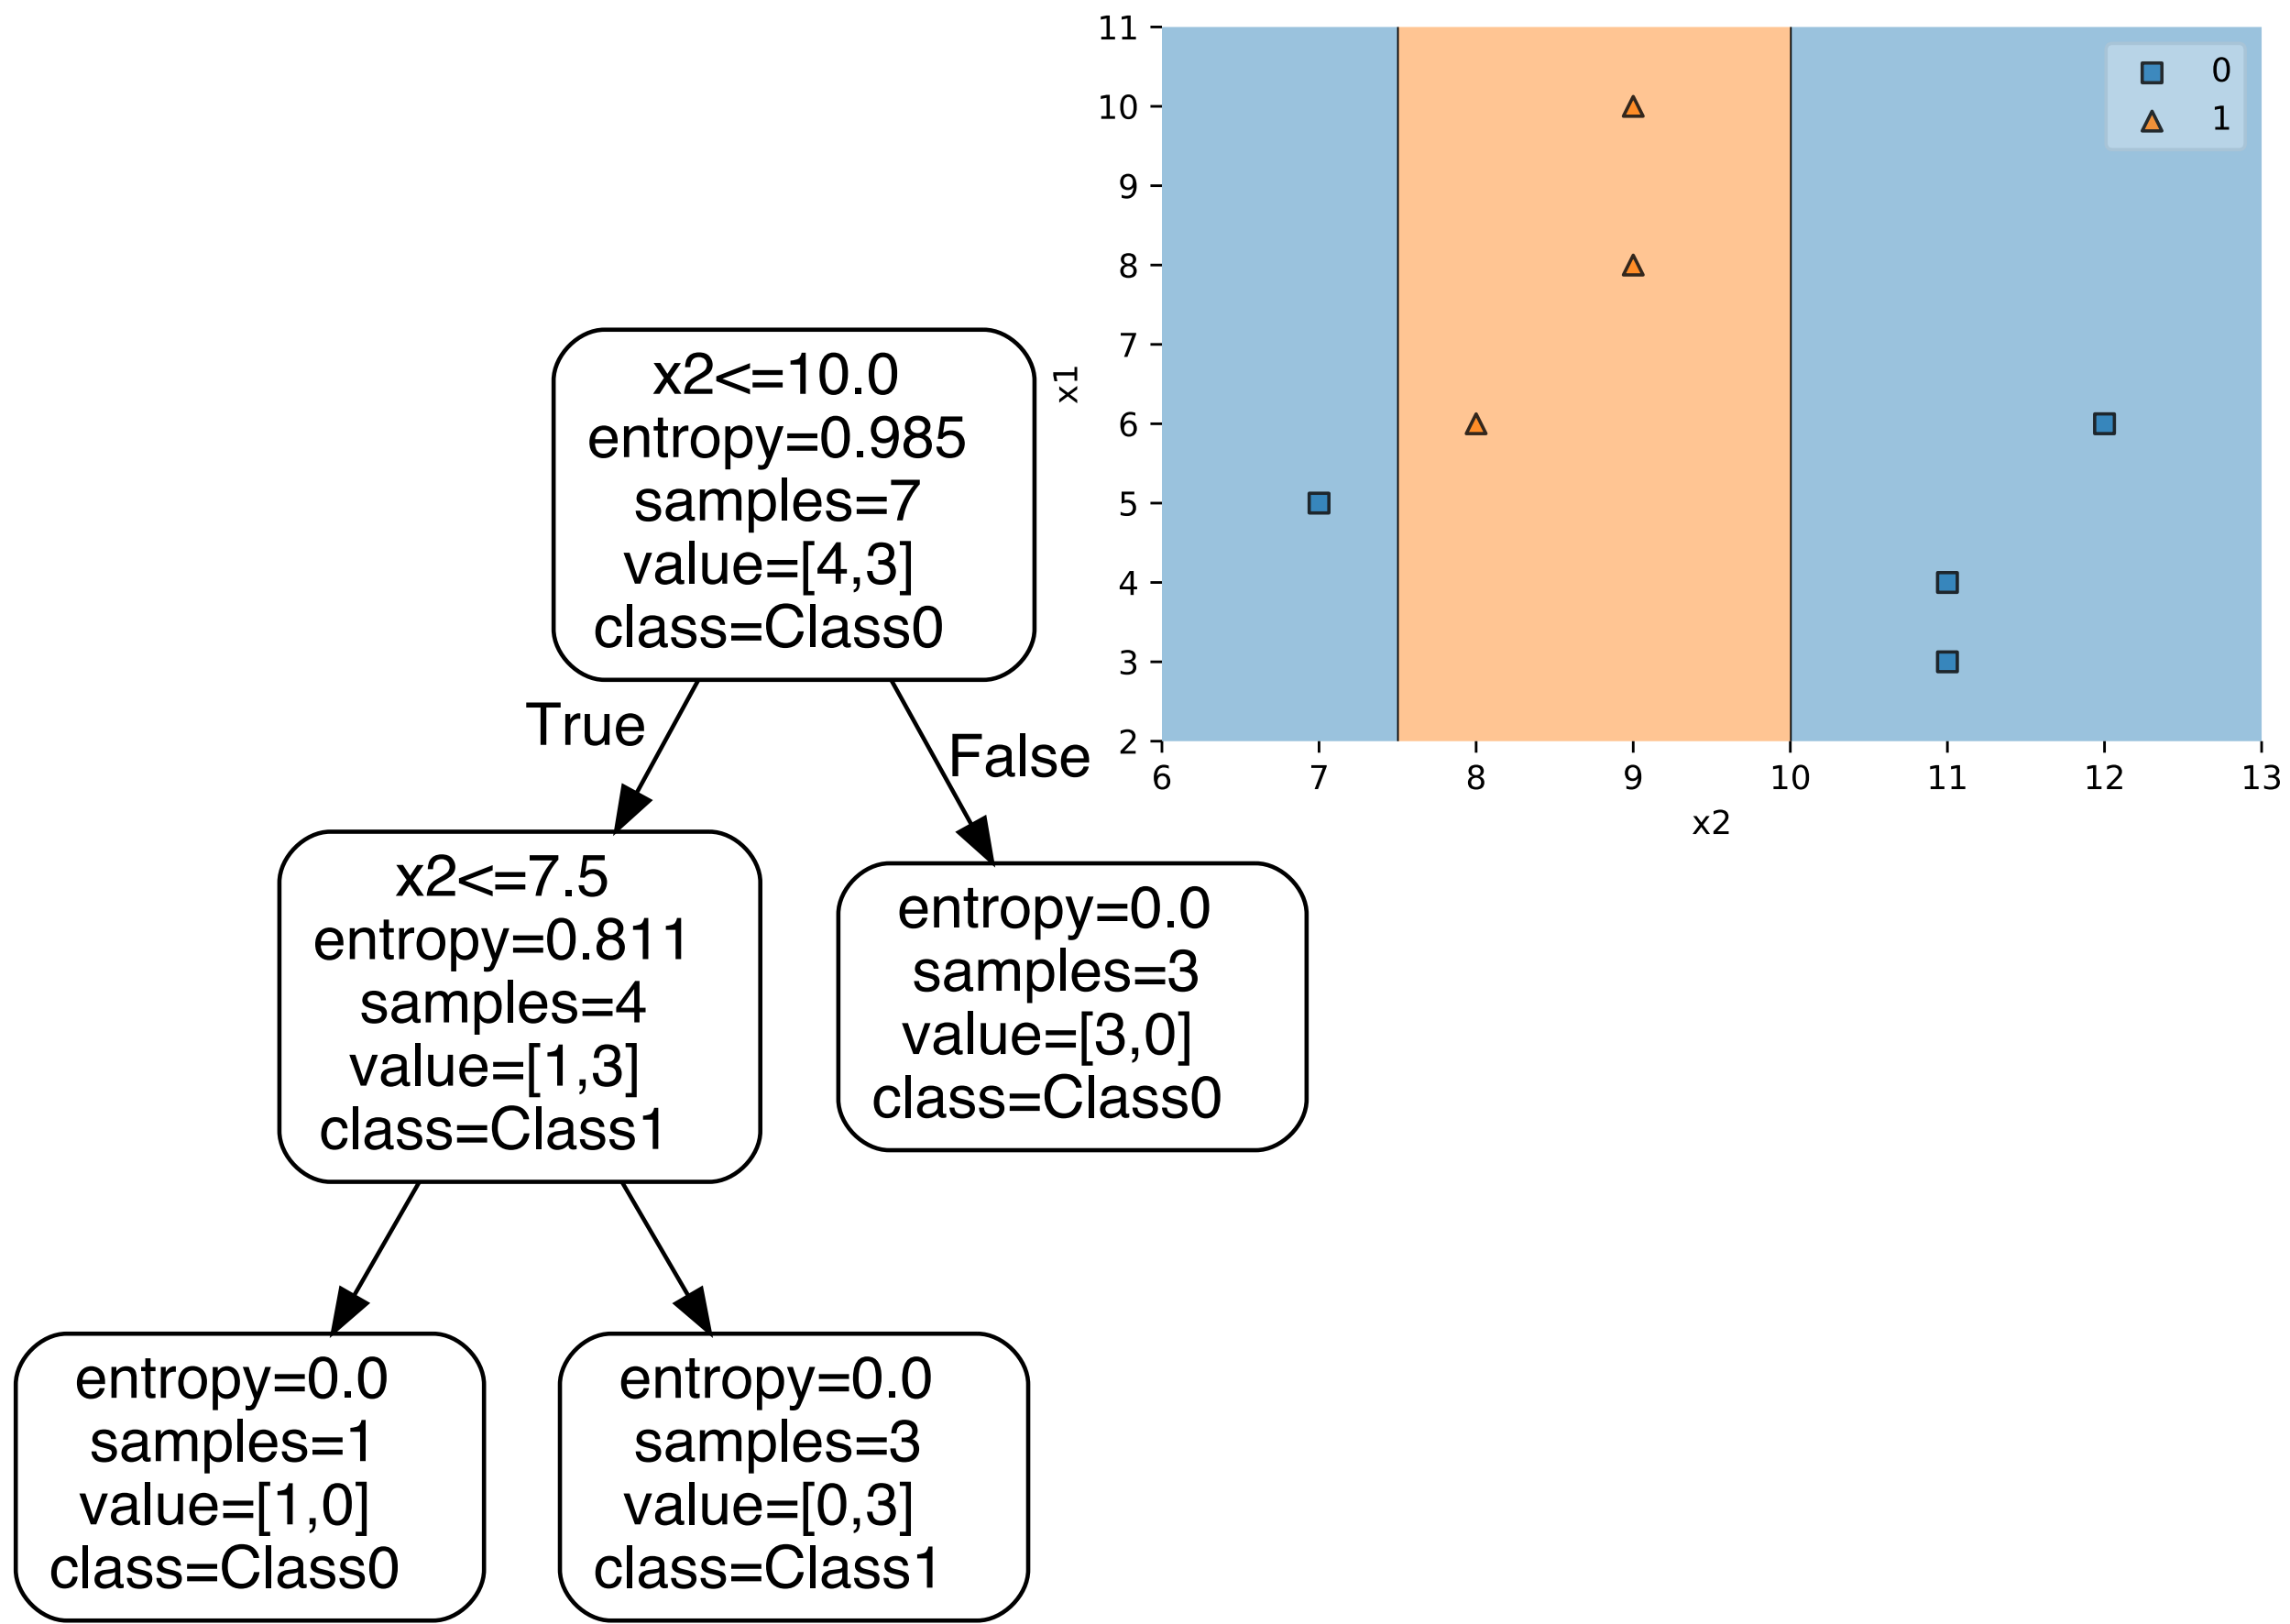
$x_2 \leq 10.5$   
 entropy=0.811  
 samples=4  
 value=[1,3]  
 class=Class1

entropy=0.0  
 samples=3  
 value=[0,3]  
 class=Class1

entropy=0.0  
 samples=1  
 value=[1,0]  
 class=Class0







# The Splitting Criterion

# Information Gain

$$GAIN(\mathcal{D}, x_j) = H(\mathcal{D}) - \sum_{v \in Values(x_j)} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} H(\mathcal{D}_v)$$




# (Shannon) Entropy

$$H = - \sum_i p(i | x_j) \log_2(p(i | x_j))$$

# (Shannon) Entropy

$$H = - \sum_i p(i | x_j) \log_2(p(i | x_j))$$



# Information Gain

$$GAIN(\mathcal{D}, x_j) = H(\mathcal{D}) - \sum_{v \in Values(x_j)} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} H(\mathcal{D}_v)$$

# Gini Impurity

$$Gini = 1 - \sum_i (p(i | x_j))^2$$


# Misclassification Error

$$ERR = \frac{1}{n} \sum_{i=1}^n L(\hat{y}^{[i]}, y^{[i]}),$$

$$L(\hat{y}, y) = \begin{cases} 0 & \text{if } \hat{y} = y, \\ 1 & \text{otherwise.} \end{cases}$$

# Misclassification Error

$$ERR = 1 - \max_i(p(i | x_j))$$




# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
2. Recursive algorithms & Big-O
3. Types of decision trees
4. Splitting criteria
- 5. Gini & Entropy vs misclassification error**
6. Improvements & dealing with overfitting
7. Code example





# **Why Growing Decision Trees via Entropy instead of Misclassification Error?**

# Why Growing Decision Trees via Entropy instead of Misclassification Error?

$$GAIN(\mathcal{D}, x_j) = I(\mathcal{D}) - \sum_{v \in Values(x_j)} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} I(\mathcal{D}_v)$$

# Entropy

$$H = - \sum_i p(i | x_j) \log_2(p(i | x_j))$$


# Gini Impurity


$$Gini = 1 - \sum_i (p(i | x_j))^2$$

# Misclassification Error

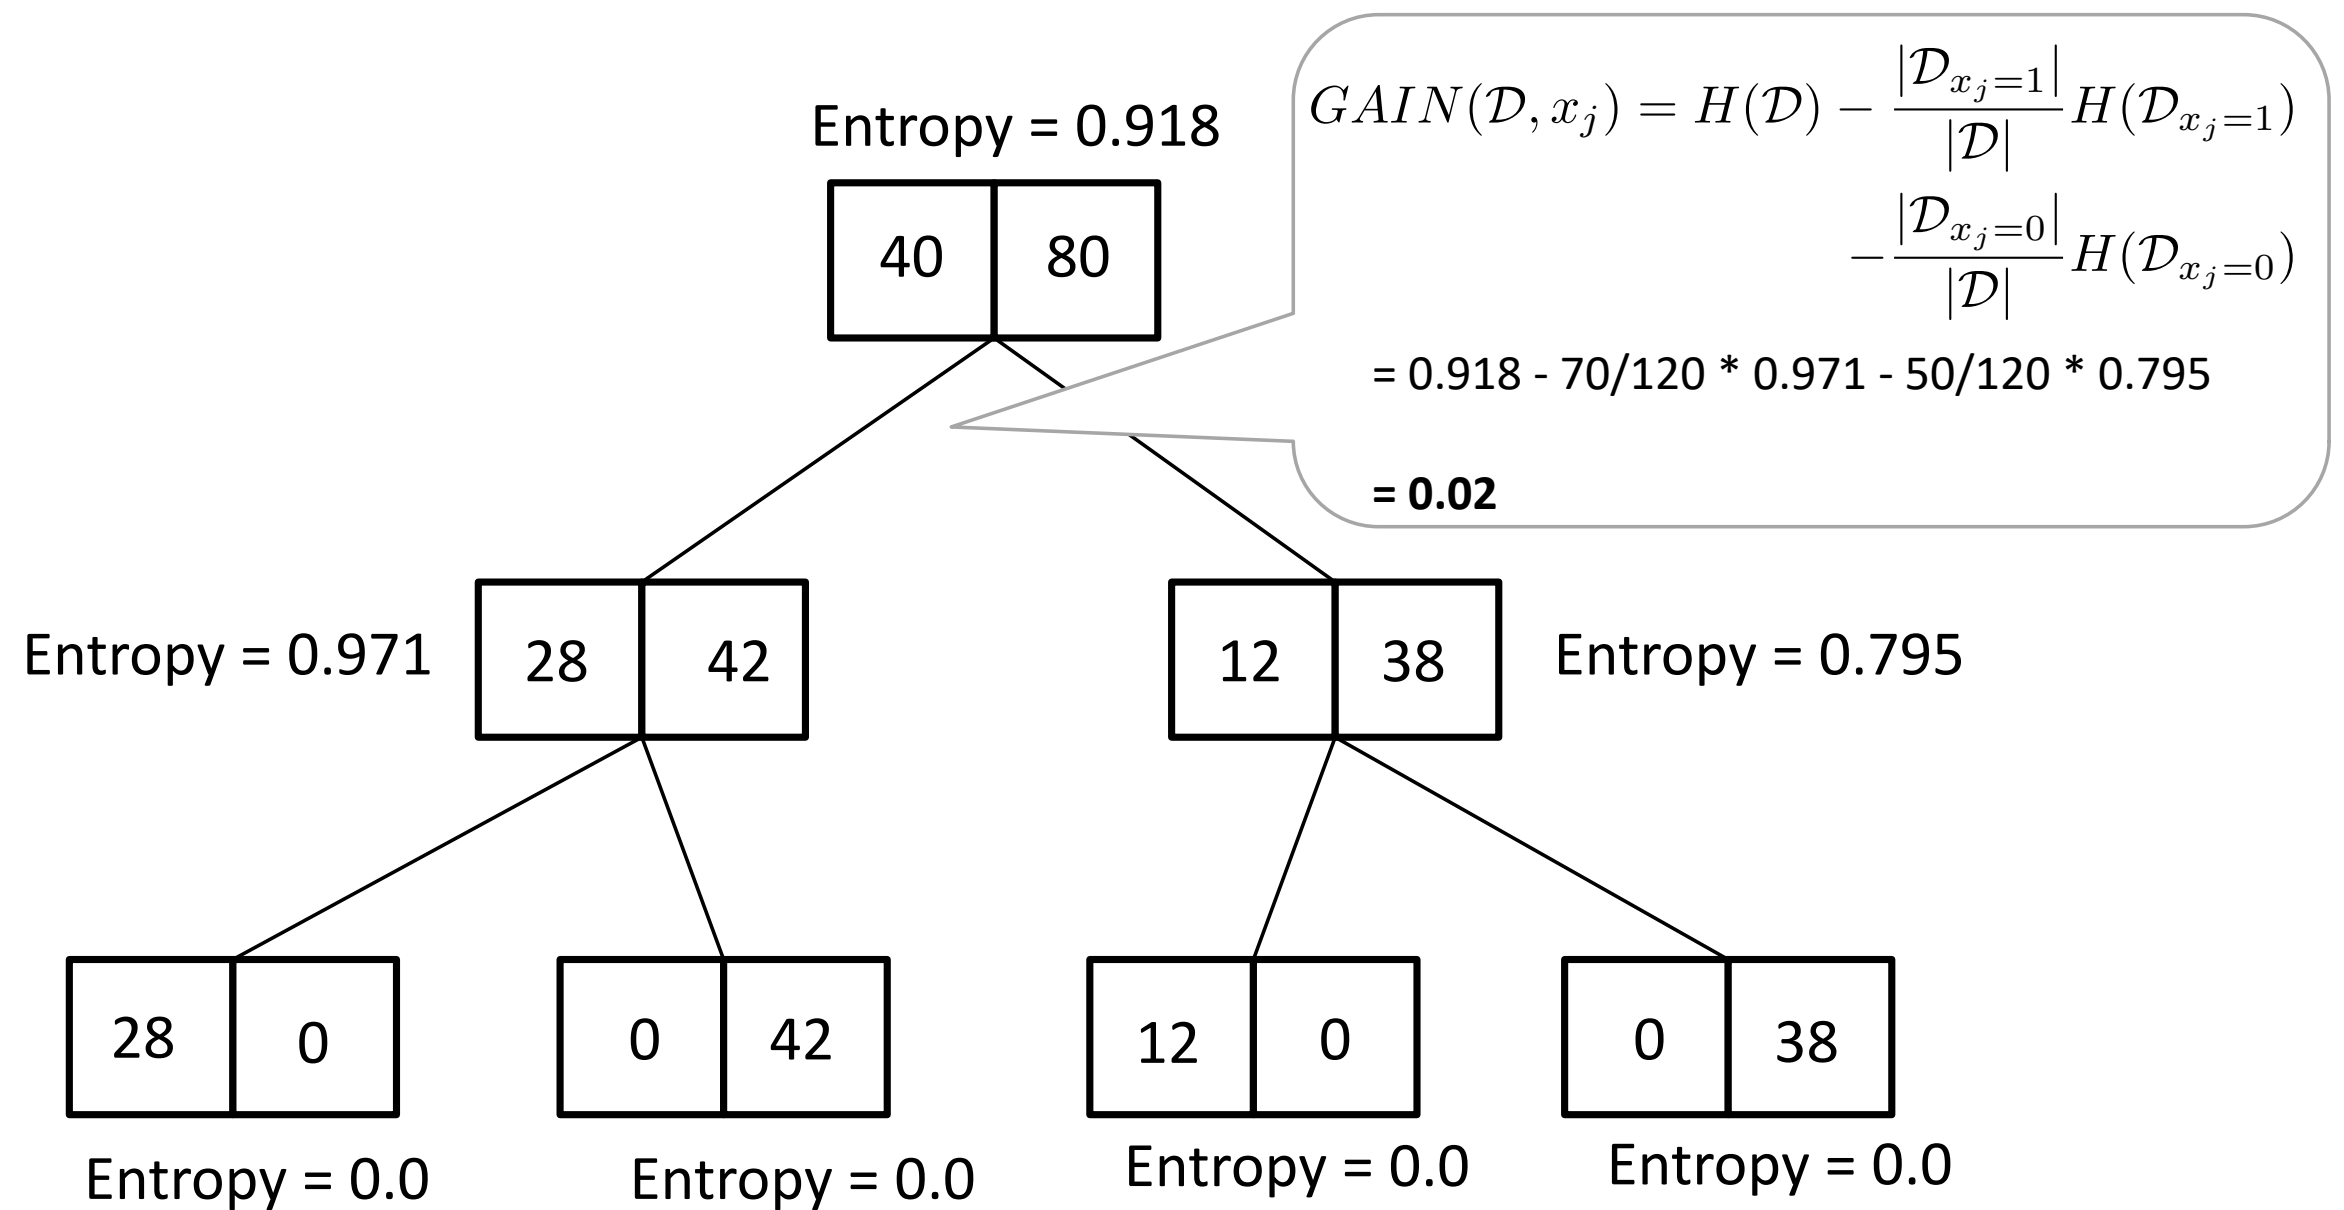
$$ERR = \frac{1}{n} \sum_{i=1}^n L(\hat{y}^{[i]}, y^{[i]}),$$

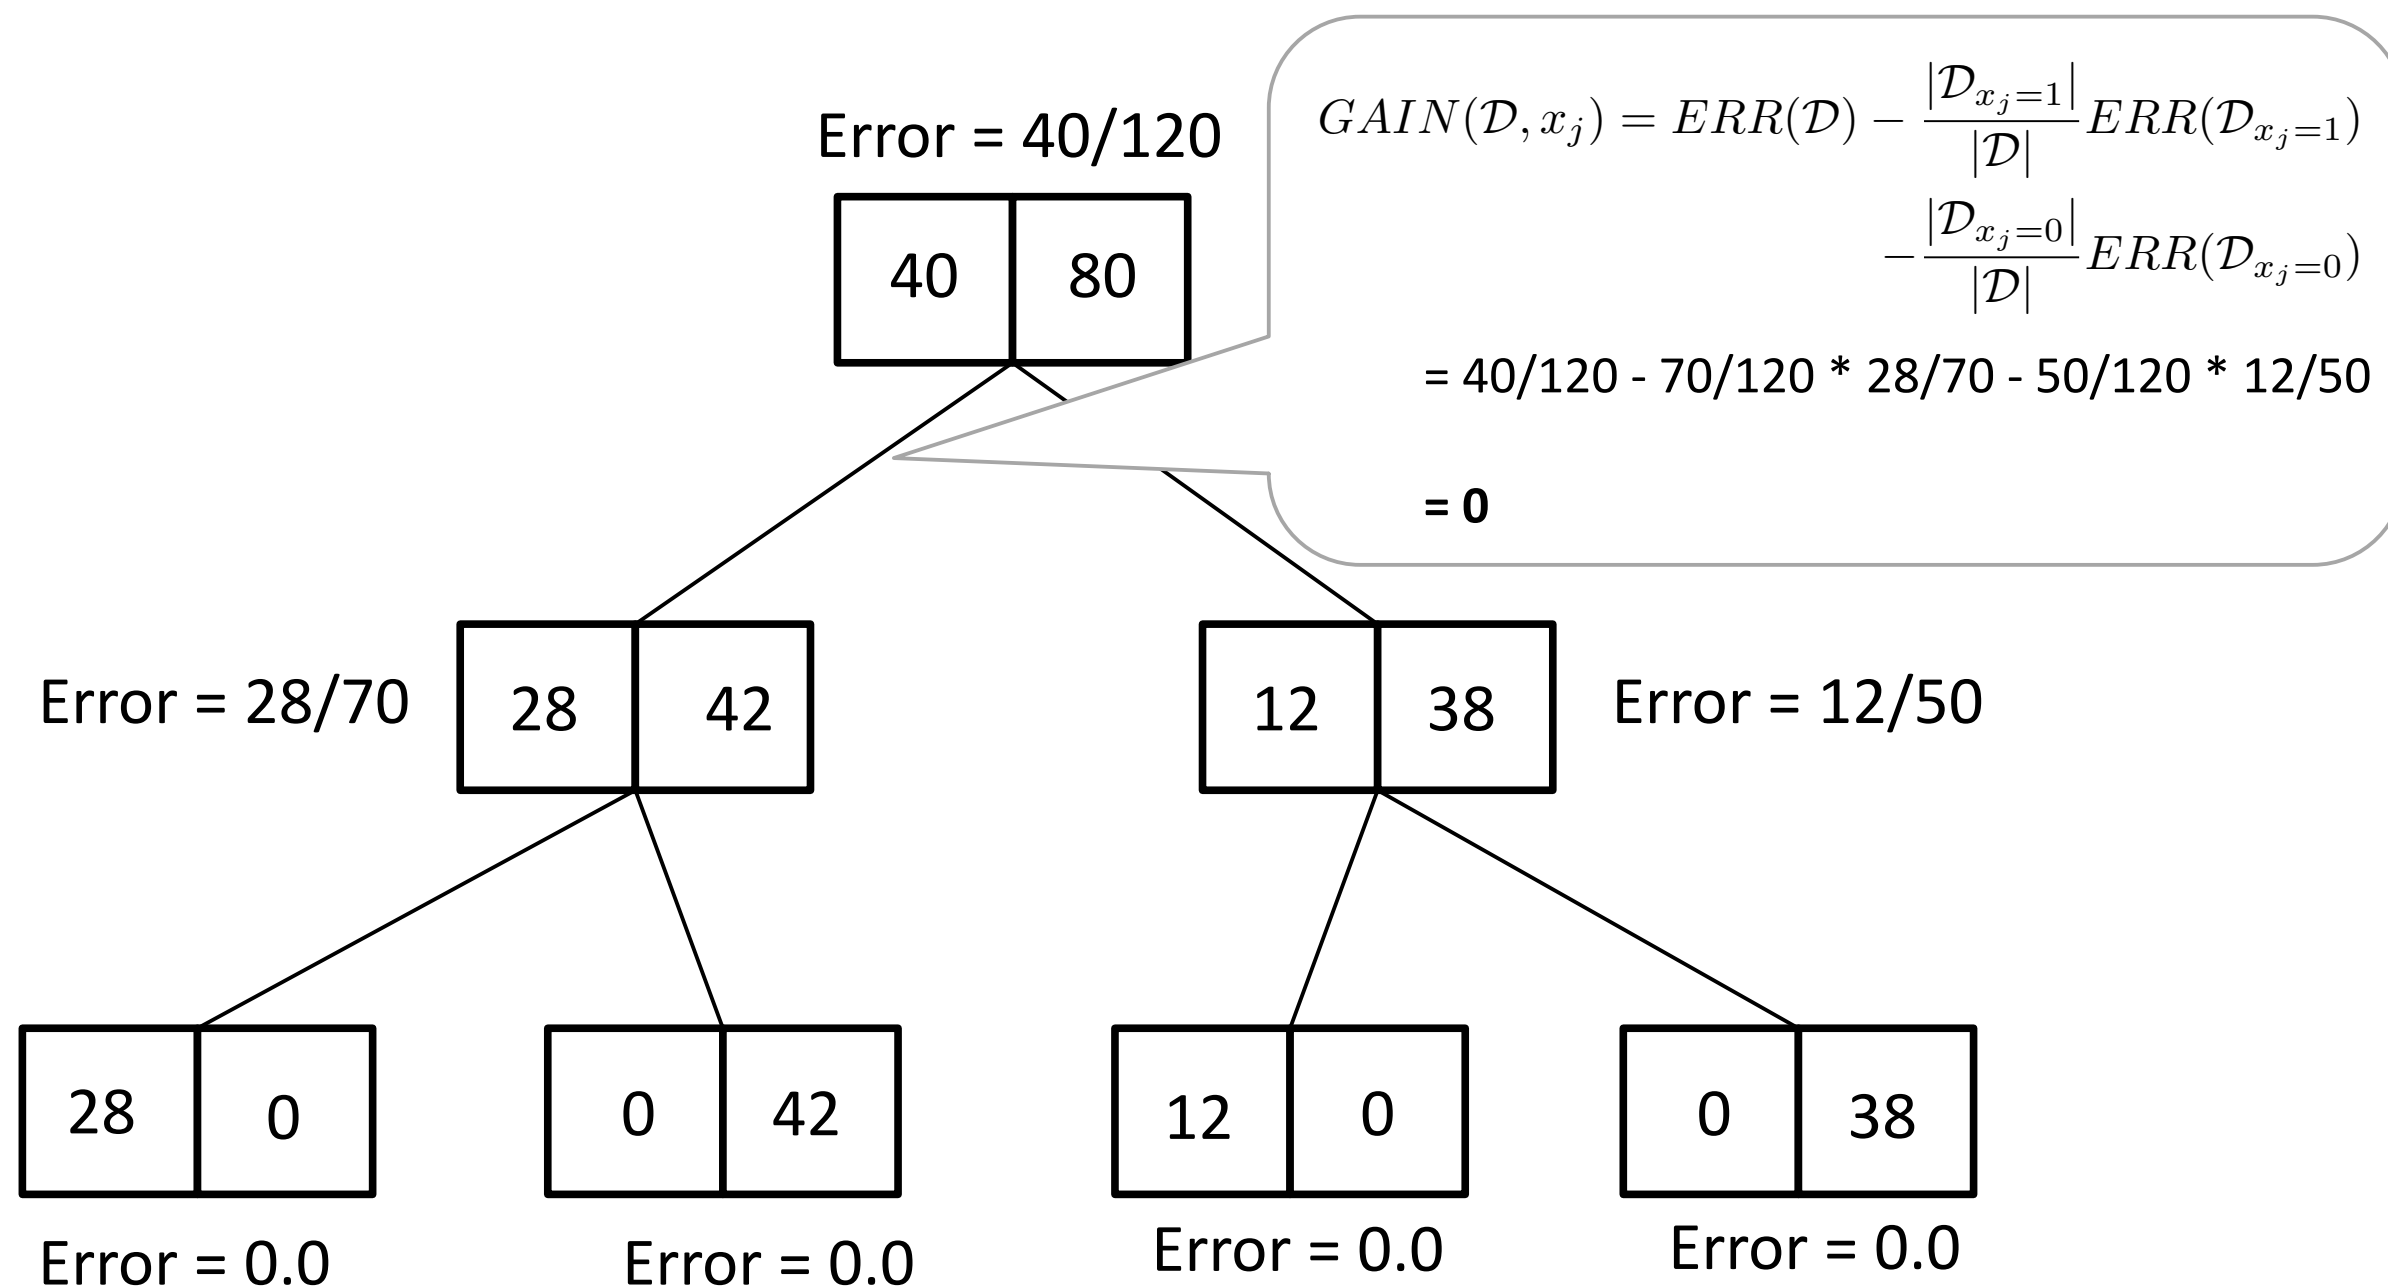
$$L(\hat{y}, y) = \begin{cases} 0 & \text{if } \hat{y} = y, \\ 1 & \text{otherwise.} \end{cases}$$














# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
2. Recursive algorithms & Big-O
3. Types of decision trees
4. Splitting criteria
5. Gini & Entropy vs misclassification error
- 6. Improvements & dealing with overfitting**
7. Code example

# Gain Ratio

$$\textit{GainRatio}(\mathcal{D}, x_j) = \frac{\textit{Gain}(\mathcal{D}, x_j)}{\textit{SplitInfo}(\mathcal{D}, x_j)}$$


Quinlan 1986

where the split information measures the entropy of the feature:

$$\textit{SplitInfo}(\mathcal{D}, x_j) = - \sum_{v \in x_j} \frac{|\mathcal{D}_v|}{|\mathcal{D}|} \log_2 \frac{|\mathcal{D}_v|}{|\mathcal{D}|}$$


Penalizes splitting categorical attributes with many values (e.g., think date column, or really bad: row ID) via the split information

# Shortcomings



## How would the decision tree split look like?

# Overfitting



# Pre-Pruning


- Set a depth cut-off (maximum tree depth) *a priori*
- Cost-complexity pruning: , where is an impurity measure, is a tuning parameter, and is the total number of nodes.
- Stop growing if split is not statistically significant (e.g.,  $\chi^2$  test)
- Set a minimum number of data points for each node



# Post-Pruning

- Grow full tree first, then remove nodes, in C4.5
- Reduced-error pruning, remove nodes via validation set eval. (problematic for limited data)
- Can also convert trees to rules first and then prune the rules

# Post-Pruning



# Regression Trees

# Decision Tree Summary: Pros and Cons

- (+) Easy to interpret and communicate
- (+) Can represent "complete" hypothesis space
- (-) Easy to overfit
- (-) Elaborate pruning required
- (-) Expensive to just fit a "diagonal line"
- (-) Output range is bounded (dep. on training examples) in regression trees

# Lecture 6: Decision Trees

## Topics

1. Intro to decision trees
2. Recursive algorithms & Big-O
3. Types of decision trees
4. Splitting criteria
5. Gini & Entropy vs misclassification error
6. Improvements & dealing with overfitting
- 7. Code example**

# Demo