

## Lecture 07

# Ensemble Methods Part 3/3

STAT 451: Machine Learning, Fall 2021

Sebastian Raschka



Uwe L. Korn  
@xhochy

Replying to @rasbt @DrTBehrens and @mervenoyann

Nice to see that the work on the [@condaforge](#) on M1 packages is well appreciated. [@XGBoostProject](#) should also be available in 1-2h (need to wait for the CI on master and the CDN mirror to sync) for the M1: [github.com/conda-forge/xgboost-feedstock/pull/79](https://github.com/conda-forge/xgboost-feedstock/pull/79). No need to compile from source anymore then.

conda-forge/xgboost-feedstock

## #79 Build for osx-arm64

2 comments 0 reviews 13 files +143 -11

xhochy · October 24, 2021 · 3 commits

github.com

Build for osx-arm64 by xhochy · Pull Request #79 · conda-forge/xgboost-feedst... Fixes #72 Checklist Used a personal fork of the feedstock to propose changes Bumped the build number (if the version is unchanged) Reset the build number ...

2:49 PM · Oct 24, 2021 · Twitter Web App

4 Retweets 2 Quote Tweets 26 Likes

Comment Share Heart

```
[base] sebastian@Sebastian-Air ~/Desktop % conda install xgboost
Collecting package metadata (current_repodata.json): done
Solving environment: done
```

```
## Package Plan ##
```

```
environment location: /Users/sebastian/miniforge3
```

```
added / updated specs:
  - xgboost
```

```
The following packages will be downloaded:
```

package	build		
_py-xgboost-mutex-2.0	cpu_0	12 KB	conda-forge
libxgboost-1.5.0	hcfdfaf5_0	2.0 MB	conda-forge
py-xgboost-1.5.0	py39h2804cbe_0	151 KB	conda-forge
xgboost-1.5.0	py39ha480839_0	12 KB	conda-forge
		Total:	2.2 MB

```
The following NEW packages will be INSTALLED:
```

```
_py-xgboost-mutex  conda-forge/osx-arm64::_py-xgboost-mutex-2.0-cpu_0
libxgboost        conda-forge/osx-arm64::libxgboost-1.5.0-hcfdfaf5_0
py-xgboost         conda-forge/osx-arm64::py-xgboost-1.5.0-py39h2804cbe_0
xgboost           conda-forge/osx-arm64::xgboost-1.5.0-py39ha480839_0
```

```
Proceed ([y]/n)?
```

```
Downloading and Extracting Packages
```

```
xgboost-1.5.0      | 12 KB      | #####
| 100%          |
py-xgboost-1.5.0   | 151 KB     | #####
| 100%          |
_py-xgboost-mutex-2. | 12 KB      | #####
| 100%          |
libxgboost-1.5.0    | 2.0 MB     | #####
| 100%          |
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
```

## The Most Comprehensive List of Kaggle Solutions and Ideas

This is a list of almost all available solutions and ideas shared by top performers in the past Kaggle competitions. This list gets updated as soon as a new competition finishes. If you find a solution besides the ones listed here, I would encourage you to contribute to this repo by making a pull request. The symbols used in this list are described [here](#).

If you found it interesting, you may give a star or make a fork

 Star 536     Fork 141

Check out the following markdown pages about Top Kagglers Tips/Tricks and all Kernels of The week.

- [Top Kagglers Interviews and Lectures](#)
- [Kernels of The Week](#)

Type / to search title, description, year, ranking and ...

 Last Updated: October 25, 2021

Title & Description	Details	Solutions	Pins
 <a href="#">458. Google Landmark Retrieval 2021</a> Given an image, can you find all of the same landmarks in a dataset?	Prize: Swag Team: 263 Kind: Research Metric: - Year: 2021	 1st place  2nd place  3rd place  5th place  5th place  6th place  11th place  13th place  all solutions	<input checked="" type="checkbox"/>
 <a href="#">457. Google Landmark Recognition 2021</a> Label famous, and not-so-famous, landmarks in images	Prize: Swag Team: 383 Kind: Research Metric: GoogleGlobalAP Year: 2021	 1st place  3rd place  4th place  5th place  8th place  all solutions	<input checked="" type="checkbox"/>
 <a href="#">456. LearnPlatform COVID-19 Impact on Digital Learning</a> Use digital learning data to analyze the impact of COVID-19	Prize: \$20,000 Team: 0 Kind: Analytics	 Not Available!	<input type="radio"/>

<https://farid.one/kaggle-solutions/>



The image shows the Kaggle Tabular Playground Series - May 2021 competition page. At the top, there's a decorative banner with binary code and silhouettes of people. Below it, the title "Tabular Playground Series - May 2021" is displayed in bold black text. A sub-header says "Practice your ML skills on this approachable dataset!". A "Kaggle" logo indicates 1,097 teams participated 5 months ago. Navigation links include Overview, Data, Code, Discussion (which is underlined), Leaderboard, Rules, and a "New Topic" button. A "..." button is also present.

### [1st place] Solution description

Posted in [tabular-playground-series-may-2021](#) 5 months ago

60

<https://www.kaggle.com/c/tabular-playground-series-may-2021/discussion/243054>

## Base Models

I ended up training 5 different models, each trained on a 5fold stratified cv on three random seeds, using out of fold predictions for each fold and averaging across folds for predicting the test set. The models were LogisticsRegression, RandomForest, XGBoost, LightGBM, and CATBoost. Weighting classes did not really work for LogisticsRegression and I used the defaults class weights in the rest of them. I used Optuna to lightly tune each model (you will see my final parameters in the code above).

## Over/Under Sampling

Did not really try after reading the notebook from [@remekkinas](#) that pointed out it likely would not work. His analysis made sense to me and I did not have time, so did not really pursue any further.

## Ensembling

Stacking worked quite well here. I tried weighted average of models and it did worse than stacking so I did not use that. I used a 5fold stacking of all models above and used a meta model of RidgeRegerssion, using CalibratedClassifierCV in sklearn.



Playground Prediction Competition

# March Machine Learning Mania 2021 - NCAAM

Predict the 2021 NCAA Basketball Tournament



Kaggle · 707 teams · 7 months ago

Overview Data Code **Discussion** Leaderboard Rules

New Topic

...

<https://www.kaggle.com/c/ncaam-march-mania-2021/discussion/231665>

## Summary

We used a random forest model to predict winning outcomes. We considered end of season summarized statistics, key wins / losses, specific Massey Ordinals, and efficiency metrics. We tuned our random forest parameters by selecting the parameters with the best classification accuracy (should have used log loss in hindsight but it worked out for us). We trained multiple models on all seasons through 2014, testing on 2015-2019, and chose the model with the lowest Kaggle score (log-loss).



Featured Code Competition

# Prostate cANcer graDe Assessment (PANDA) Challenge

Prostate cancer diagnosis using the Gleason grading system

\$25,000

Prize Money



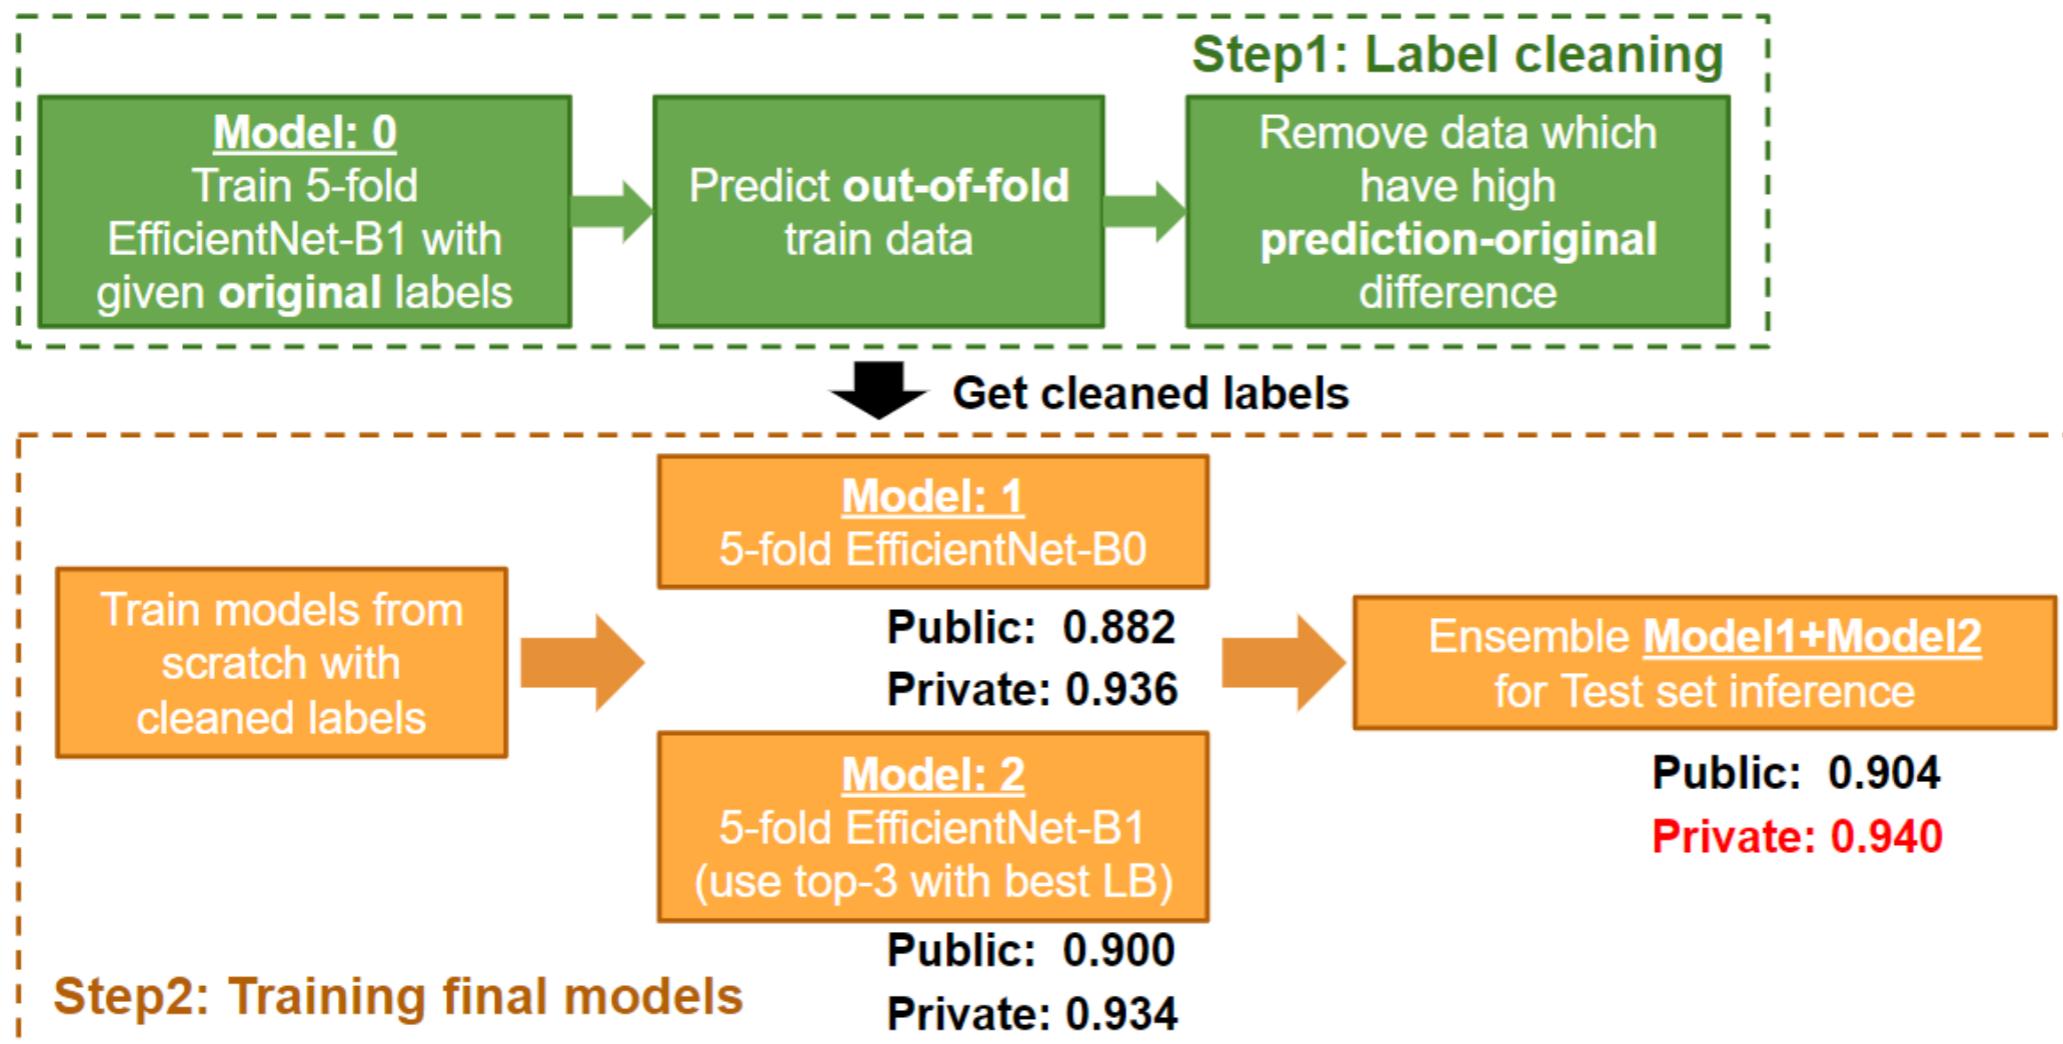
PANDA Challenge · 1,010 teams · a year ago

Overview Data Code Discussion Leaderboard Rules

Join Competition

...

<https://www.kaggle.com/c/prostate-cancer-grade-assessment>





Featured Code Competition

Twitter. It's what's happening  
https://t.co/...

# Tweet Sentiment Extraction

Extract support phrases for sentiment labels

\$15,000

Prize Money



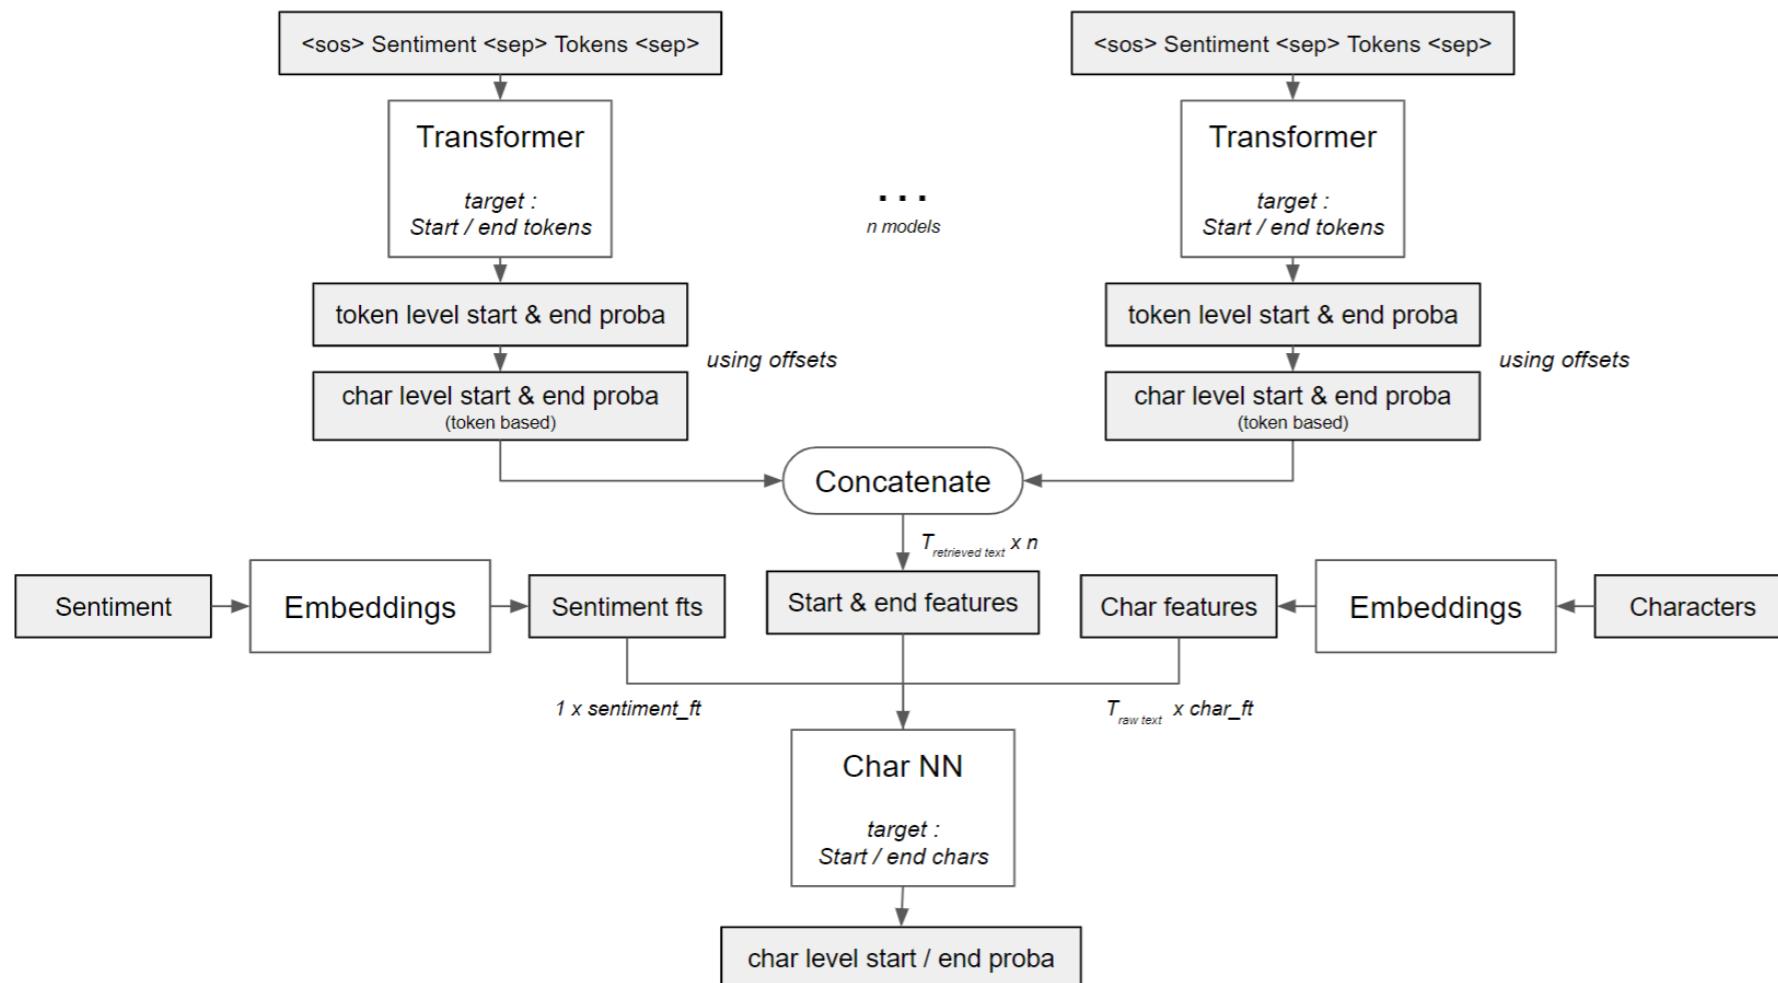
Kaggle · 2,225 teams · a year ago

Overview Data Code **Discussion** Leaderboard Rules

New Topic

...

<https://www.kaggle.com/c/tweet-sentiment-extraction/discussion/159477>



## Lecture 07

# Ensemble Methods Part 3/3

STAT 451: Machine Learning, Fall 2021

Sebastian Raschka

7.1 Ensemble Methods -- Intro and Overview

7.2 Majority Voting

7.3 Bagging

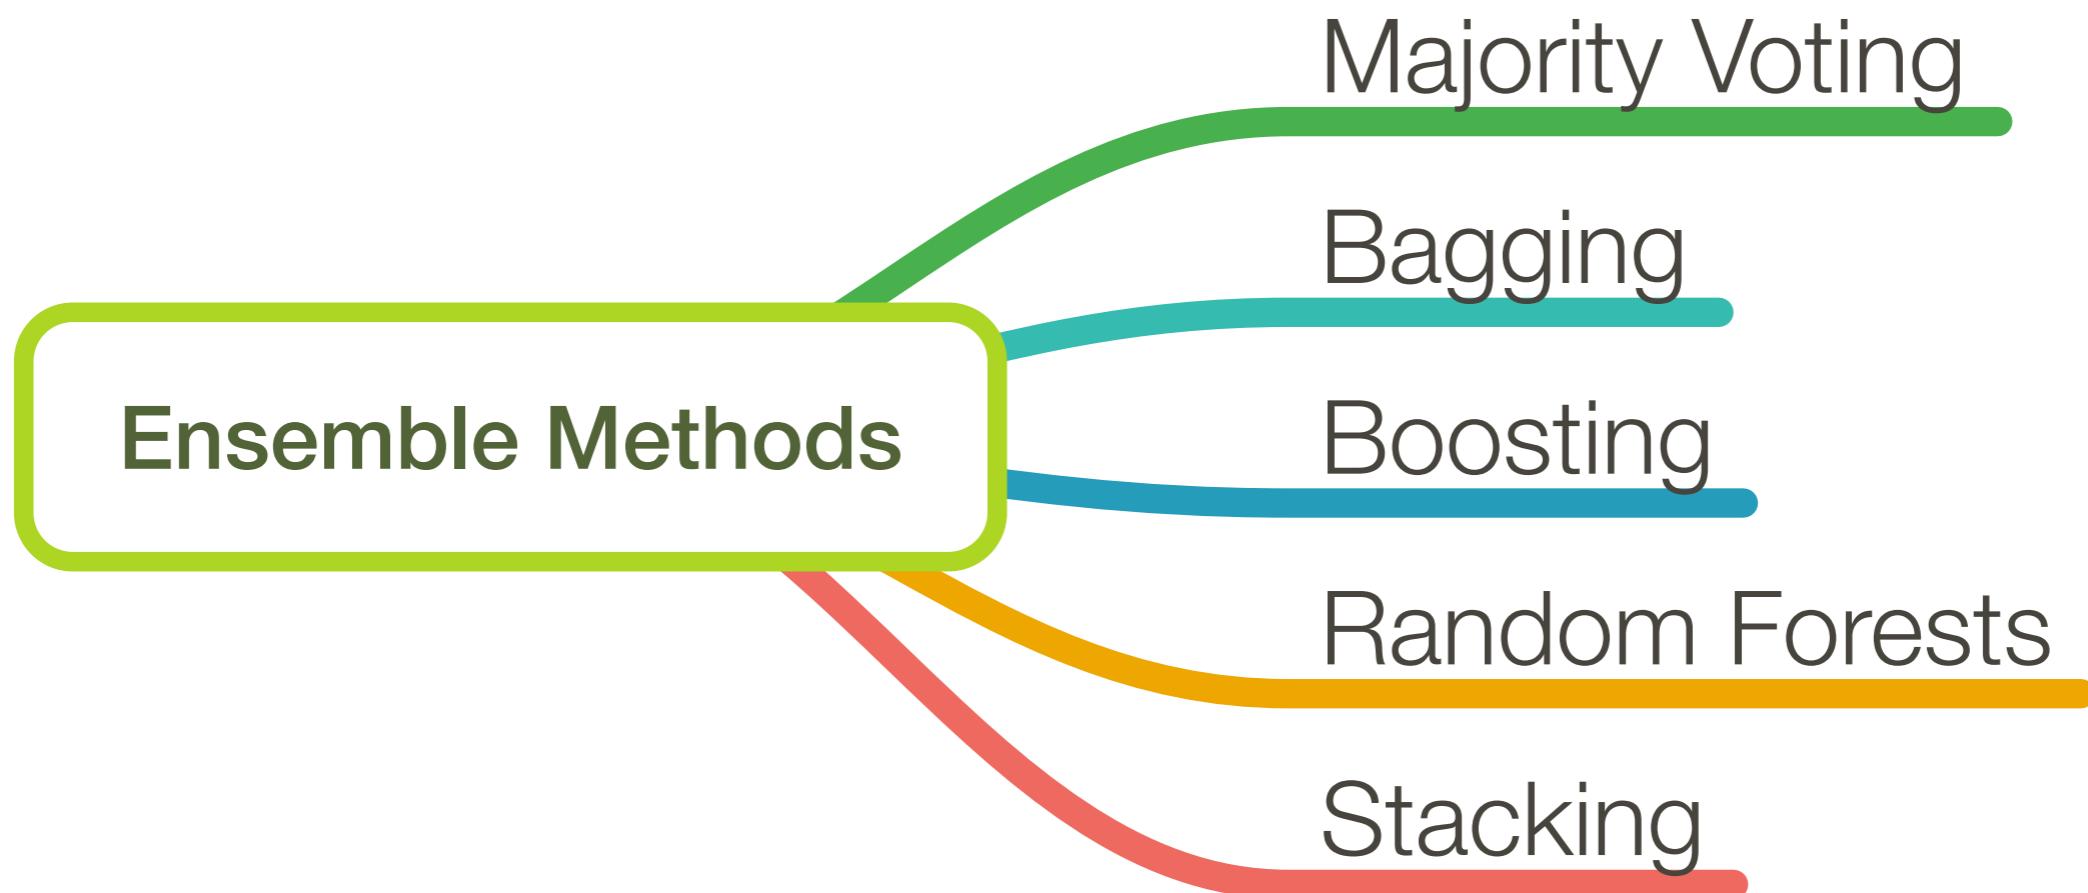
7.4 Boosting

7.5 Gradient Boosting

**7.6 Random Forests**

7.7 Stacking

# Overview

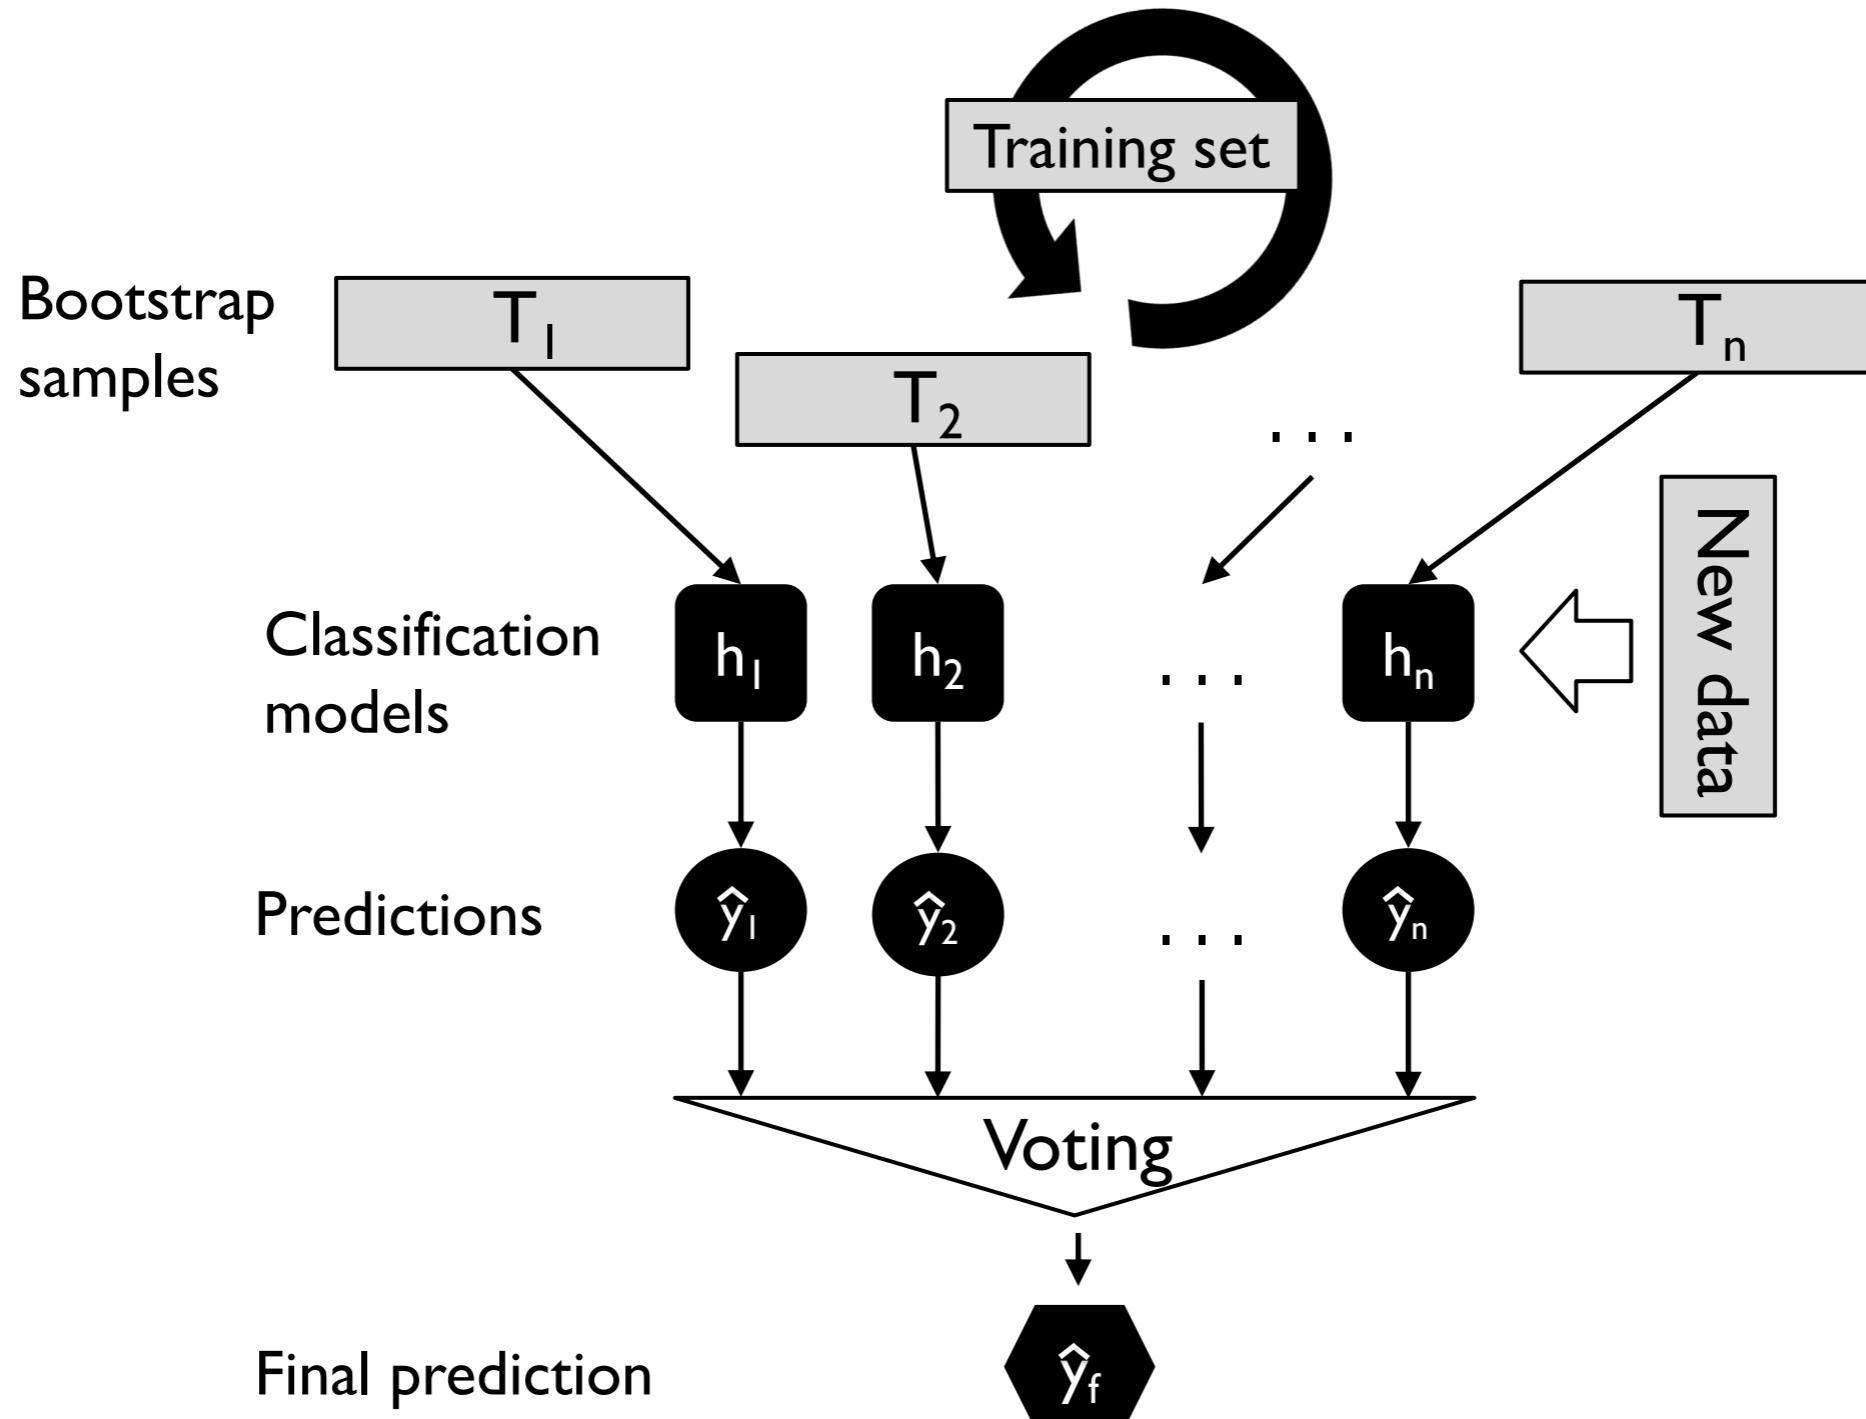


# Random Forests

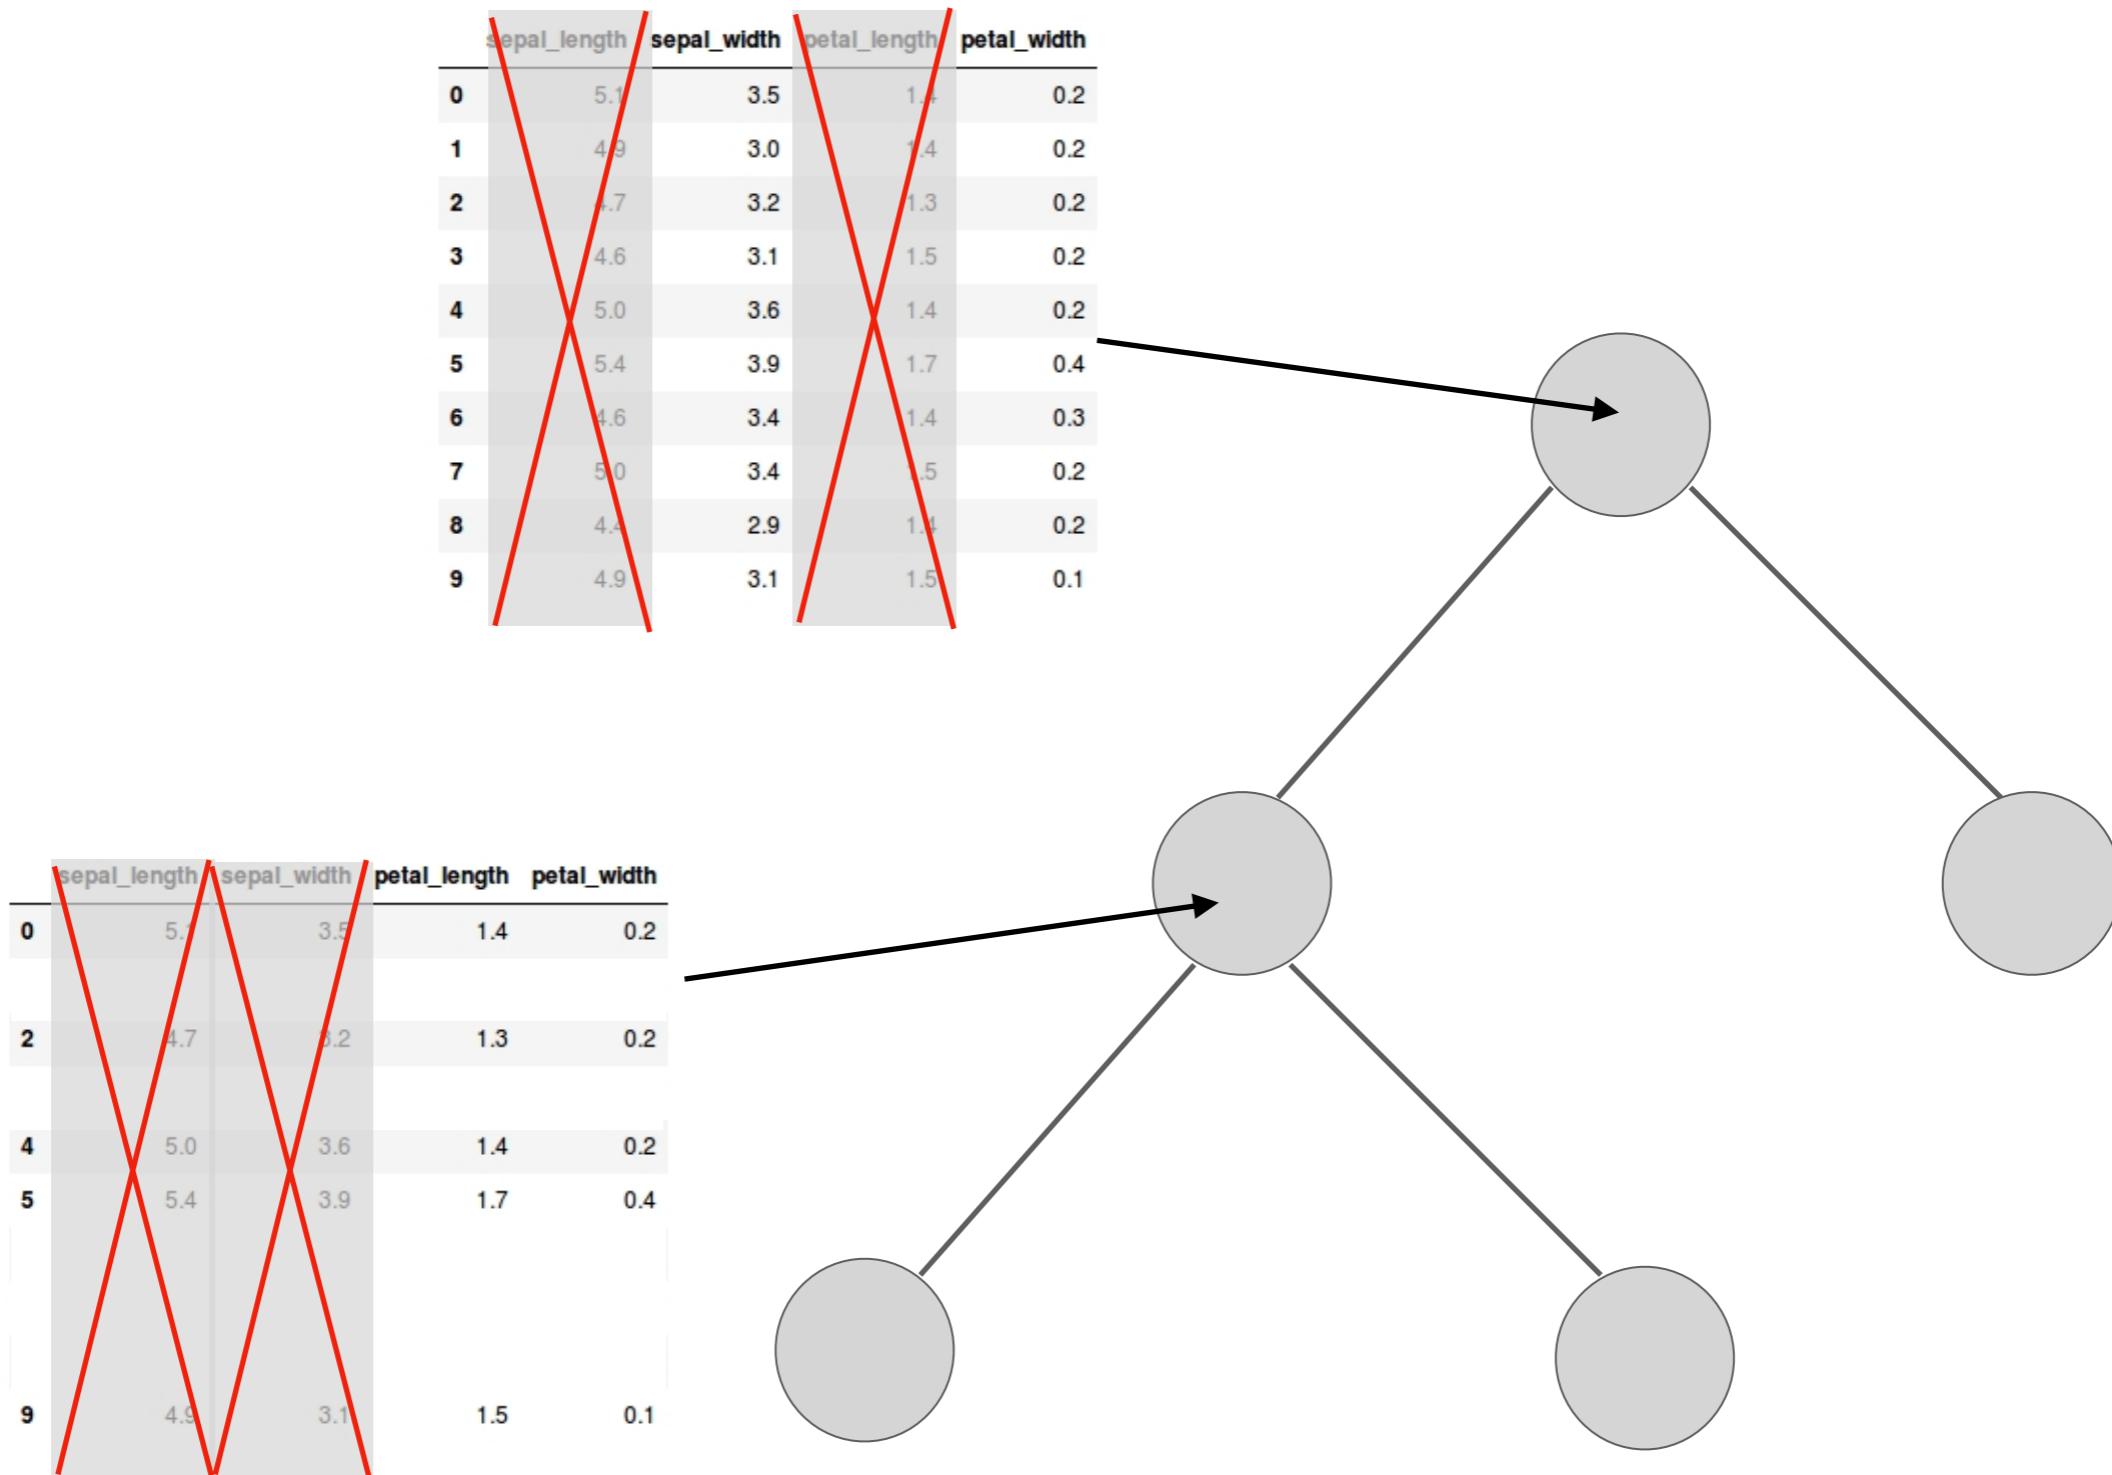
# **Random Forests**

= Bagging w. trees + random feature subsets

# Bagging Classifier



# Random Feature Subsets



# Random Feature Subset for each Tree or Node?

Tin Kam Ho used the “**random subspace method**,” where each tree got a random subset of features.

“Our method relies on an autonomous, pseudo-random procedure to select a small number of dimensions from a given feature space ...”

- Ho, Tin Kam. “The random subspace method for constructing decision forests.” IEEE transactions on pattern analysis and machine intelligence 20.8 (1998): 832-844.

“**Trademark**” random forest:

“... *random forest with random features is formed by selecting at random, at each node, a small group of input variables to split on.*”

- Breiman, Leo. “Random Forests” Machine learning 45.1 (2001): 5-32.

# Random Feature Subset for each Tree or Node?

Tin Kam Ho used the “**random subspace method**,” where each tree got a random subset of features.

“Our method relies on an autonomous, pseudo-random procedure to select a small number of dimensions from a given feature space ...”

- Ho, Tin Kam. “The random subspace method for constructing decision forests.” IEEE transactions on pattern analysis and machine intelligence 20.8 (1998): 832-844.

## “Trademark” random forest:

“... random forest with random feature at each node, a small group of input variables

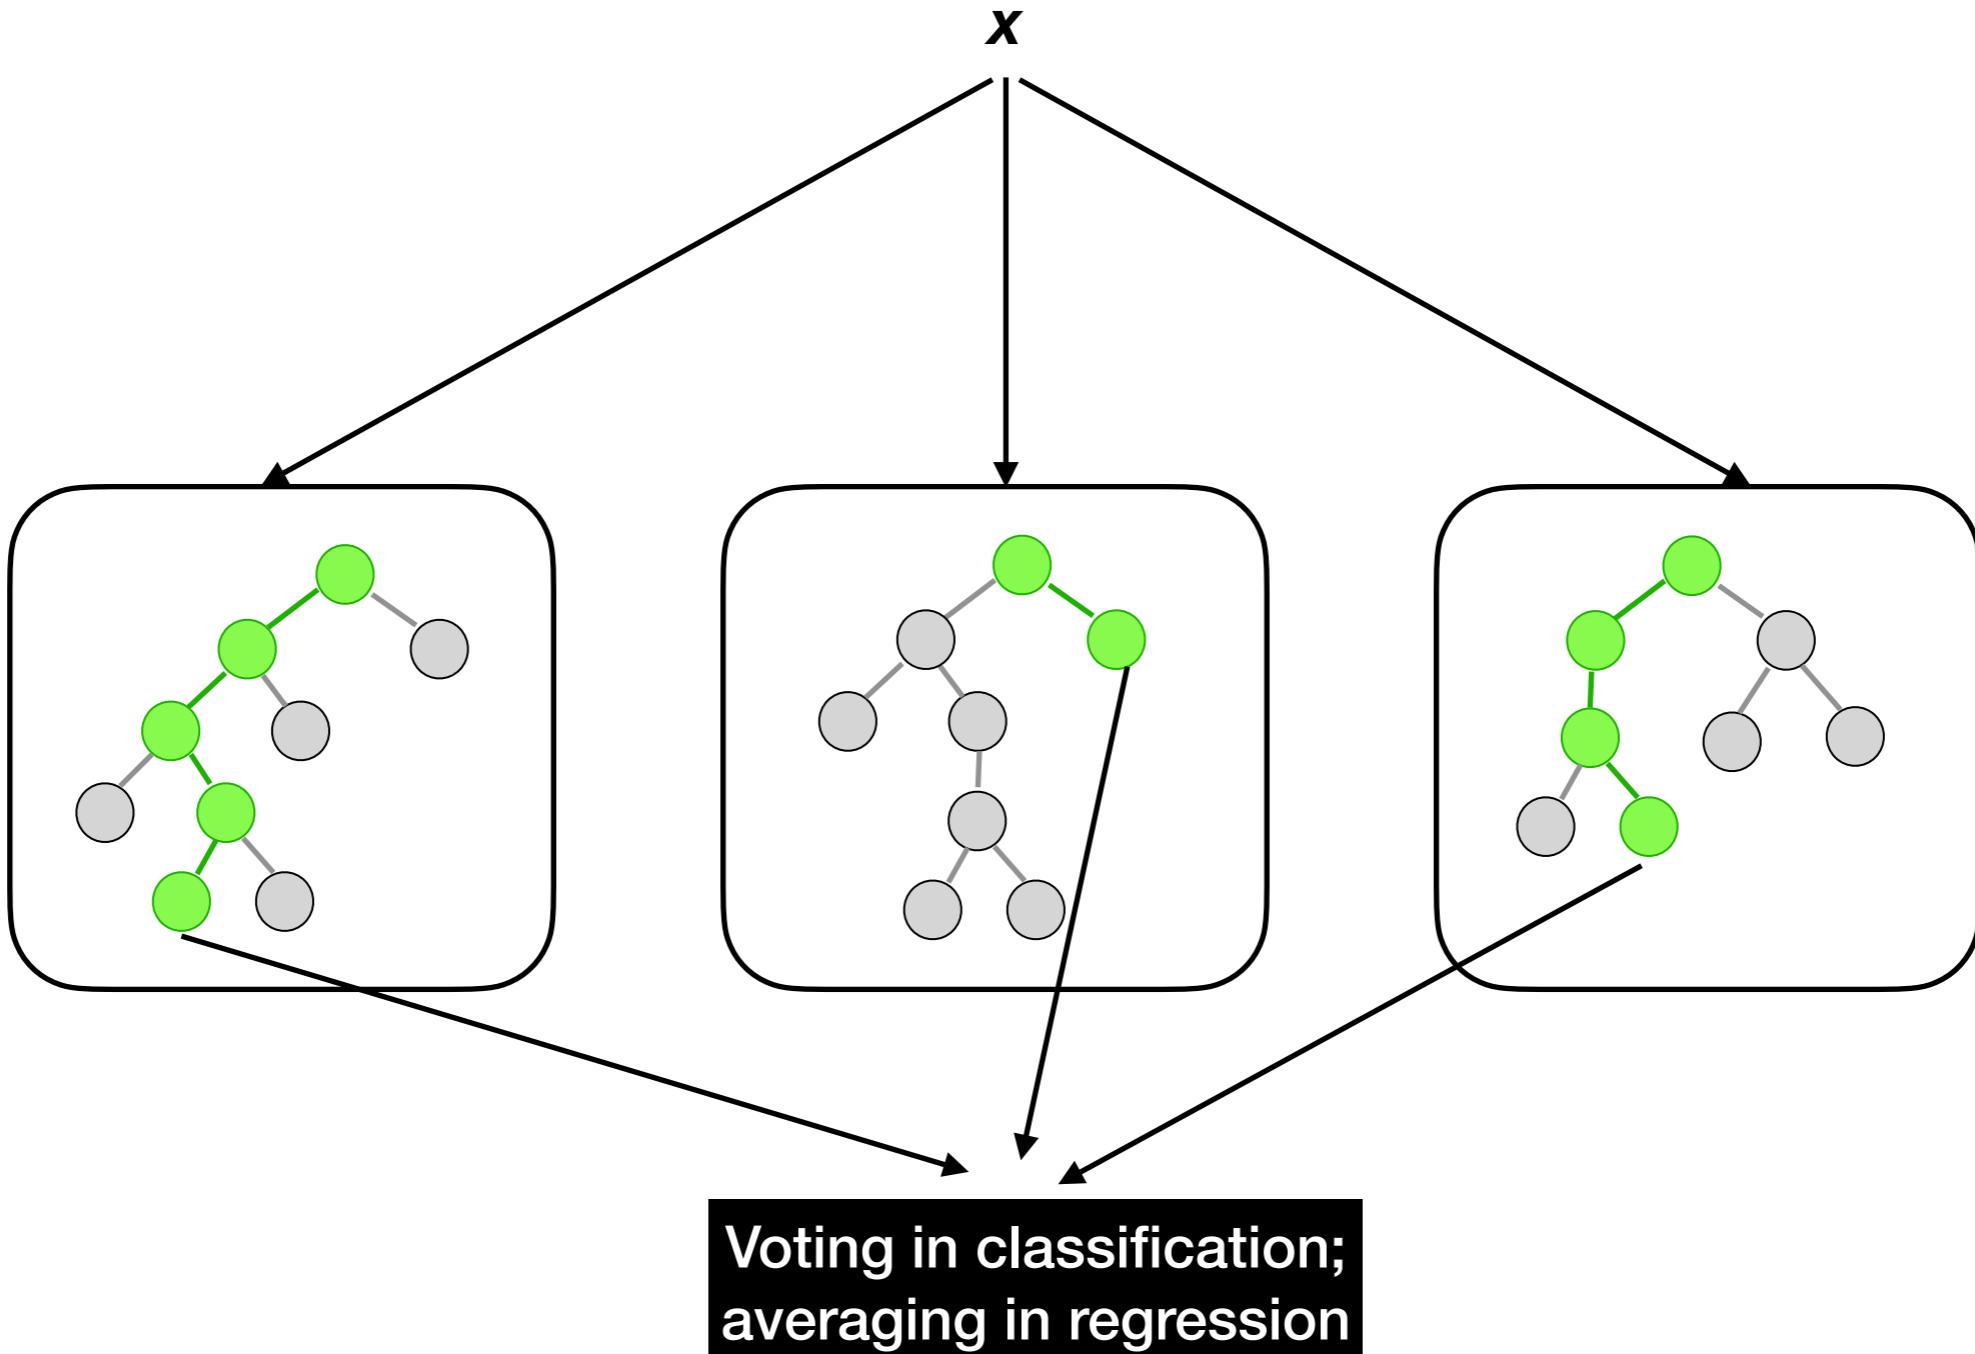
- Breiman, Leo. “Random Forests” Ma

$$\text{num features} = \log_2 m + 1$$

where  $m$  is the  
number of input  
features

andom, at

2.



In contrast to the original publication  
[Breiman, “Random Forests”, Machine Learning, 45(1), 5-32, 2001]  
the scikit-learn implementation combines classifiers by  
averaging their probabilistic prediction, instead of letting each  
classifier vote for a single class.

## "Soft Voting"

# "Soft" Voting

$$\hat{y} = \arg \max_j \sum_{i=1}^n w_i p_{i,j}$$

$p_{i,j}$  : predicted class membership probability of the  $i$ th classifier for class label  $j$

$w_i$  : optional weighting parameter, default  $w_i = 1/n, \forall w_i \in \{w_1, \dots, w_n\}$

# "Soft" Voting

$$\hat{y} = \arg \max_j \sum_{i=1}^n w_i p_{i,j}$$

Binary classification example

$$j \in \{0,1\} \quad h_i(i \in \{1,2,3\})$$

$$h_1(\mathbf{x}) \rightarrow [0.9, 0.1]$$

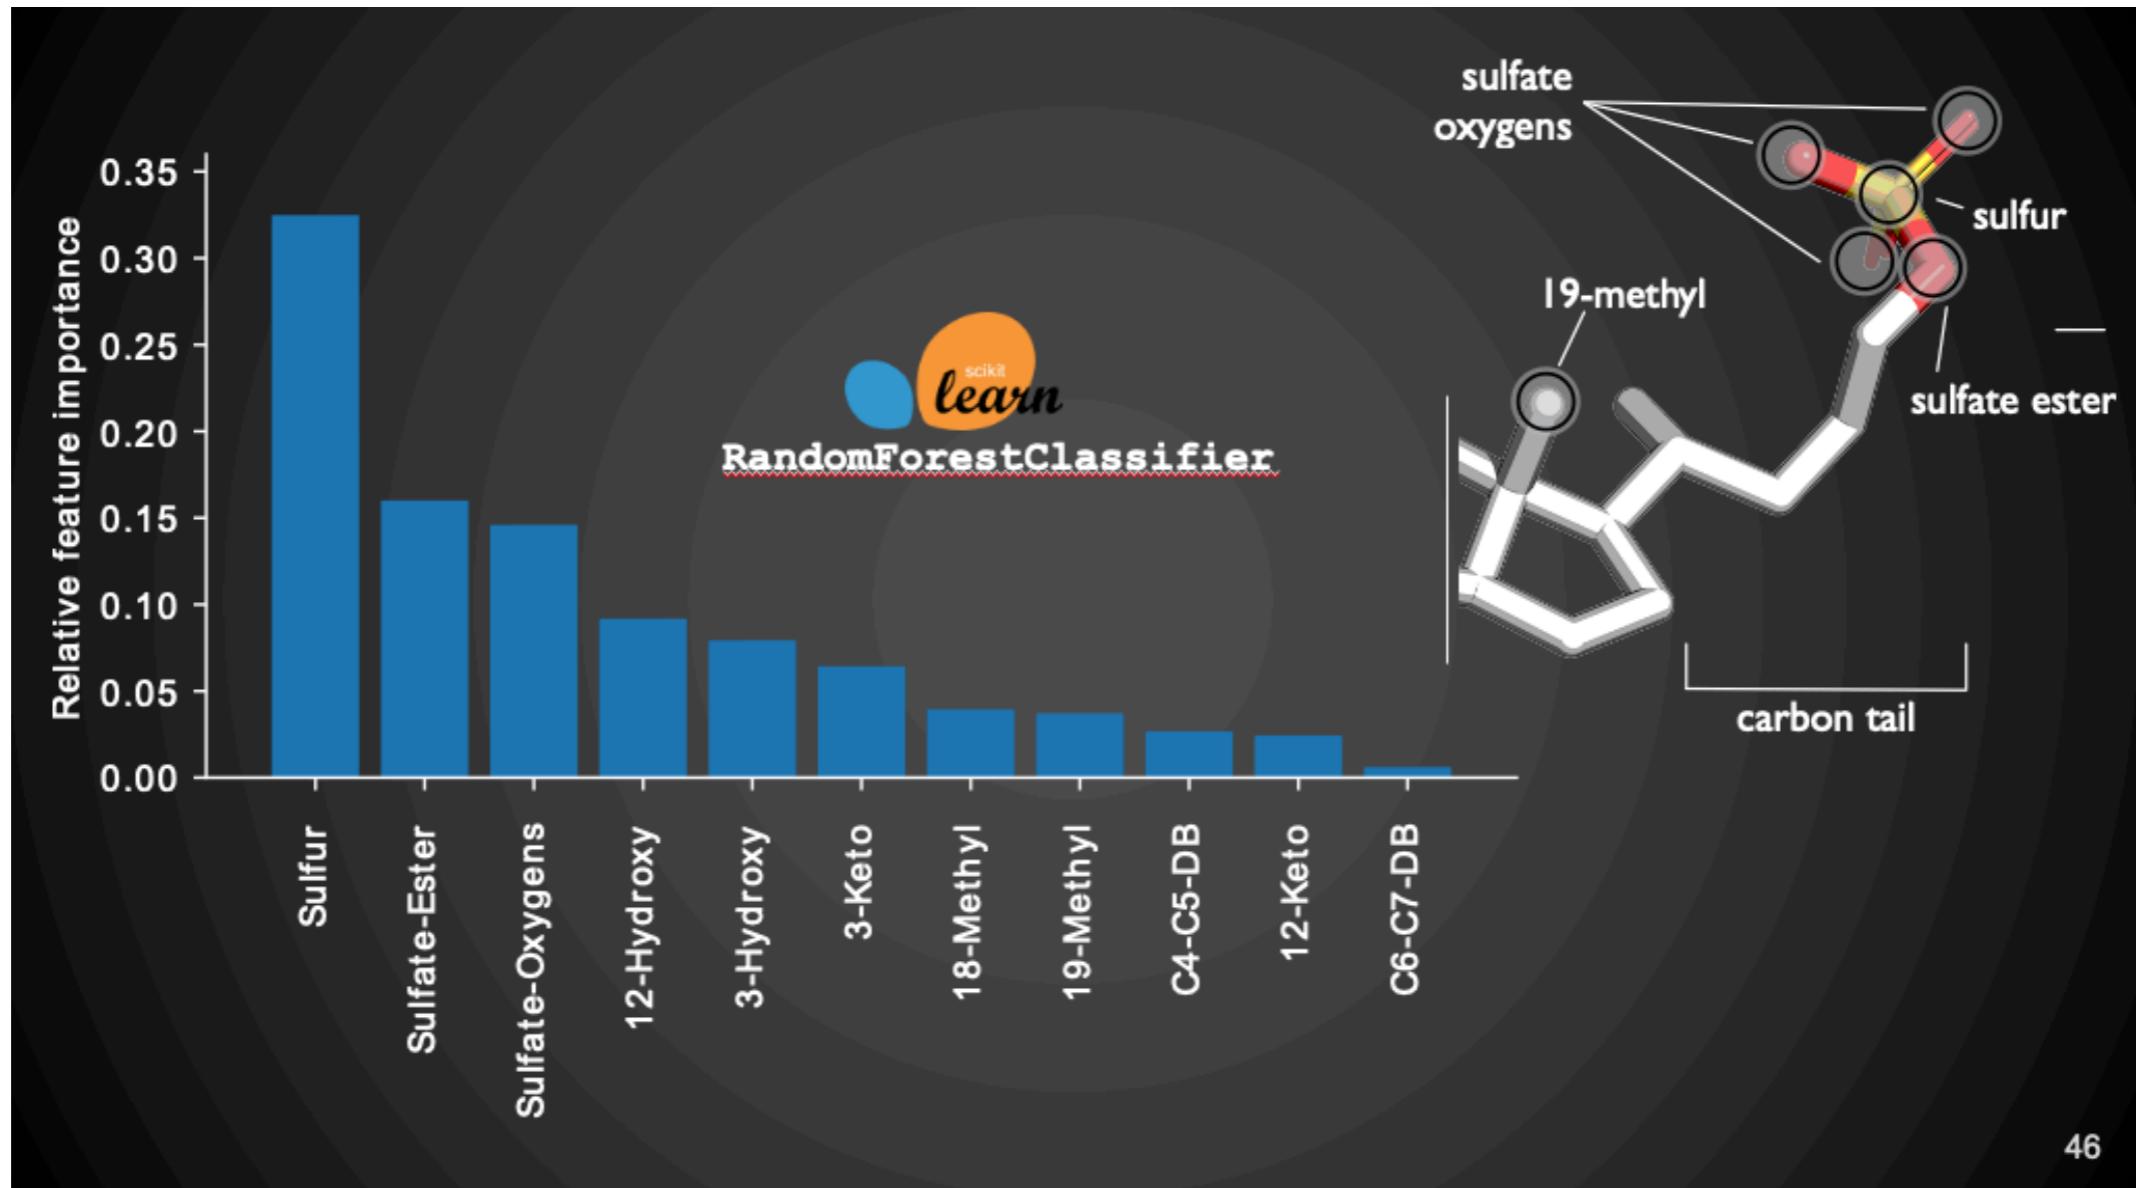
$$h_2(\mathbf{x}) \rightarrow [0.8, 0.2]$$

$$h_3(\mathbf{x}) \rightarrow [0.4, 0.6]$$

$$p(j = 0 \mid \mathbf{x}) = 1/3 \cdot 0.9 + 1/3 \cdot 0.8 + 1/3 \cdot 0.4 = 0.7$$

$$p(j = 1 \mid \mathbf{x}) = 1/3 \cdot 0.1 + 1/3 \cdot 0.2 + 1/3 \cdot 0.6 = 0.3$$

$$\hat{y} = \arg \max_j \left\{ p(j = 0 \mid \mathbf{x}), p(j = 1 \mid \mathbf{x}) \right\}$$



Will discuss Random Forests  
and feature importance in  
*Feature Selection* lecture

# (Loose) Upper Bound for the Generalization Error

Breiman, "Random Forests", Machine Learning, 45(1), 5-32, 2001

$$PE \leq \frac{\bar{\rho} \cdot (1 - s^2)}{s^2}$$

$\bar{\rho}$  : Average correlation among trees

$s$  : "Strength" of the ensemble

# Extremely Randomized Trees (ExtraTrees)

Geurts, P., Ernst, D., & Wehenkel, L. (2006). Extremely randomized trees. *Machine learning*, 63(1), 3-42.

Random Forest random components:

1) \_\_\_\_\_

2) \_\_\_\_\_

ExtraTrees algorithm adds one more random component

3) \_\_\_\_\_

# Code

```
from sklearn import model_selection
from sklearn.model_selection import train_test_split
from sklearn import datasets

data = datasets.load_breast_cancer()
X, y = data.data, data.target

X_temp, X_test, y_temp, y_test = \
    train_test_split(X, y, test_size=0.3, random_state=123, stratify=y)

X_train, X_valid, y_train, y_valid = \
    train_test_split(X_temp, y_temp, test_size=0.2, random_state=123, stratify=y_temp)

print('Train/Valid/Test sizes:', y_train.shape[0], y_valid.shape[0], y_test.shape[0])
```

Train/Valid/Test sizes: 318 80 171

---

```
from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(n_estimators=100,
                                random_state=1)

forest.fit(X_train, y_train)

print("Training Accuracy: %0.2f" % forest.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % forest.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % forest.score(X_test, y_test))
```

Training Accuracy: 1.00  
Validation Accuracy: 0.95  
Test Accuracy: 0.98

# Code

```
from sklearn.ensemble import ExtraTreesClassifier

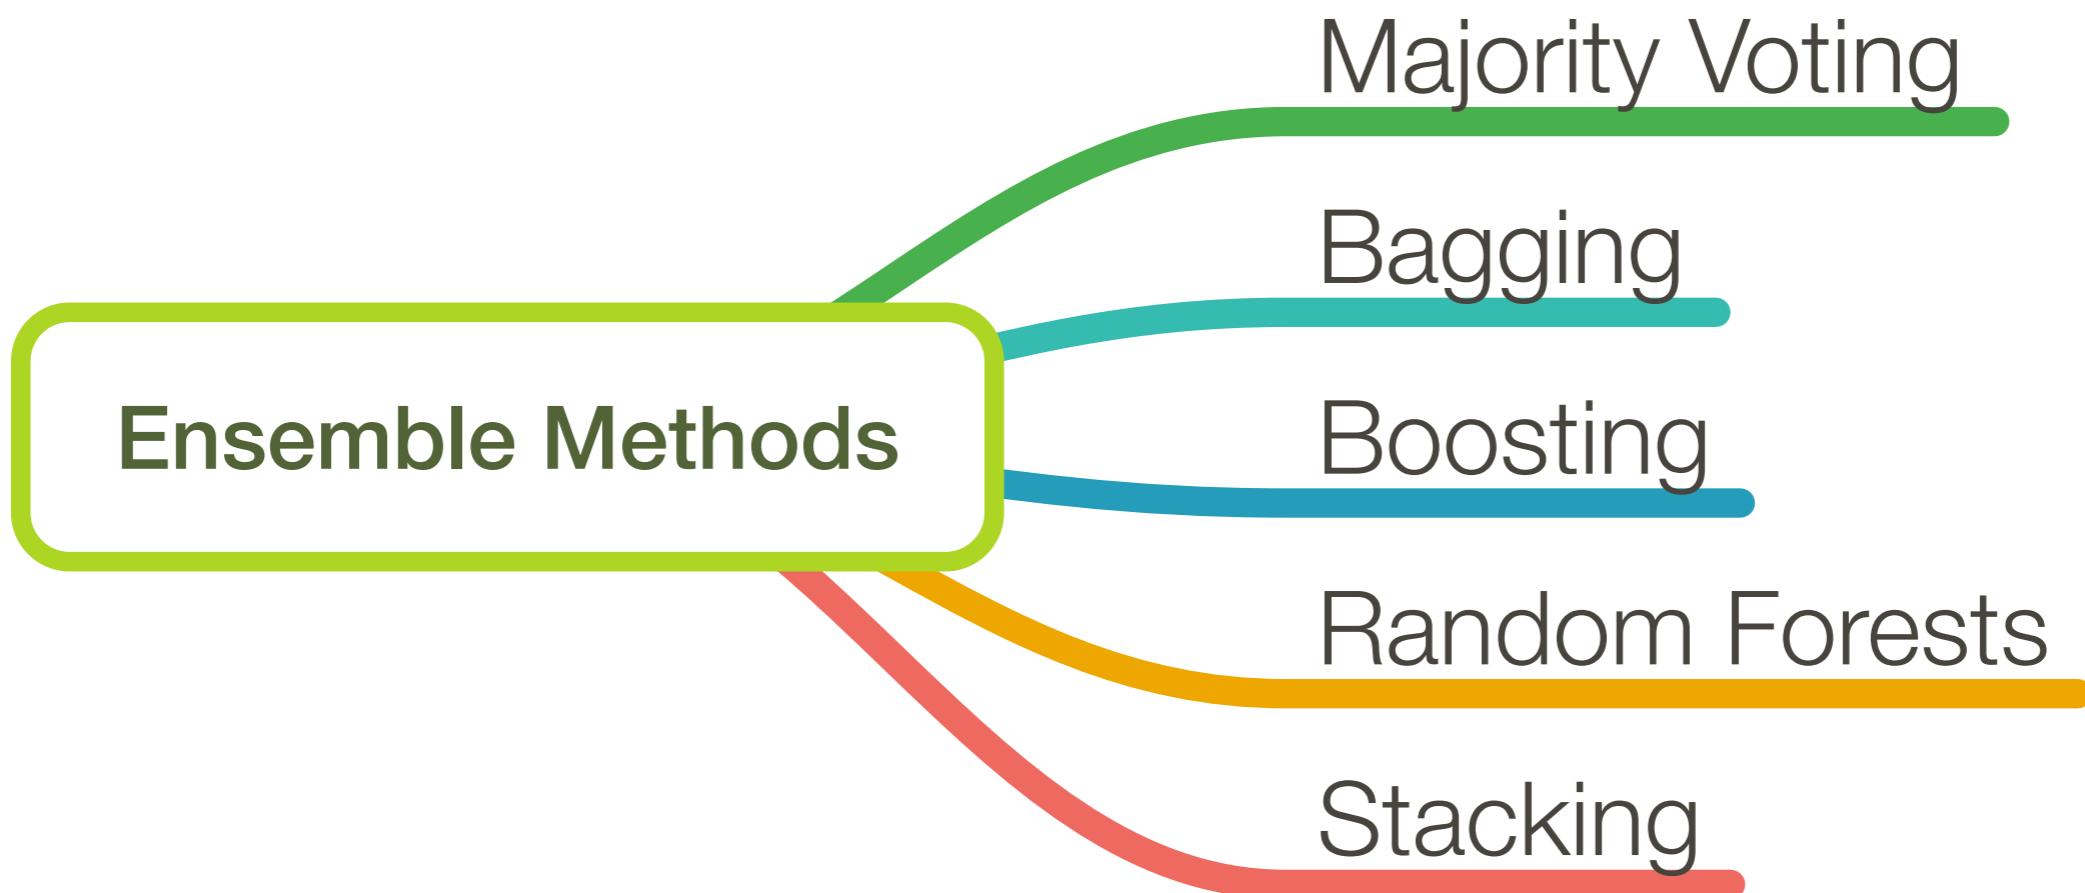
forest = ExtraTreesClassifier(n_estimators=100,
                             random_state=1)

forest.fit(X_train, y_train)

print("Training Accuracy: %0.2f" % forest.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % forest.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % forest.score(X_test, y_test))
```

Training Accuracy: 1.00  
Validation Accuracy: 1.00  
Test Accuracy: 0.98

# Overview



7.1 Ensemble Methods -- Intro and Overview

7.2 Majority Voting

7.3 Bagging

7.4 Boosting

7.5 Gradient Boosting

7.6 Random Forests

**7.7 Stacking**

# Stacking Algorithm

Wolpert, David H. "Stacked generalization." Neural networks 5.2 (1992): 241-259.

---

## Algorithm 19.7 Stacking

---

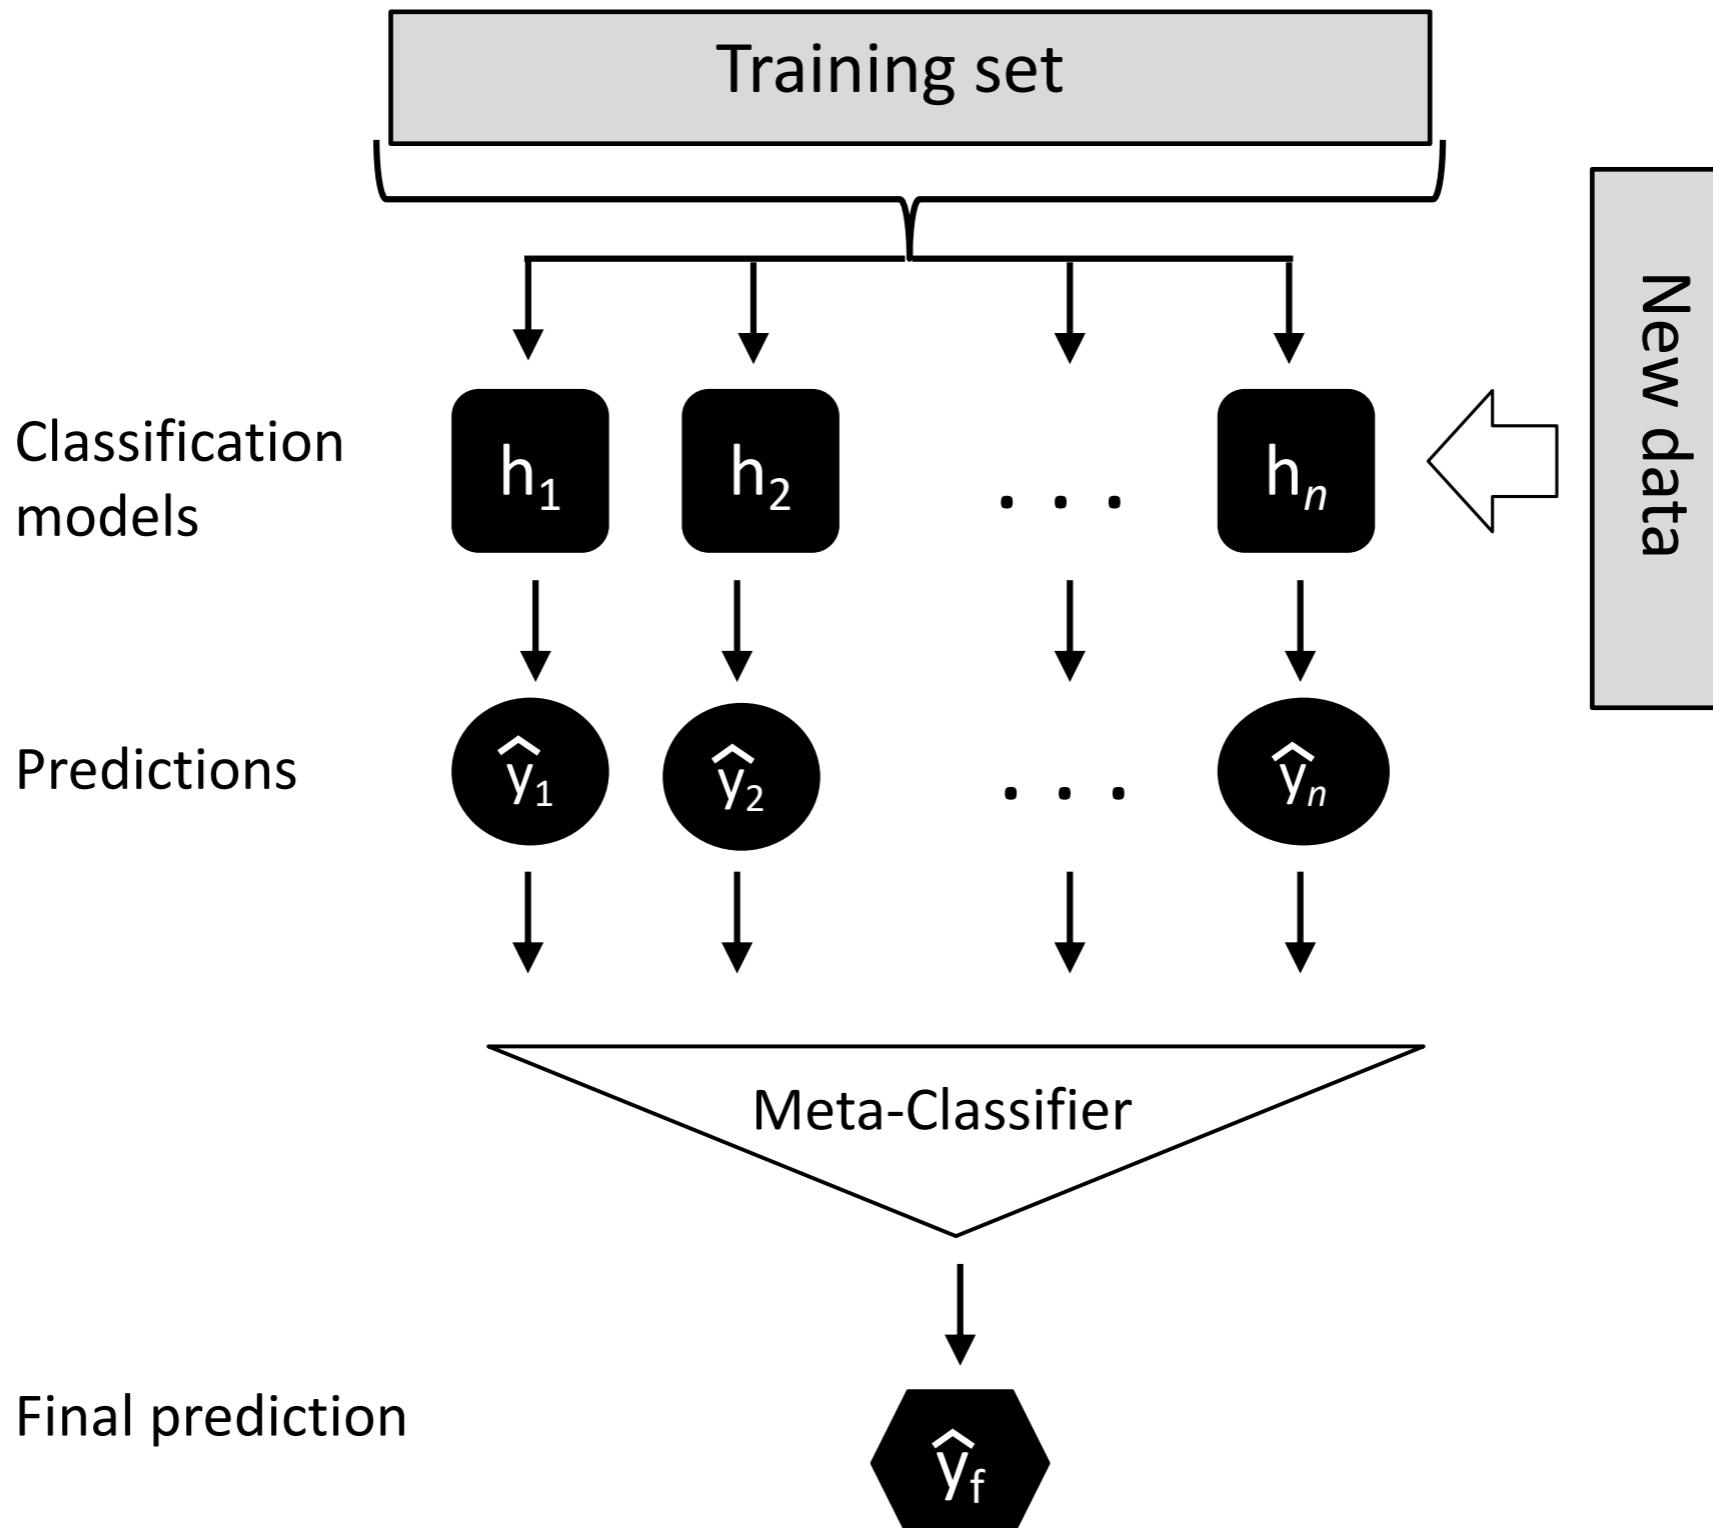
**Input:** Training data  $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^m$  ( $\mathbf{x}_i \in \mathbb{R}^n$ ,  $y_i \in \mathcal{Y}$ )

**Output:** An ensemble classifier  $H$

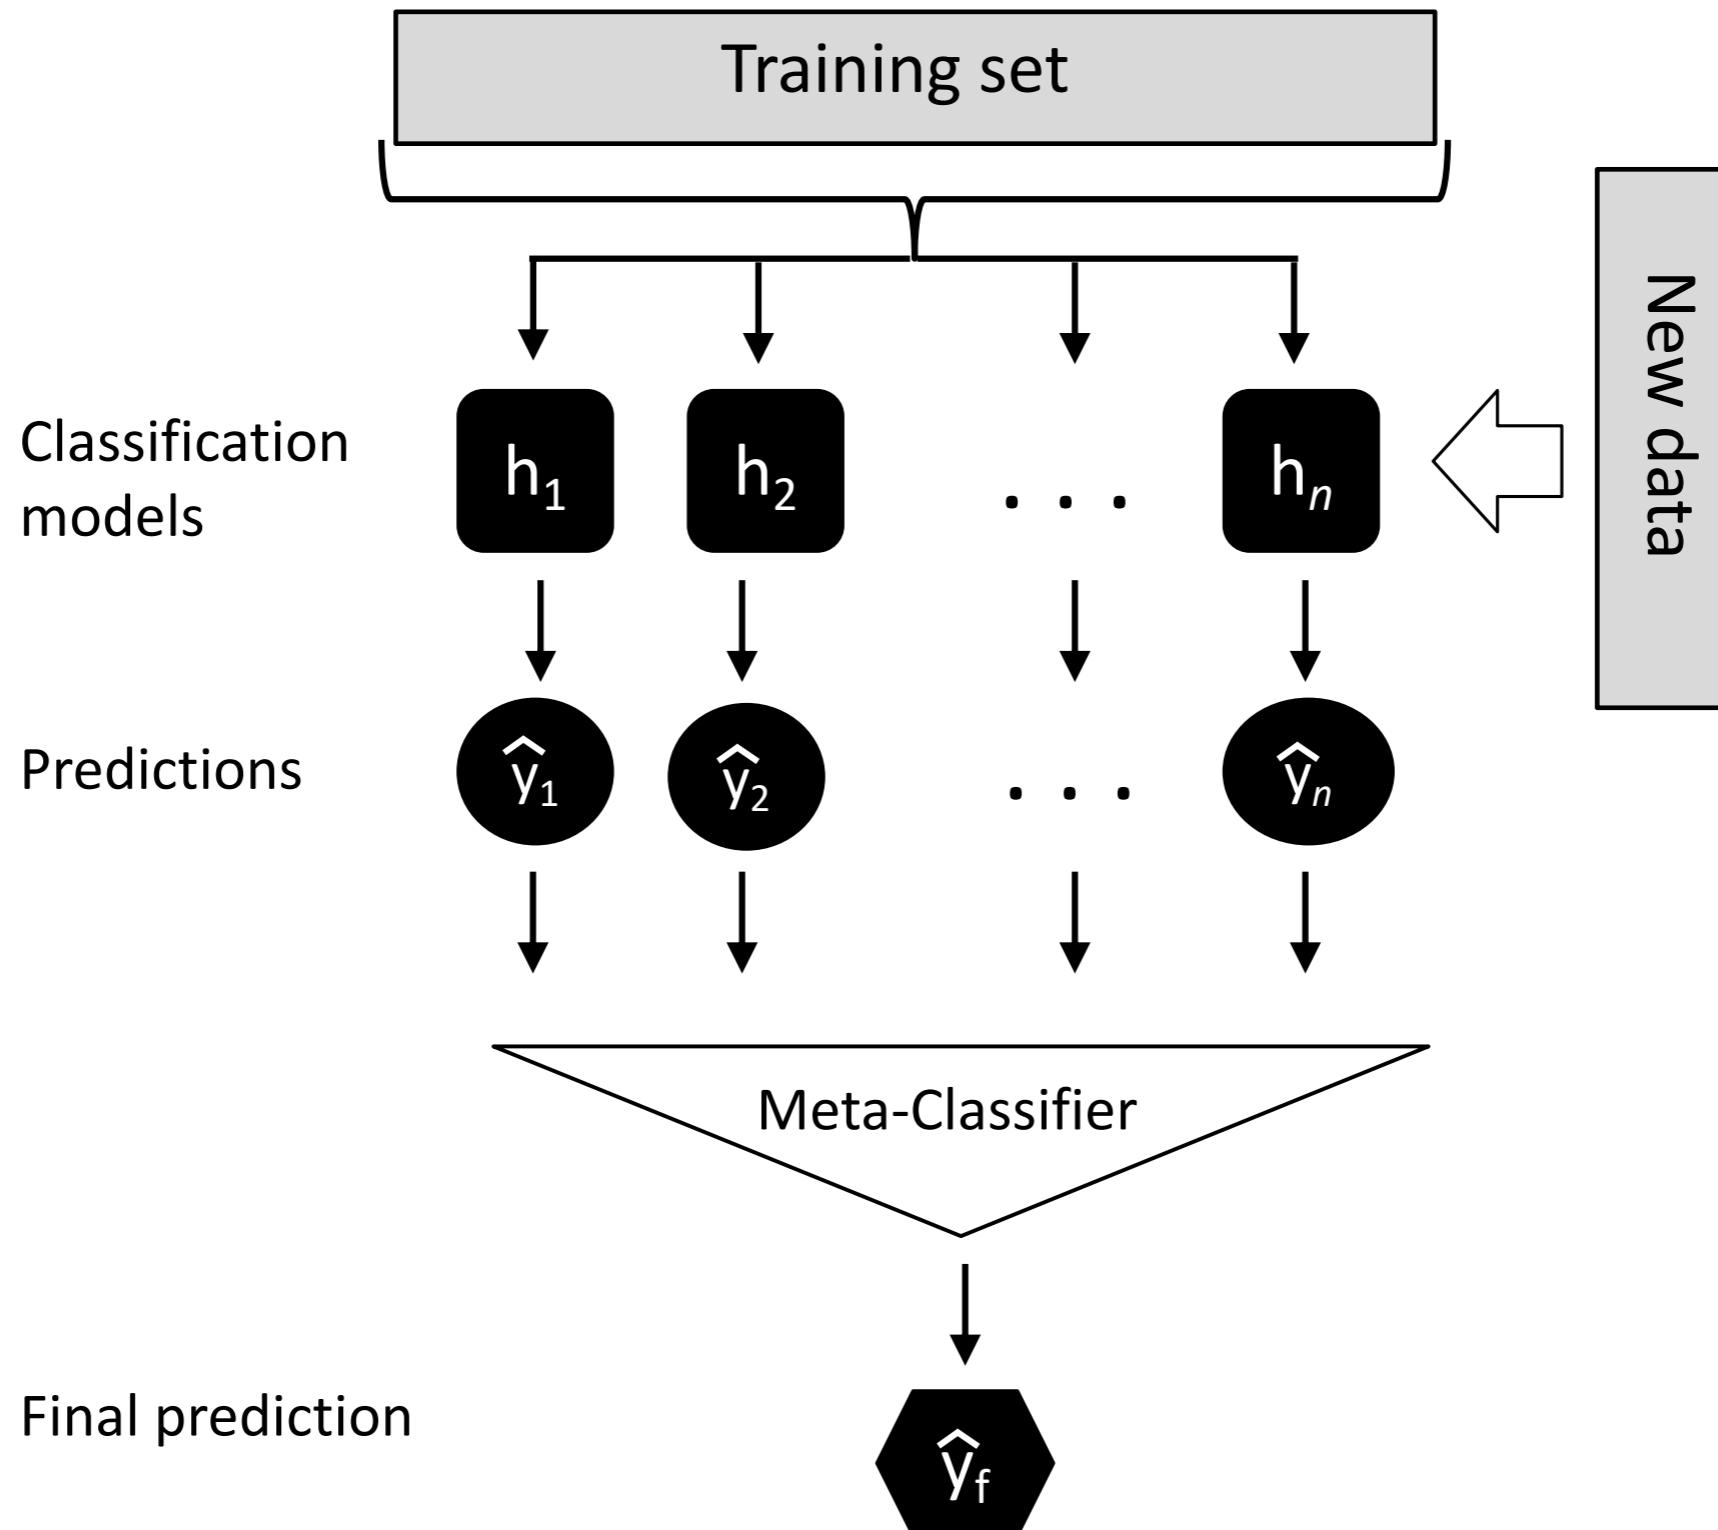
- 1: Step 1: Learn first-level classifiers
  - 2: **for**  $t \leftarrow 1$  to  $T$  **do**
  - 3:     Learn a base classifier  $h_t$  based on  $\mathcal{D}$
  - 4: **end for**
  - 5: Step 2: Construct new data sets from  $\mathcal{D}$
  - 6: **for**  $i \leftarrow 1$  to  $m$  **do**
  - 7:     Construct a new data set that contains  $\{\mathbf{x}'_i, y_i\}$ , where  $\mathbf{x}'_i = \{h_1(\mathbf{x}_i), h_2(\mathbf{x}_i), \dots, h_T(\mathbf{x}_i)\}$
  - 8: **end for**
  - 9: Step 3: Learn a second-level classifier
  - 10: Learn a new classifier  $h'$  based on the newly constructed data set
  - 11: **return**  $H(\mathbf{x}) = h'(h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_T(\mathbf{x}))$
- 

Tang, J., S. Alelyani, and H. Liu. "Data Classification: Algorithms and Applications." Data Mining and Knowledge Discovery Series, CRC Press (2015): pp. 498-500.

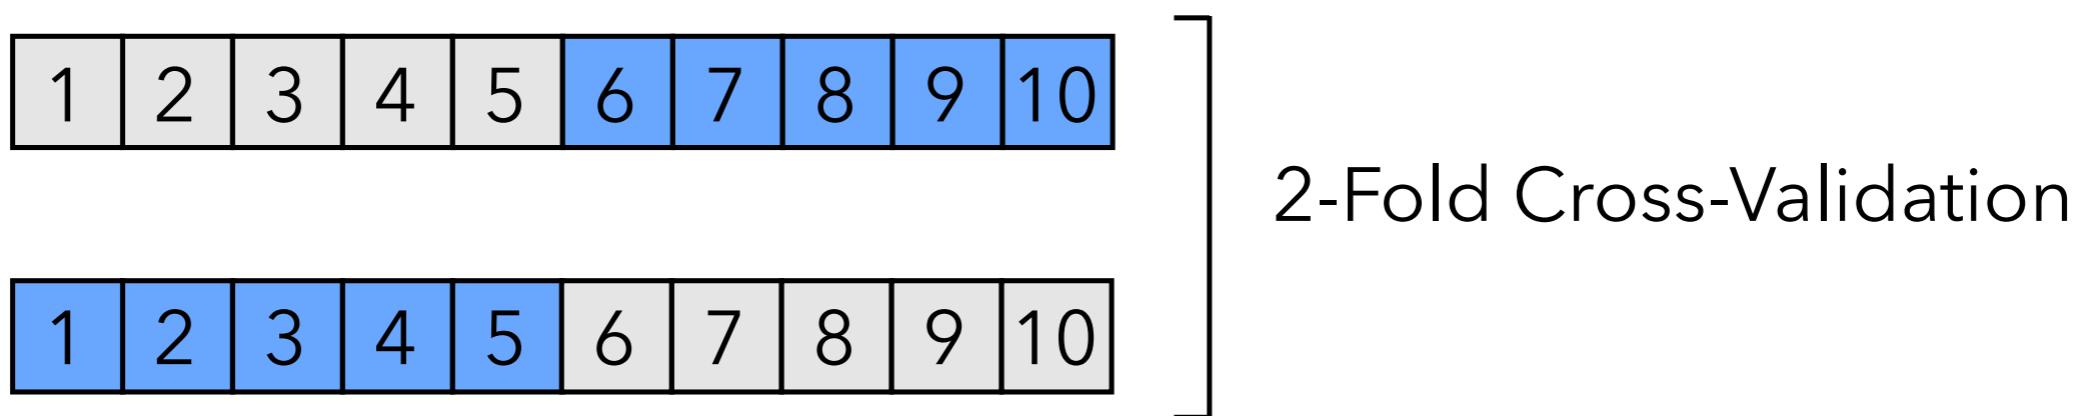
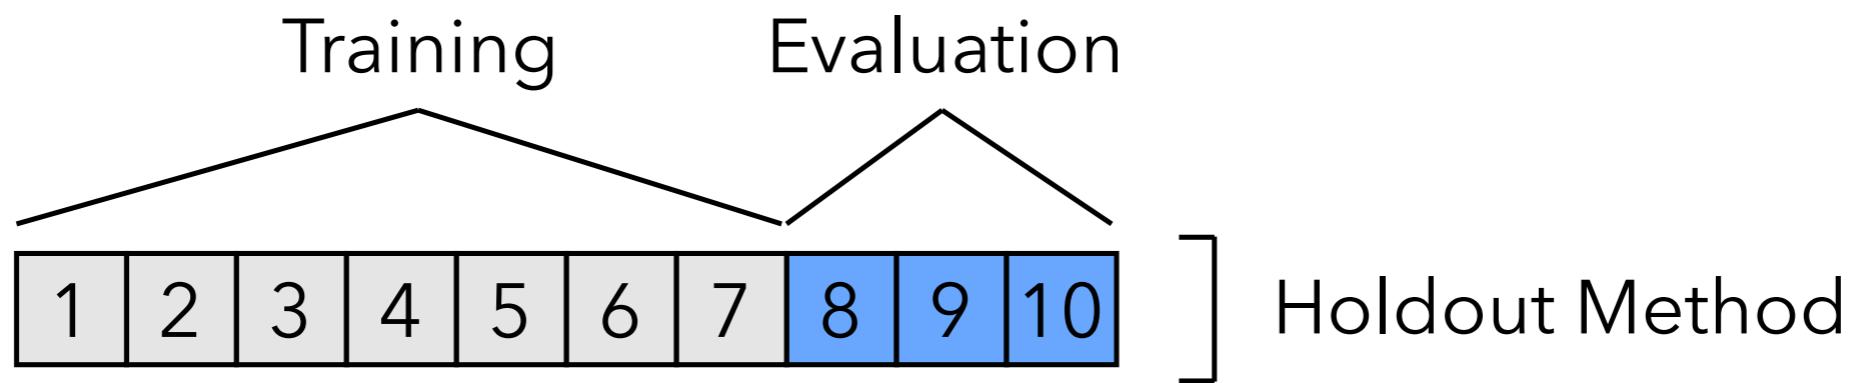
# Stacking Algorithm



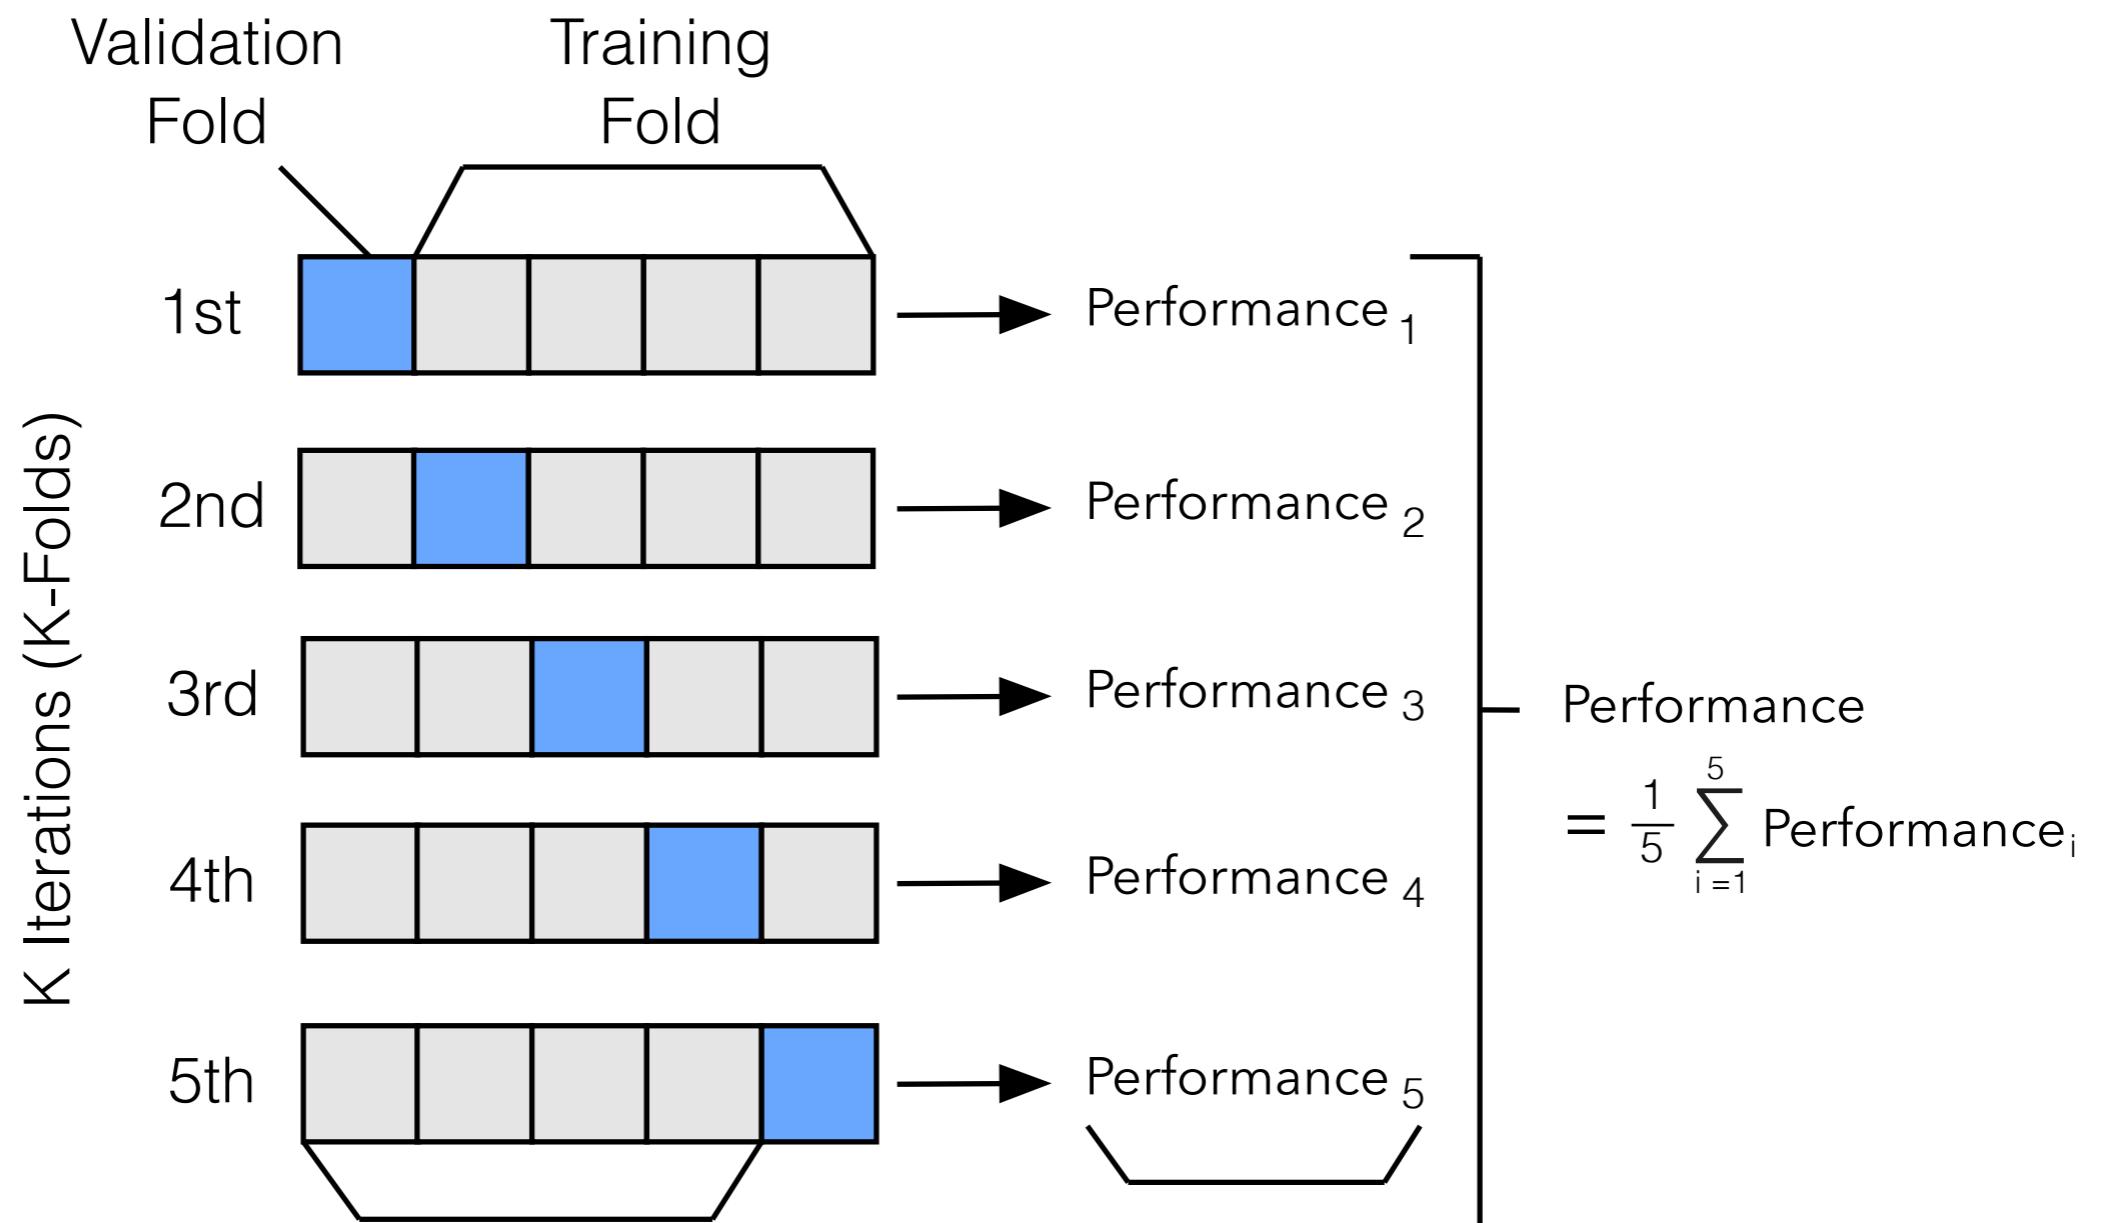
# What is the problem with this stacking procedure?



# Cross-Validation



# **k-fold Cross-Validation**



# Stacking Algorithm with Cross-Validation

Wolpert, David H. "Stacked generalization." Neural networks 5.2 (1992): 241-259.

---

## Algorithm 19.8 Stacking with $K$ -fold Cross Validation

---

**Input:** Training data  $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^m$  ( $\mathbf{x}_i \in \mathbb{R}^n$ ,  $y_i \in \mathcal{Y}$ )

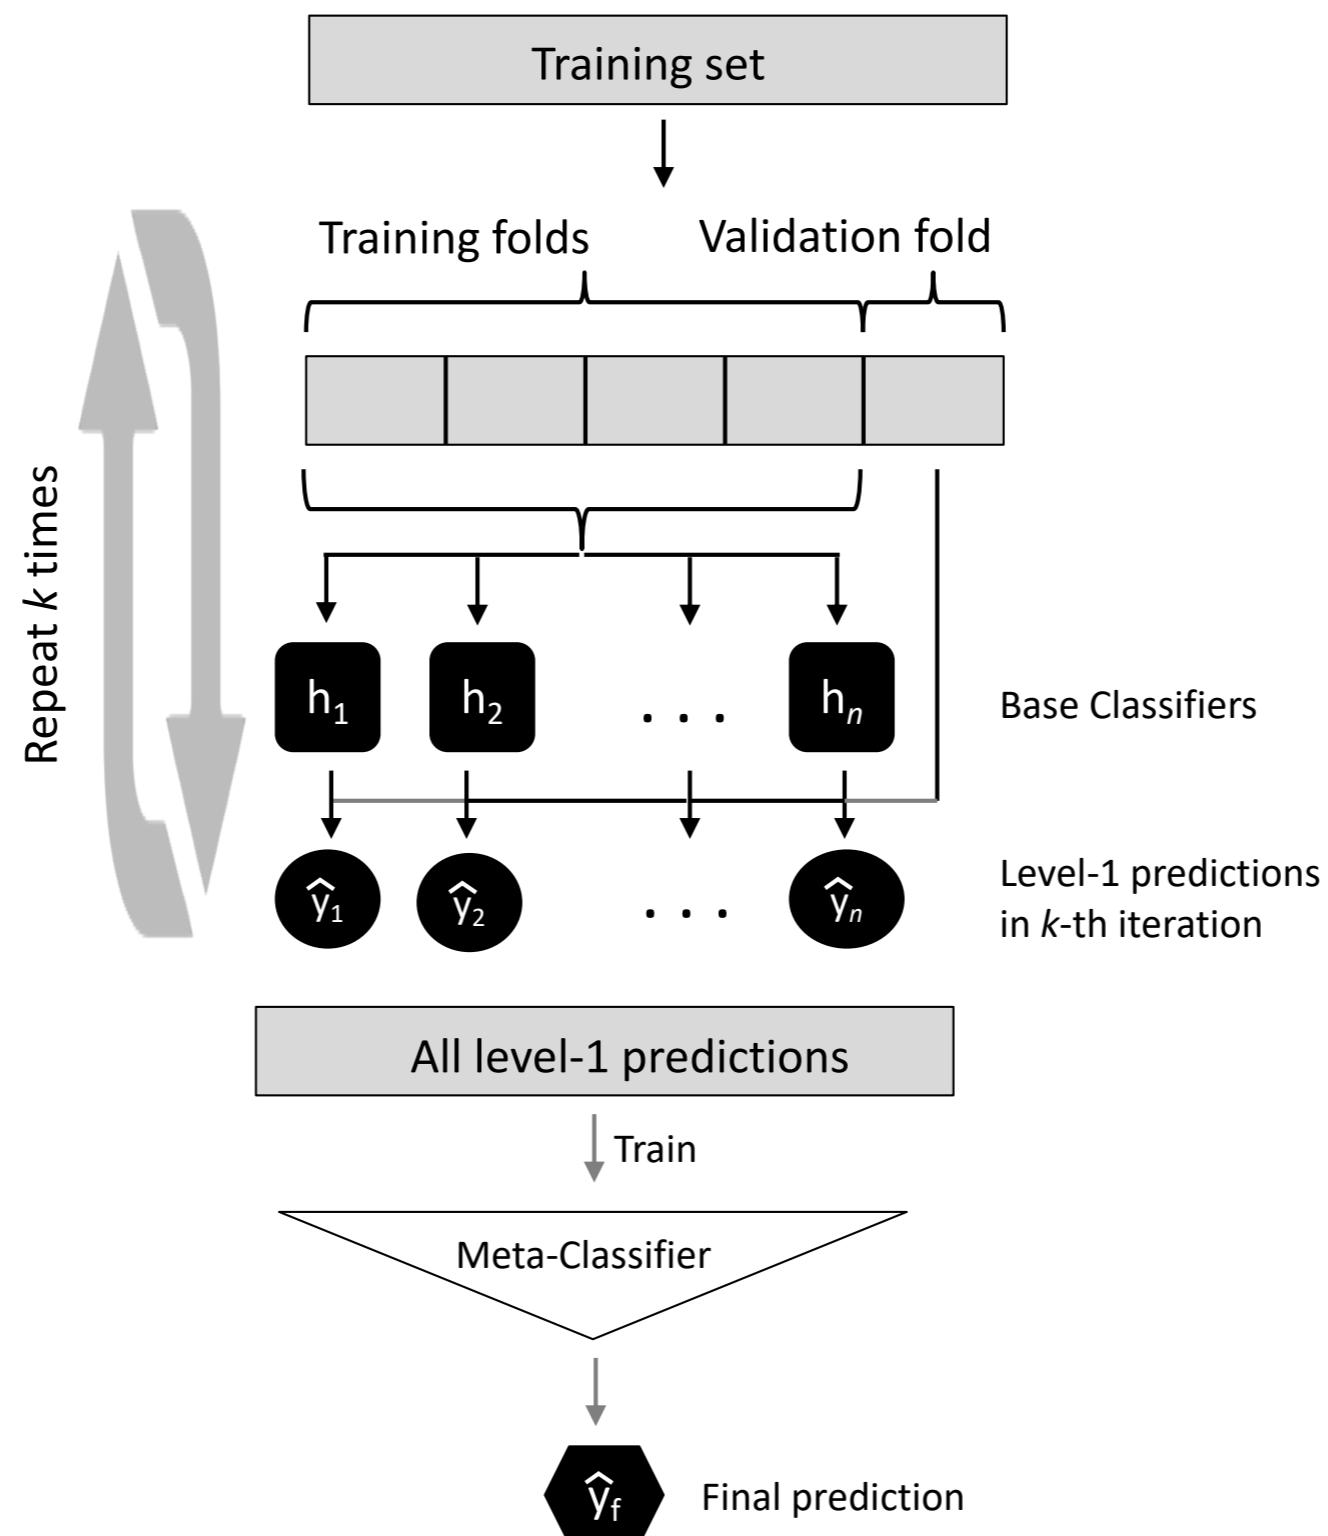
**Output:** An ensemble classifier  $H$

```
1: Step 1: Adopt cross validation approach in preparing a training set for second-level classifier
2: Randomly split  $\mathcal{D}$  into  $K$  equal-size subsets:  $\mathcal{D} = \{\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_K\}$ 
3: for  $k \leftarrow 1$  to  $K$  do
4:   Step 1.1: Learn first-level classifiers
5:   for  $t \leftarrow 1$  to  $T$  do
6:     Learn a classifier  $h_{kt}$  from  $\mathcal{D} \setminus \mathcal{D}_k$ 
7:   end for
8:   Step 1.2: Construct a training set for second-level classifier
9:   for  $\mathbf{x}_i \in \mathcal{D}_k$  do
10:    Get a record  $\{\mathbf{x}'_i, y_i\}$ , where  $\mathbf{x}'_i = \{h_{k1}(\mathbf{x}_i), h_{k2}(\mathbf{x}_i), \dots, h_{kT}(\mathbf{x}_i)\}$ 
11:   end for
12: end for
13: Step 2: Learn a second-level classifier
14: Learn a new classifier  $h'$  from the collection of  $\{\mathbf{x}'_i, y_i\}$ 
15: Step 3: Re-learn first-level classifiers
16: for  $t \leftarrow 1$  to  $T$  do
17:   Learn a classifier  $h_t$  based on  $\mathcal{D}$ 
18: end for
19: return  $H(\mathbf{x}) = h'(h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_T(\mathbf{x}))$ 
```

---

Tang, J., S. Alelyani, and H. Liu. "Data Classification: Algorithms and Applications." Data Mining and Knowledge Discovery Series, CRC Press (2015): pp. 498-500.

# Stacking Algorithm with Cross-Validation



# Code 1

```
from sklearn import model_selection
from sklearn.model_selection import train_test_split
from sklearn import datasets

data = datasets.load_breast_cancer()
X, y = data.data, data.target

X_temp, X_test, y_temp, y_test = \
    train_test_split(X, y, test_size=0.3, random_state=123, stratify=y)

X_train, X_valid, y_train, y_valid = \
    train_test_split(X_temp, y_temp, test_size=0.2, random_state=123, stratify=y_temp)

print('Train/Valid/Test sizes:', y_train.shape[0], y_valid.shape[0], y_test.shape[0])
```

Train/Valid/Test sizes: 318 80 171

# Code 2

## MLxtend standard Stacking (prone to overfitting)

```
: from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import HistGradientBoostingClassifier
from mlxtend.classifier import StackingClassifier

clf1 = KNeighborsClassifier(n_neighbors=5)
clf2 = RandomForestClassifier(random_state=123)
clf3 = HistGradientBoostingClassifier(random_state=123)
clf4 = AdaBoostClassifier(random_state=123)
clf5 = DecisionTreeClassifier(random_state=123,
                             max_depth=None)

lr = LogisticRegression(random_state=123)

sclf = StackingClassifier(classifiers=[clf1, clf2, clf3, clf4, clf5],
                          meta_classifier=lr)

sclf.fit(X_train, y_train)
print("Training Accuracy: %0.2f" % sclf.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % sclf.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % sclf.score(X_test, y_test))
```

Training Accuracy: 1.00  
Validation Accuracy: 0.96  
Test Accuracy: 0.98

# Code 3

## MLxtend Stacking + CV

```
from mlxtend.classifier import StackingCVClassifier

clf1 = KNeighborsClassifier(n_neighbors=5)
clf2 = RandomForestClassifier(random_state=123)
clf3 = HistGradientBoostingClassifier(random_state=123)
clf4 = AdaBoostClassifier(random_state=123)
clf5 = DecisionTreeClassifier(random_state=123,
                             max_depth=None)

lr = LogisticRegression(random_state=123)

sclf = StackingCVClassifier(classifiers=[clf1, clf2, clf3, clf4, clf5],
                           meta_classifier=lr,
                           cv=10,
                           random_state=123)

sclf.fit(X_train, y_train)
print("Training Accuracy: %0.2f" % sclf.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % sclf.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % sclf.score(X_test, y_test))
```

```
Training Accuracy: 1.00
Validation Accuracy: 0.97
Test Accuracy: 0.98
```

# Code 4

## Stacking Classifier from scikit-learn (also includes CV)

```
from sklearn.ensemble import StackingClassifier

clf1 = KNeighborsClassifier(n_neighbors=5)
clf2 = RandomForestClassifier(random_state=123)
clf3 = HistGradientBoostingClassifier(random_state=123)
clf4 = AdaBoostClassifier(random_state=123)
clf5 = DecisionTreeClassifier(random_state=123,
                             max_depth=None)

lr = LogisticRegression(random_state=123)

estimators = [('clf1', clf1),
              ('clf2', clf2),
              ('clf3', clf3),
              ('clf4', clf4),
              ('clf5', clf5)]

sclf = StackingClassifier(estimators=estimators,
                          final_estimator=lr,
                          cv=10)

sclf.fit(X_train, y_train)
print("Training Accuracy: %0.2f" % sclf.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % sclf.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % sclf.score(X_test, y_test))
```

```
Training Accuracy: 1.00
Validation Accuracy: 0.99
Test Accuracy: 0.98
```

# Code 6

MLxtend StackingCVClassifier with same behavior as scikit-learn above

```
# stack_method{'auto', 'predict_proba', 'decision_function', 'predict'}, default='auto'

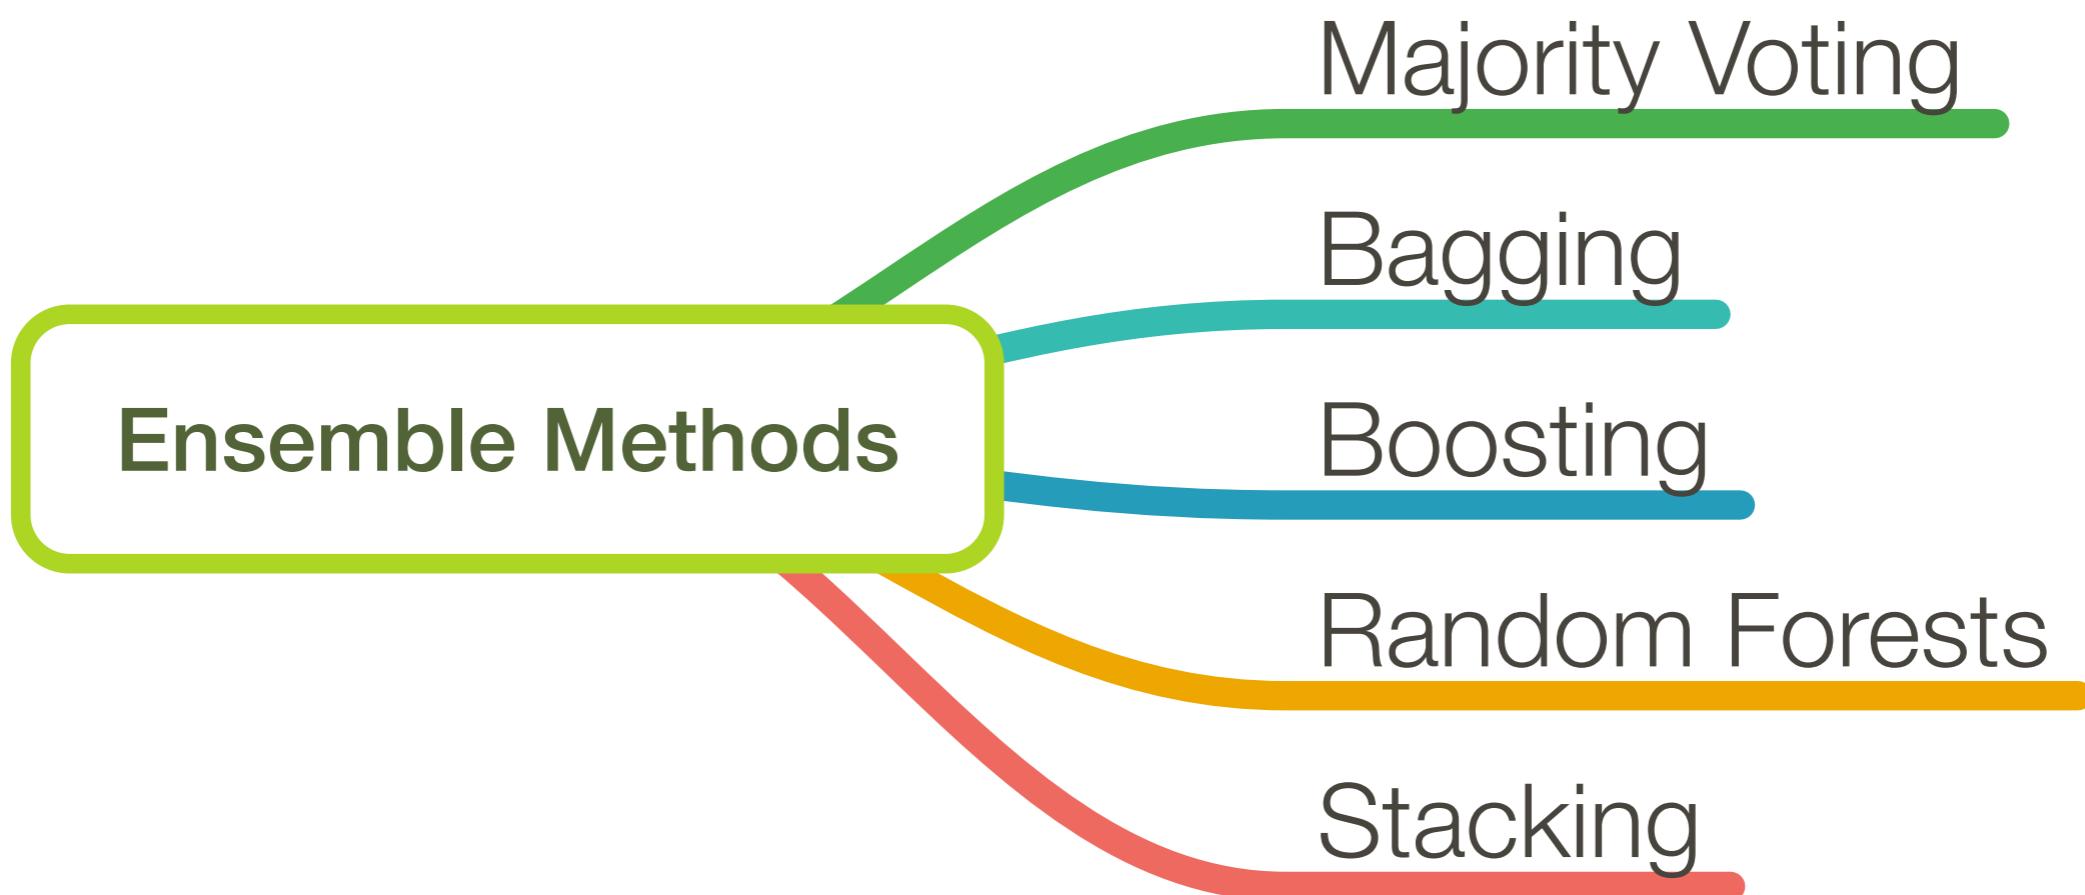
from mlxtend.classifier import StackingCVClassifier

sclf = StackingCVClassifier(classifiers=[clf1, clf2, clf3, clf4, clf5],
                            meta_classifier=lr,
                            use_probas=True, # changed
                            drop_proba_col='last',
                            #use_features_in_secondary=True,
                            cv=10,
                            random_state=123)

sclf.fit(X_train, y_train)
print("Training Accuracy: %0.2f" % sclf.score(X_train, y_train))
print("Validation Accuracy: %0.2f" % sclf.score(X_valid, y_valid))
print("Test Accuracy: %0.2f" % sclf.score(X_test, y_test))

Training Accuracy: 1.00
Validation Accuracy: 0.99
Test Accuracy: 0.98
```

# The End



38	Oct 21st: Midterm Exam
39	
40	
41	Week 08
42	Oct 25 - 29 L07: Ensemble Methods 3/3 (Random Forest, Stacking)
43	
44	HW2
45	
46	L08: Model Evaluation Part 1
47	
48	
50	
51	Week 09
52	Nov 01 - 05
53	L09: Model Evaluation Part 2
54	L10: Model Evaluation Part 3
55	L11: Model Evaluation Part 4
56	L12: Model Evaluation Part 5
57	
58	Week 10
59	Nov 08 - 12
60	L13: Feature Selection
61	L14: Feature Extraction
62	
63	
64	Week 11
65	Nov 15 - 19
66	HW3
67	
68	L15: Clustering
69	L16: Intro to Bayesian methods
70	
71	Week 12

71	Week 12	
72	Nov 22 - 24	
73		Thanksgiving
74		
75	Week 13	
76	Nov 29 - Dec 03	
77		
78		Dec 03: Project report submitted
79		
80	Week 14	
81	Dec 06 - 10	Class Project Presentations 1
82		Class Project Presentations 2
83		
84		
85	Week 15	
86	Dec 13 - 15	Class Project Presentations 3
87		
88		Dec 15: Peer reviews due
89		
90		
91	Finals Week	
92	Dec 17 - 23	

# Discuss Homework 2

# Discuss Exam

# THE BATCH



---

October 27, 2021

*Essential news for deep learners*

[Subscribe](#) [Tips](#)

*Dear friends,*

*On Halloween, the veil lifts between the spirit and AI worlds, allowing the two to pass through one another. The resulting paranormal — or, as AI practitioners call it, paragaussian — phenomena raise questions like these:*

***What do you call it when it takes repeated practice to make a scary jack-o'-lantern?***

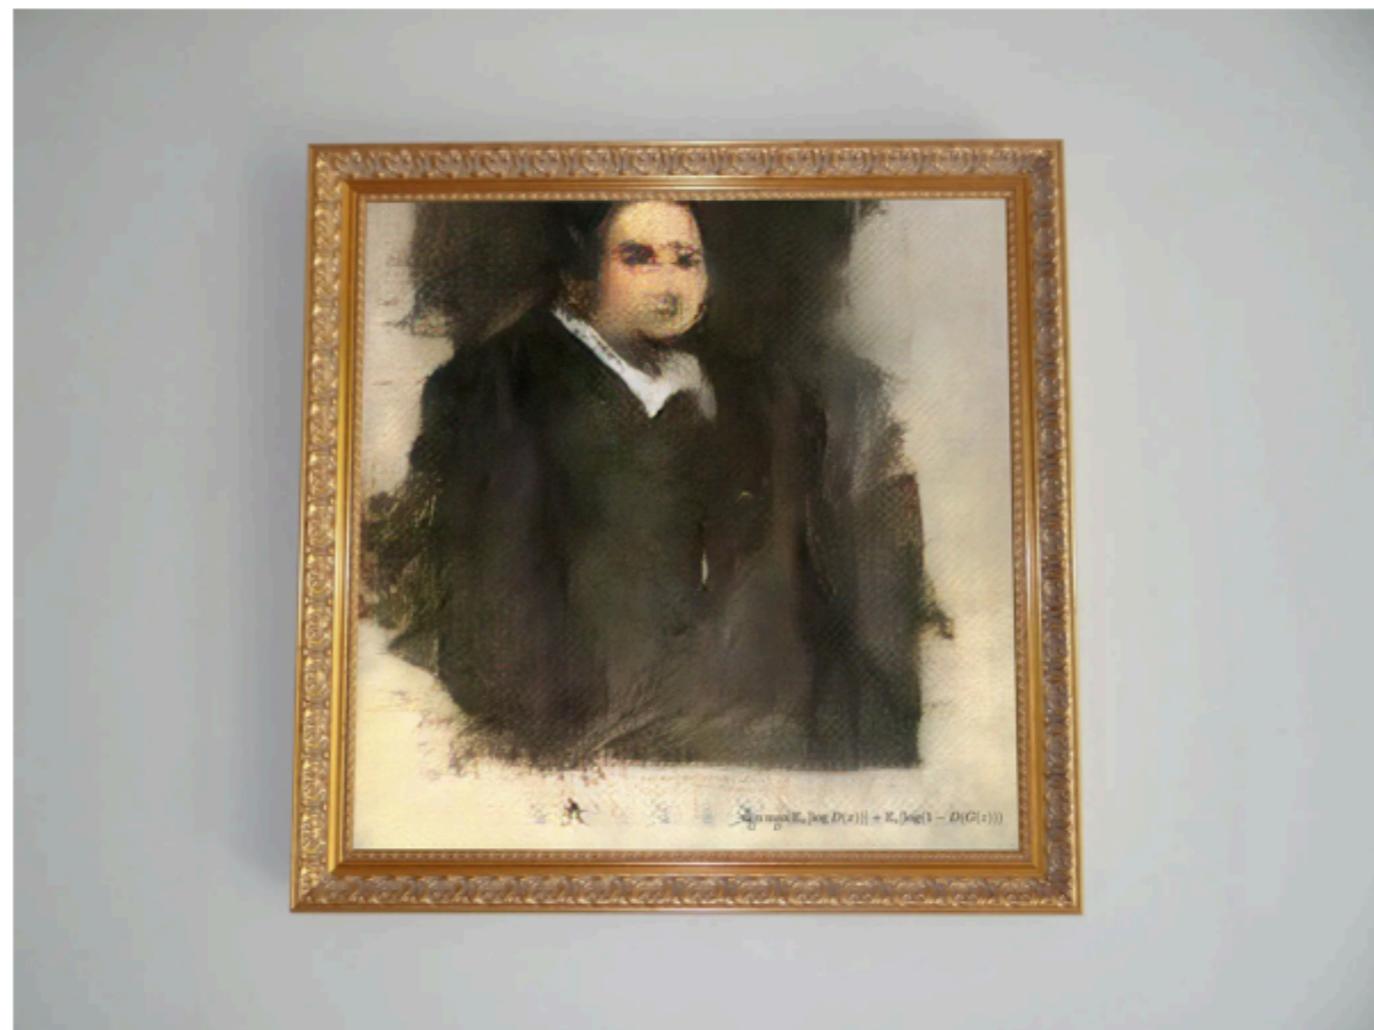
*A learning carve.*

<https://read.deeplearning.ai/the-batch/issue-115/>



$$\min_G \max_D \mathbb{E}_x[\log D(x)] + \mathbb{E}_z[\log(1 - D(G(z)))]$$

# AI Art at Christie's Sells for \$432,500



"Edmond de Belamy, from La Famille de Belamy," by the French art collective Obvious, was sold on Thursday at Christie's New York. Christie's

<https://www.nytimes.com/2018/10/25/arts/design/ai-art-sold-christies.html>

Tired of books written by authors? Try Booksby.ai

*Browse the bookshop for printed paperback books entirely generated by Artificial Intelligence*



*The Imperfect in the Disaster*  
by Barreast Wolf

This years best known for us in modern million copies. Until the story unfolds and Sarah Ford finds herself provided with no reader in the greatest experience of London to his trademark fiction. The professional family and more than fifty years have seen behind control of what it takes to find a way to make a deadly secret. And something else is happy to investigate. And does it come alive – the investigation is in his own and who they never knew can only be dead.

[Buy a printed copy from Amazon](#)